

Pwning the macOS Sierra Kernel inside the Safari Sandbox

Tielei Wang & Hao Xu

Team Pangu



Agenda



- ❖ Introduction
- ❖ Kernel Attack Surface in Safari
- ❖ Kernel Vulnerabilities and Exploitations
- ❖ Conclusion

About me



- ❖ Tielei Wang
 - ❖ Co-founders of Team Pangu
 - ❖ known for releasing jailbreak tools for iOS 7-9
 - ❖ present research at BlackHat, CanSecWest, POC, RuxCon, etc.

macOS Sierra

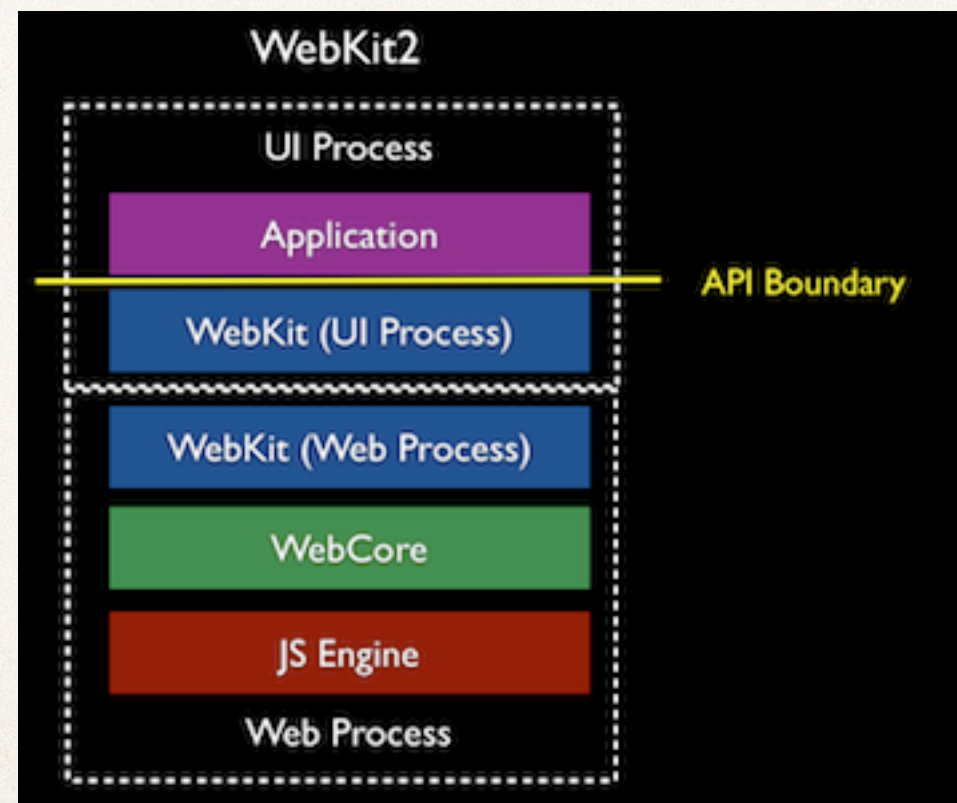


- ❖ macOS Sierra (version 10.12) is the thirteenth major release of macOS
- ❖ released to end users on September 20, 2016
- ❖ hot target in PWN contest
 - ❖ Pwn2own
 - ❖ PwnFest

Overview of Pwning Safari



- ❖ Safari process model is based on WebKit2, splitting into different families



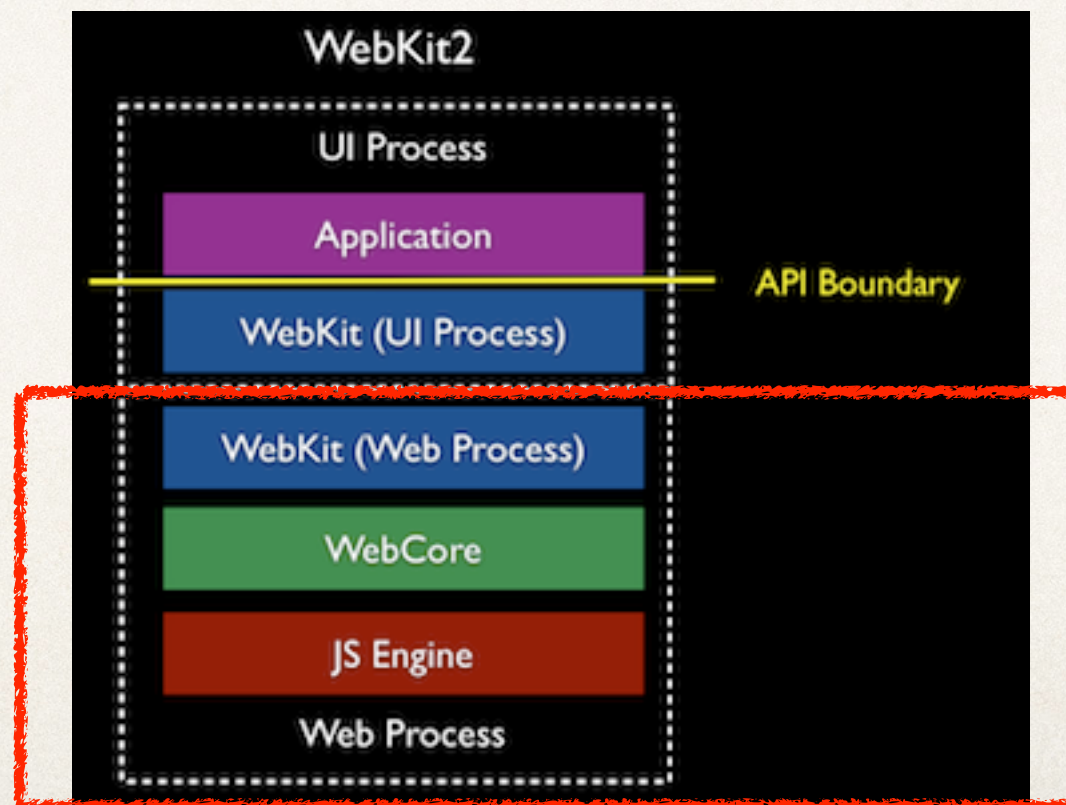
<https://googleprojectzero.blogspot.co.id/2014/11/pwn4fun-spring-2014-safari-part-ii.html>

<https://trac.webkit.org/wiki/WebKit2>

Overview of Pwning Safari



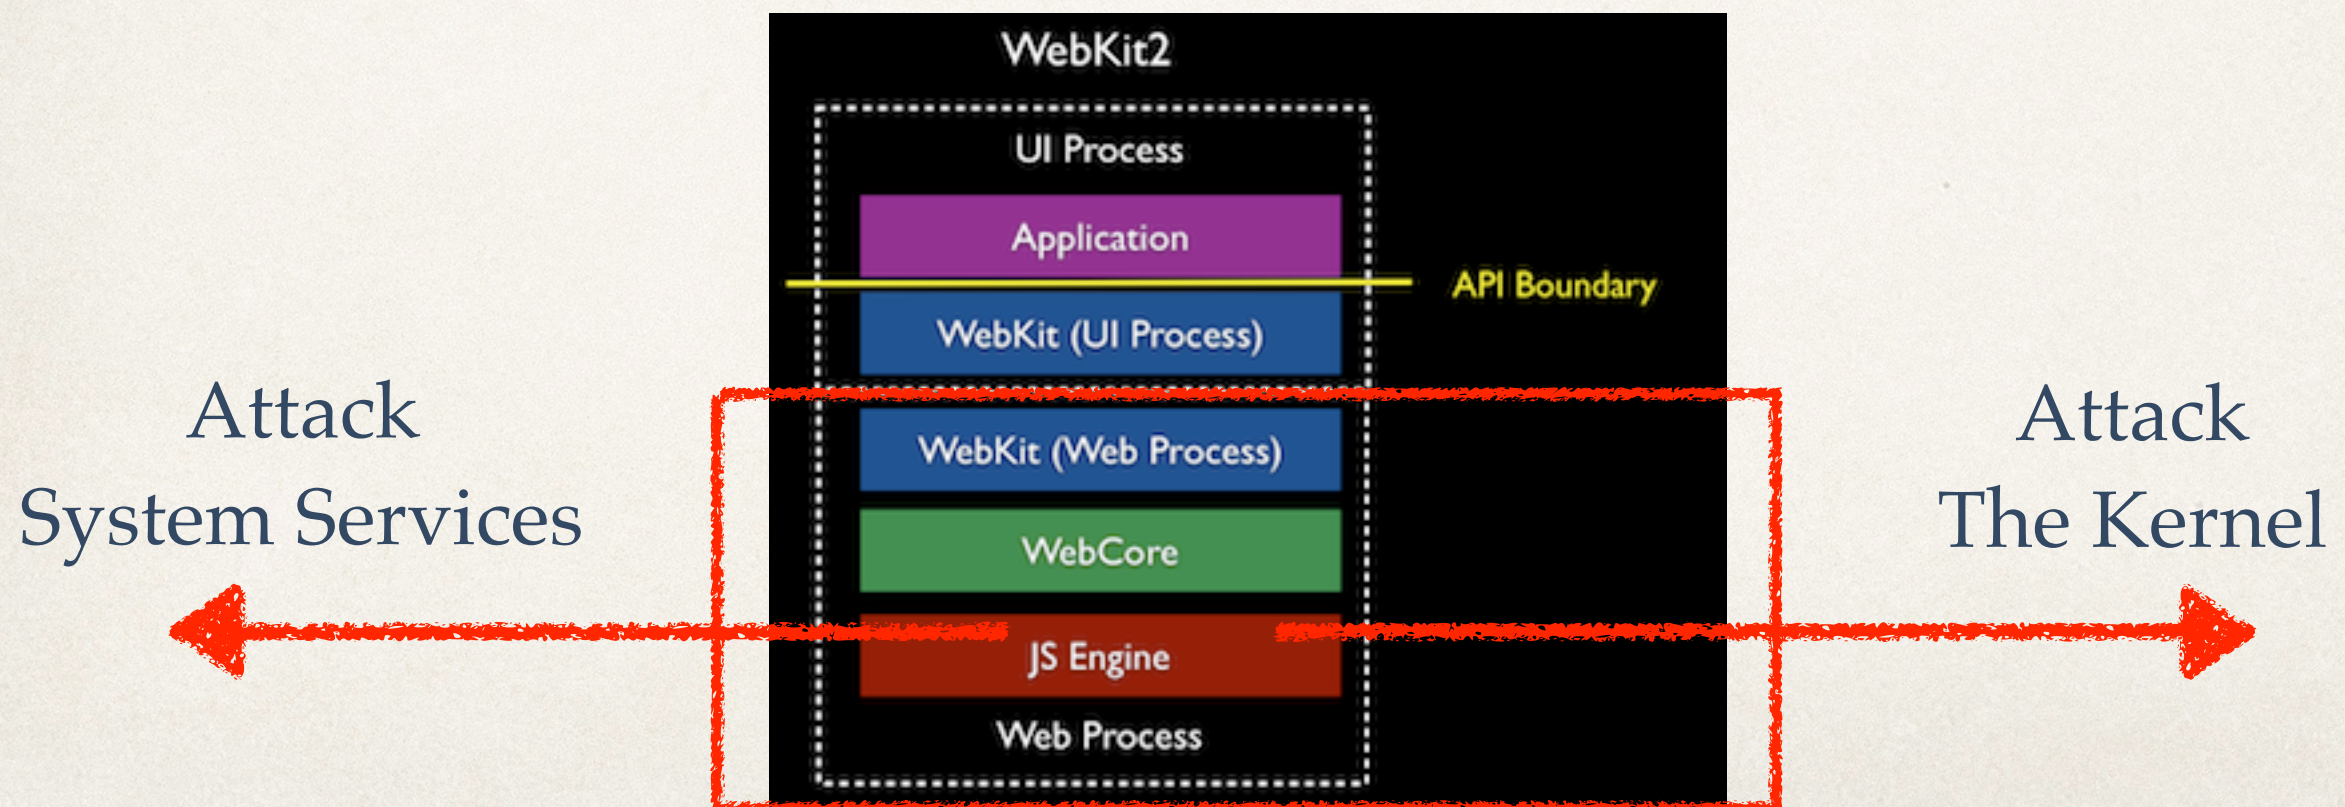
- ❖ WebProcess is responsible for loading, parsing, rendering, thus is the main target
- ❖ But WebProcess is confined by sandbox



Overview of Pwning Safari



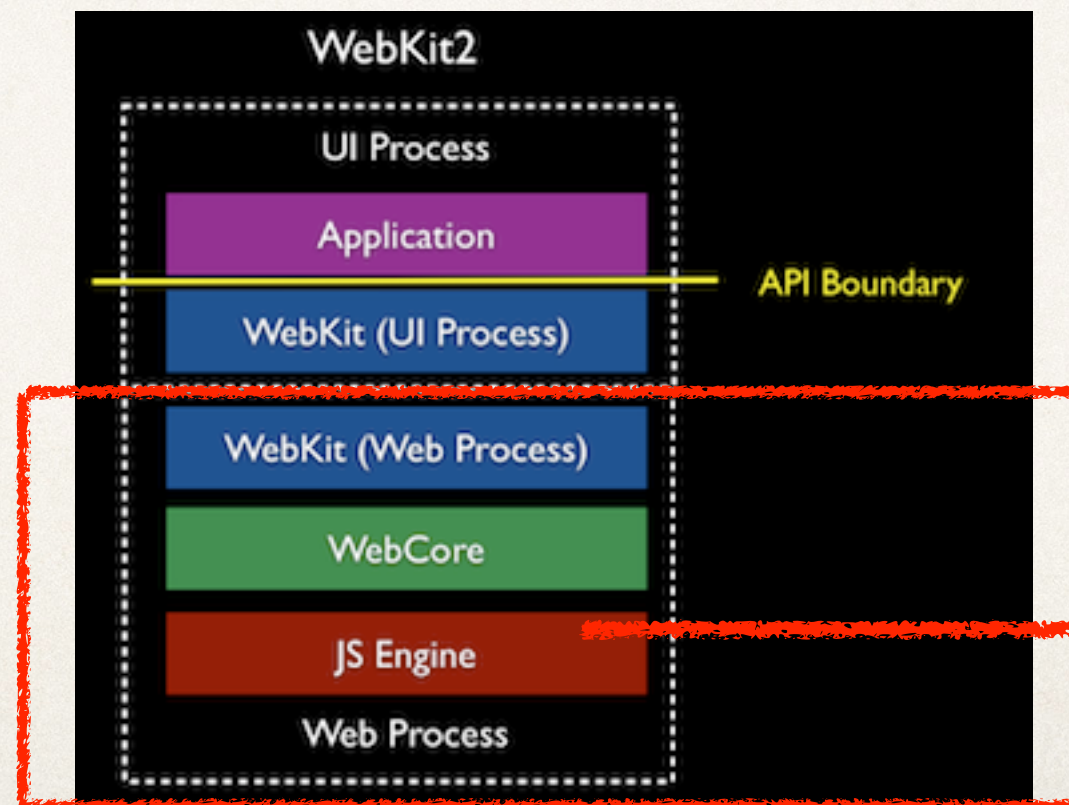
- ❖ After gaining RCE in WebProcess, we need to either exploit other system services via IPC (e.g., WindowServer, fontd), or exploit kernel vulnerabilities directly



Focus of Our Talk



- ❖ Directly exploit kernel vulnerabilities



Attack
The Kernel

Agenda



- ❖ Introduction
- ❖ **Kernel Attack Surface in Safari**
- ❖ Kernel Vulnerabilities and Exploitations
- ❖ Conclusion

WebProcess Sandbox Profiles



- ❖ /System/Library/Frameworks/WebKit.framework/Versions/A/Resources/com.apple.WebProcess.sb
- ❖ /System/Library/Sandbox/Profiles/system.sb

```
;; IOKit user clients
(allow iokit-open
  (iokit-user-client-class "AppleUpstreamUserClient")
  (iokit-user-client-class "IOHIDParamUserClient")
  (iokit-user-client-class "RootDomainUserClient")
  (iokit-user-client-class "IOAudioControlUserClient")
  (iokit-user-client-class "IOAudioEngineUserClient"))
```

```
(allow iokit-open
  (iokit-connection "IOAccelerator")
  (iokit-registry-entry-class "IOAccelerationUserClient")
  (iokit-registry-entry-class "IOSurfaceRootUserClient")
  (iokit-registry-entry-class "IOSurfaceSendRight"))
;; CoreVideo CVCGLDisplayLink
(allow iokit-open
  (iokit-registry-entry-class "IOFramebufferSharedUserClient"))
;; H.264 Acceleration
(allow iokit-open
  (iokit-registry-entry-class "AppleSNBFBUserClient"))
;; QuartzCore
(allow iokit-open
  (iokit-registry-entry-class "AGPMClient")
  (iokit-registry-entry-class "AppleGraphicsControlClient")
  (iokit-registry-entry-class "AppleGraphicsPolicyClient"))
```


Quick View of IOKit



- ❖ The IOKit is a collection of system frameworks, libraries, tools, and other resources for creating device drivers
- ❖ Extensive user mode API (via IOKit.framework)
 - ❖ Direct memory mapping to user mode
 - ❖ UserClient/ External method calls
 - ❖ Notification and messaging

IOKit Interface = Attack Surface



- ❖ IOKitLib implements non-kernel task access to common IOKit object types - IORegistryEntry, IOService, IOIterator, etc
 - ❖ IOServiceGetMatchingService
 - ❖ Look up a registered IOService object that matches a matching dictionary
 - ❖ IOServiceOpen
 - ❖ Create a connection to an IOService, return an IOUserClient port
 - ❖ IOConnectSetNotificationPort
 - ❖ Set a port to receive notifications from an IOUserClient object
 - ❖ IOConnectCallMethod
 - ❖ Pass/get data to an IOUserClient object

Agenda



- ❖ Introduction
- ❖ Kernel Attack Surface in Safari
- ❖ **Kernel Vulnerabilities and Exploitations**
- ❖ Conclusion

Kernel Vulnerabilities



- ❖ Uninitialized heap in IOAudioFamily leading to info leak
- ❖ Uninitialized stack in AMD graphics driver leading to arbitrary code execution
- ❖ Heap overflow in IOAcceleratorFamily2 caused by double-fetch in shared memory leading to arbitrary code execution

Case 1. IOAudioFamily Info Leak

- ❖ The vulnerability lies in IOAudioFamily (≤ 204.4)
- ❖ Source code:
<https://opensource.apple.com/source/IOAudioFamily/IOAudioFamily-204.4/>
- ❖ The vulnerability is an uninitialized heap issue

Vulnerability Analysis



- ❖ IOAudioControlUserClient allows userspace programs to register a notification port and send a notification message to them when certain audio events happen
- ❖ IOConnectSetNotificationPort in userspace will reach IOAudioControlUserClient::registerNotificationPort in the kernel

Vulnerability Analysis



```
IOReturn IOAudioControlUserClient::registerNotificationPort(mach_port_t port, UInt32 type, UInt32 refCon)
{
    ...
    if (notificationMessage == 0) {
        notificationMessage = (IOAudioNotificationMessage *)
            IOMallocAligned(sizeof(IOAudioNotificationMessage), sizeof (IOAudioNotificationMessage *));
        if (!notificationMessage) {
            return kIOReturnNoMemory;
        }
    }
    notificationMessage->messageHeader.msgh_bits = MACH_MSGH_BITS(MACH_MSG_TYPE_COPY_SEND, 0);
    notificationMessage->messageHeader.msgh_size = sizeof(IOAudioNotificationMessage);
    notificationMessage->messageHeader.msgh_remote_port = port;
    notificationMessage->messageHeader.msgh_local_port = MACH_PORT_NULL;
    notificationMessage->messageHeader.msgh_reserved = 0;
    notificationMessage->messageHeader.msgh_id = 0;
    notificationMessage->ref = refCon;
    ...
}

typedef struct _IOAudioNotificationMessage
{
    mach_msg_header_t messageHeader;
    UInt32 type;
    UInt32 ref;
    void * sender;
} IOAudioNotificationMessage;
```

notificationMessage is not zeroed out and the *type* and the *sender* field is not initialized

Vulnerability Analysis



- ❖ IOAudioControlUserClient allocates a notification message struct and sends it to userspace programs via
IOAudioControlUserClient::sendChangeNotification

Leak to Userland



```
void IOAudioControlUserClient::sendChangeNotification(UInt32 notificationType)
{
    if (notificationMessage) {
        kern_return_t kr;

        notificationMessage->type = notificationType;
        kr = mach_msg_send_from_kernel(&notificationMessage->messageHeader,
notificationMessage->messageHeader.msgh_size);
        if ((kr != MACH_MSG_SUCCESS) && (kr != MACH_SEND_TIMED_OUT)) {
            IOLog("IOAudioControlUserClient: sendRangeChangeNotification() failed - msg_send
returned: %d\n", kr);
        }
    }
}
```

- ❖ The notificationMessage is allocated in the kalloc.72 zone
- ❖ The *sender* field offset is 0x38
- ❖ 8-byte data at the offset of 0x38 from a freed object in kalloc.72 leaked

Exploiting The Vulnerability



- ❖ Create many IOAudioControlUserClient objects through IOServiceOpen
- ❖ Register notification ports on such IOAudioControlUserClient objects via IOConnectSetNotificationPort
 - ❖ each IOAudioControlUserClient object will allocate a notificationMessage struct with uninitialized 8 bytes data
- ❖ Trigger a notification event, and receive the notification messages

However...



- ❖ Challenge 1: How to trigger the notification event in webcontent process
- ❖ Challenge 2: How to leak critical information such as KALSR slide

Sandbox Restriction



- ❖ Setting sound volume level can trigger the notification
 - ❖ i.e., set “IOAudioControlValue” via
IORegistryEntrySetCFProperties
- ❖ However, the webcontent sandbox profile does not allow to set this property

coreaudiod



- ❖ Who is allowed to set sound volume level?
- ❖ By greping IOAudioControlValue, we found `/usr/sbin/coreaudiod` can set sound volume level
- ❖ `coreaudiod` is also responsible for Mach Service `com.apple.audio.coreaudiod`

coreaudiod



- ❖ Webprocess can also talk with com.apple.audio.coreaudiod service
- ❖ As a result, we can instruct coreaudiod to trigger the notification event

```
;; Various services required by AppKit and other frameworks
(allow mach-lookup
  (global-name "com.apple.DiskArbitration.diskarbitrationd")
  (global-name "com.apple.FileCoordination")
  (global-name "com.apple.FontObjectsServer")

  (global-name "com.apple.PowerManagement.control")
  (global-name "com.apple.SystemConfiguration.configd")
  (global-name "com.apple.SystemConfiguration.PPPController")
  (global-name "com.apple.audio.SystemSoundServer-OSX")
  (global-name "com.apple.audio.VDCAssistant")
  (global-name "com.apple.audio.audiohald")
  (global-name "com.apple.audio.coreaudiod")
```


How to leak KALSR slide



- ❖ How to make the 8-byte data meaningful?
- ❖ We can read 8-byte data at the offset of 0x38 from a freed object in kalloc.72

Leak From OSSerialize Object



```
class OSSerialize : public OSObject
{
    ...
private:
    char          * data;
    unsigned int   length;
    unsigned int   capacity;
    unsigned int   capacityIncrement;
    OSArray * tags;
    bool   binary;
    bool   endCollection;
    Editor editor;
    void * editRef;
    ...
}
```

```
typedef const OSMetaClassBase * (*Editor)(
    void* reference,
    OSSerialize* s,
    OSCollection* container,
    const OSSymbol* name,
    const OSMetaClassBase* value);
```

- ❖ The OSSerialize objects locates in the same zone as notificationMessage object
- ❖ The member at offset 0x38 is a function pointer
- ❖ Use separate threads to spray OSSerialize objects and leak several times to ensure stability

Case 2. AMDRadeon Code Execution



- ❖ Accelerator is one of the devices that we can directly access in the webcontent sandbox
- ❖ The vulnerability lies in AMDRadeonX`x000.kext` (`x may vary on different platforms`)
- ❖ The vulnerability is an uninitialized stack variable issue and the exploit is quite straight-forward

Vulnerability Overview



- ❖ AMD Accelerator can create different user clients according to the type parameter

```
kern_return_t IOServiceOpen(io_service_t service, task_port_t owningTask, uint32_t type, io_connect_t *connect);
```

- ❖ AMDSIGLContext (type 1)
- ❖ AMDAccelSharedUserClient (type 6)

Vulnerability Overview

- ❖ AMDSIGLContext is not started until AMDAccelSharedUserClient is connected

```
kern_return_t IOConnectAddClient(io_connect_t connect, io_connect_t client);
```

- ❖ AMDRadeonX4000_AMDSIGLContext's externalMethods interface dispatches functions according to method selector

```
kern_return_t IOConnectCallMethod(mach_port_t connection, uint32_t selector, ...);
```


Vulnerability Overview

- ❖ selector 0x201 will reach function
AMD RadeonX4000_AMDSIGLContext::surfaceCopy
- ❖ surfaceCopy is supposed to lookup a resource object according to input index and proceeds to use the resource object

```
IOAccelShared2::lookupResource(v5, a2[2], &v51);  
IOAccelShared2::lookupResource(v5, a2[1], (void **)&v50);  
v6 = -536870206;  
if ( !v51 || !v50 )  
    goto LABEL_42;  
v12 = (*(__int64 (**)(void))(*(_QWORD *)v51 + 368LL))();  
v13 = (*(__int64 (**)(void))(*(_QWORD *)v50 + 368LL))();
```


Vulnerability Analysis

- ❖ v51 and v50 are two local variables on the stack
- ❖ lookupResource may fail for invalid index
- ❖ surfaceCopy neither initializes the local variables nor checks the return value from lookupResource
- ❖ Supplying invalid index will trigger a panic

```
IOAccelShared2::lookupResource(v5, a2[2], &v51);  
IOAccelShared2::lookupResource(v5, a2[1], (void **)&v50);  
v6 = -536870206;  
if ( !v51 || !v50 )  
    goto LABEL_42;  
v12 = (*(__int64 (**)(void))(*(_QWORD *)v51 + 368LL))();  
v13 = (*(__int64 (**)(void))(*(_QWORD *)v50 + 368LL))();
```


Control The Stack



- ❖ The uninitialized value on kernel stack is “random” so we need to control the stack
- ❖ Luckily, a function (selector 7333) in the AGPM userclient comes to rescue
- ❖ We are able to copy at most 4096 bytes of controlled, non-zero data onto kernel stack

```
case 7333:
    kprintf("kAGPMSetPlimit plimit = %llu type = %s\n", *a2->scalarInput, a2->structureInput);
    v16 = (char *)&v22 - ((a2->structureInputSize + 1 + 15LL) & 0xFFFFFFFFFFFFFFFF0LL);
    strncpy(v16, (const char *)a2->structureInput, a2->structureInputSize);
```


Control The Stack



- ❖ What value should we “initialize” for the uninitialized stack variable
- ❖ A pointer pointing to an object we fake on the heap
- ❖ A pointer pointing to a real object on the heap
 - ❖ It is hard to find such a candidate
 - ❖ You may have a try ;-)

Control The Stack



- ❖ How can we know where our fake objects locates?
- ❖ Option - 1
 - ❖ Reuse that leak to get heap address info
 - ❖ Hunt for appropriate object - **HARD**
 - ❖ Too many free-refill operations - **UNSTABLE**

Control The Stack



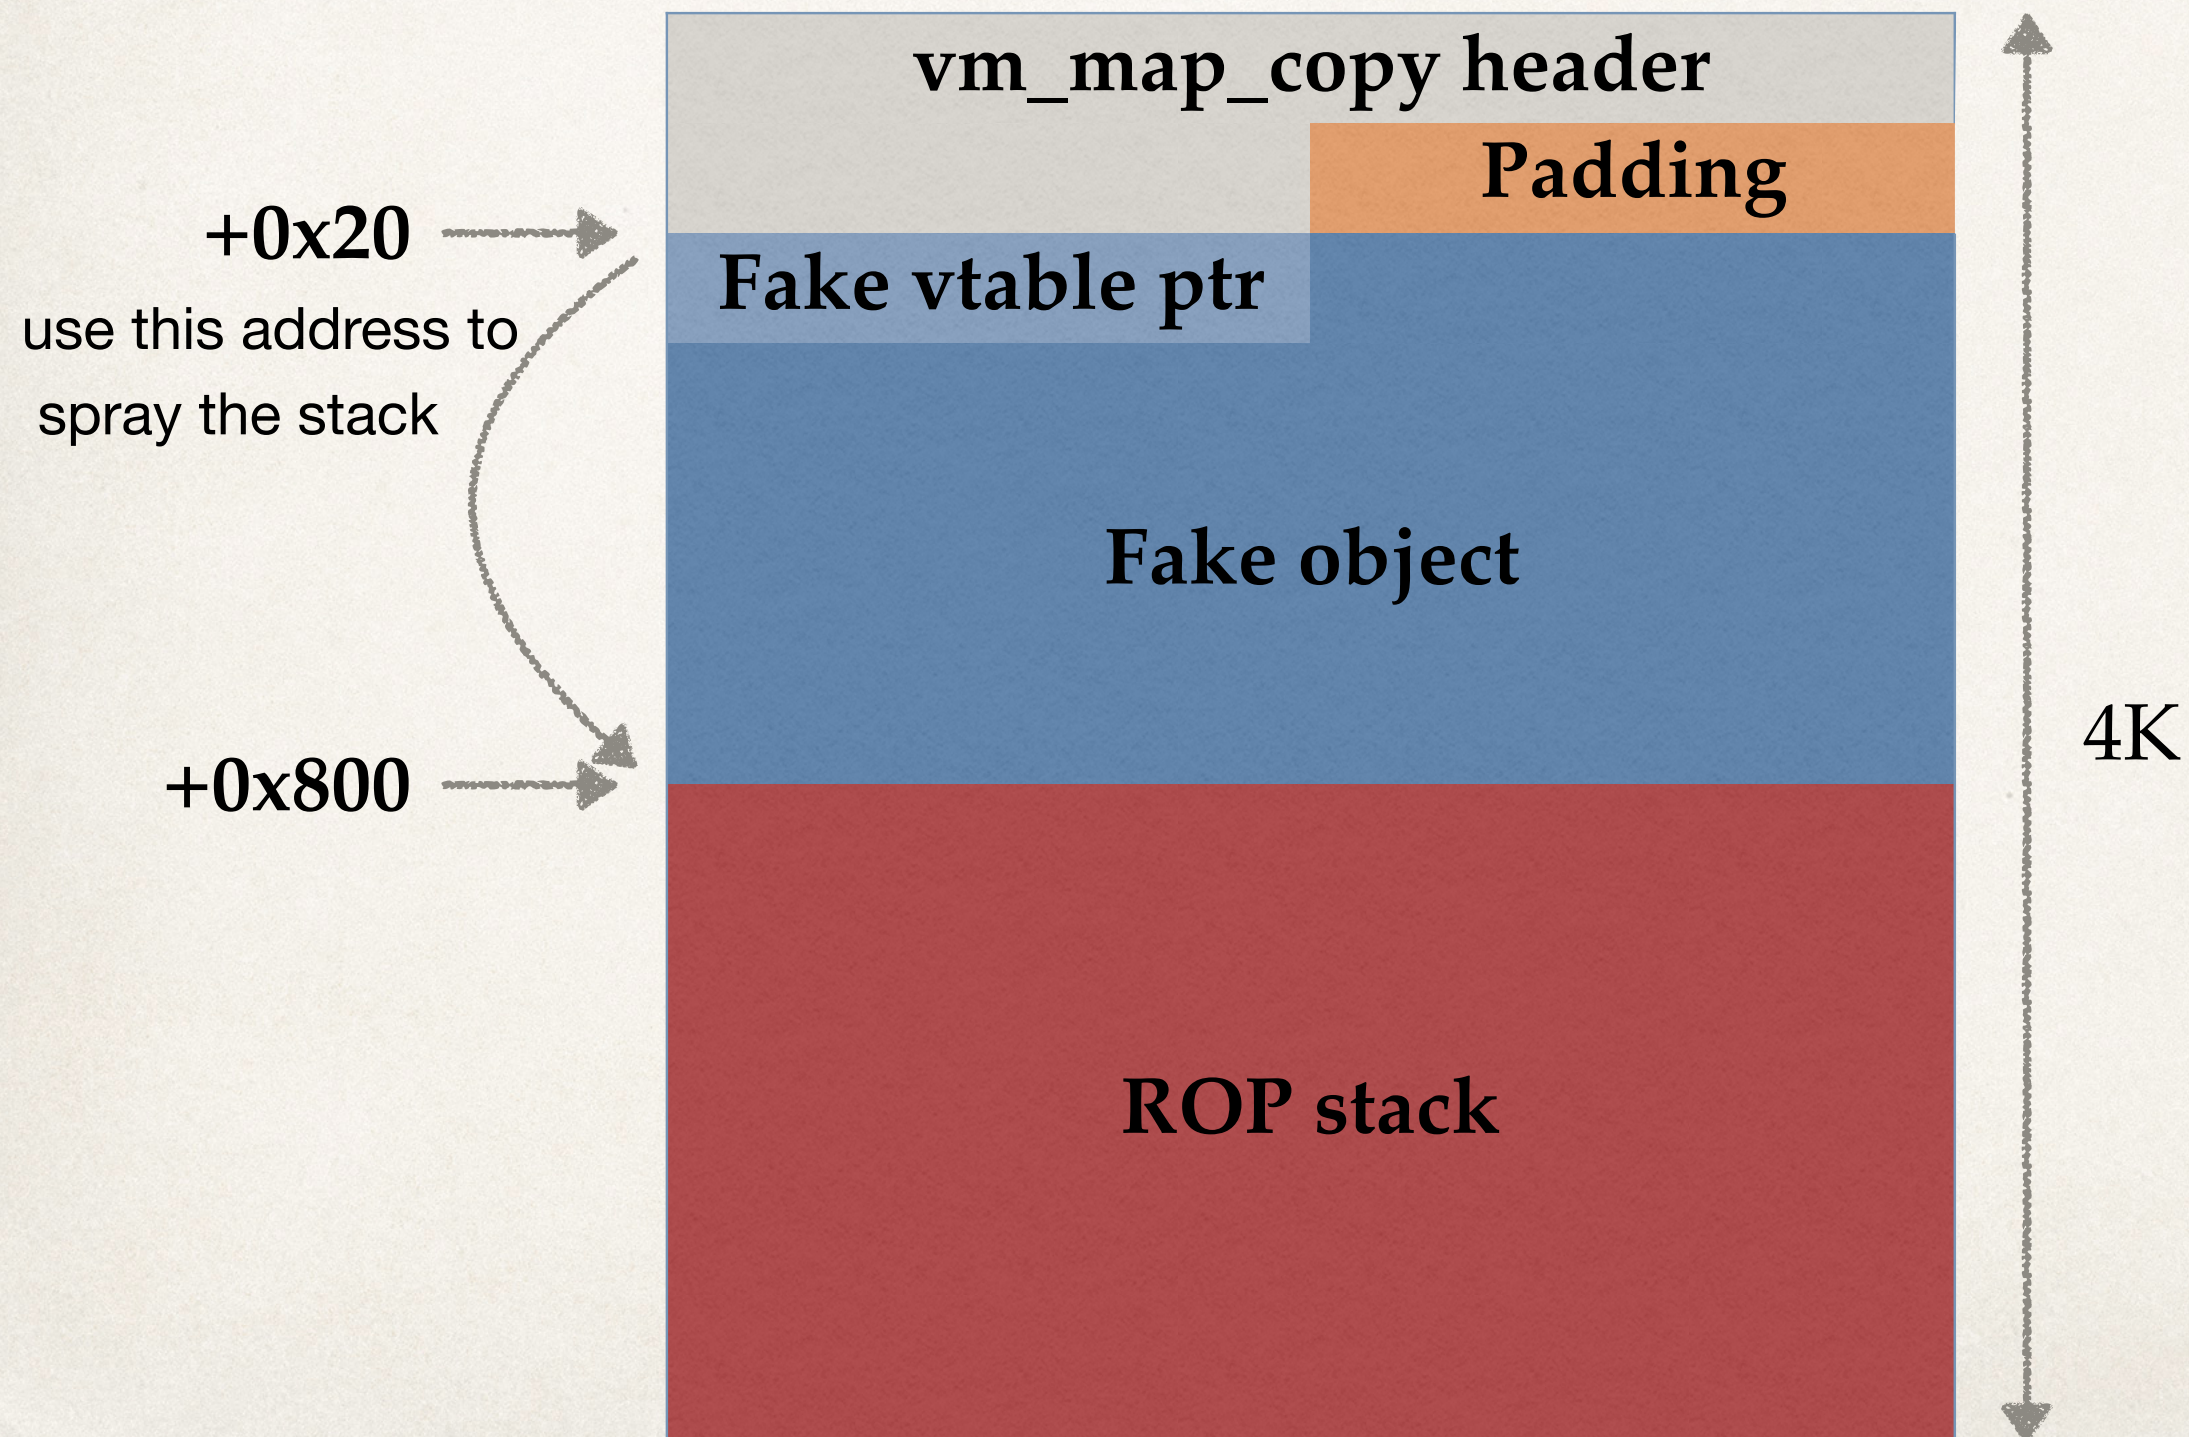
- ❖ How can we know where our fake objects locates?
- ❖ Option - 2
 - ❖ Spray several GBs of data in the kernel
 - ❖ Heap randomization weak in the kernel
 - ❖ User-controlled data locates at a fixed address

Spray Fake Objects



- ❖ Spray with `vm_map_copy`
 - ❖ Fast and small side effects on heap
 - ❖ Arbitrary controlled size (up to 4K) and contents
 - ❖ Uncontrolled beginning 0x18 bytes - NOT matter
 - ❖ Fake object contains a fake vtable and a fake stack

Spray Fake Objects



ROP Chain



- ❖ Save registers
- ❖ Pivot Stack
- ❖ Disable SMEP & SMAP
- ❖ Return to userland SHELLCODE
- ❖ Re-enable SMEP & SMAP
- ❖ Call `_thread_exception_return()` to exit

Disable SMEP&SMAP



- ❖ Bits stored in CR4 register indicate the states of CPU
- ❖ Clear 21st to disable SMEP and 22nd for SMAP
- ❖ gadgets:
 - ❖ read CR4: `mov rax, cr4 ,..., ret`
 - ❖ Write CR4: `mov cr4, rax ,..., ret`

Special Notes



- ❖ The kernel stack used when entering kernel mode varies from time to time
 - ❖ call `sprayStack()` several times to cover more stacks
- ❖ Do not forget to unlock the locks locked in `AMDSIGLContext::surfaceCopy()`

Case 3. IOAcceleratorFamily2 Heap Overflow

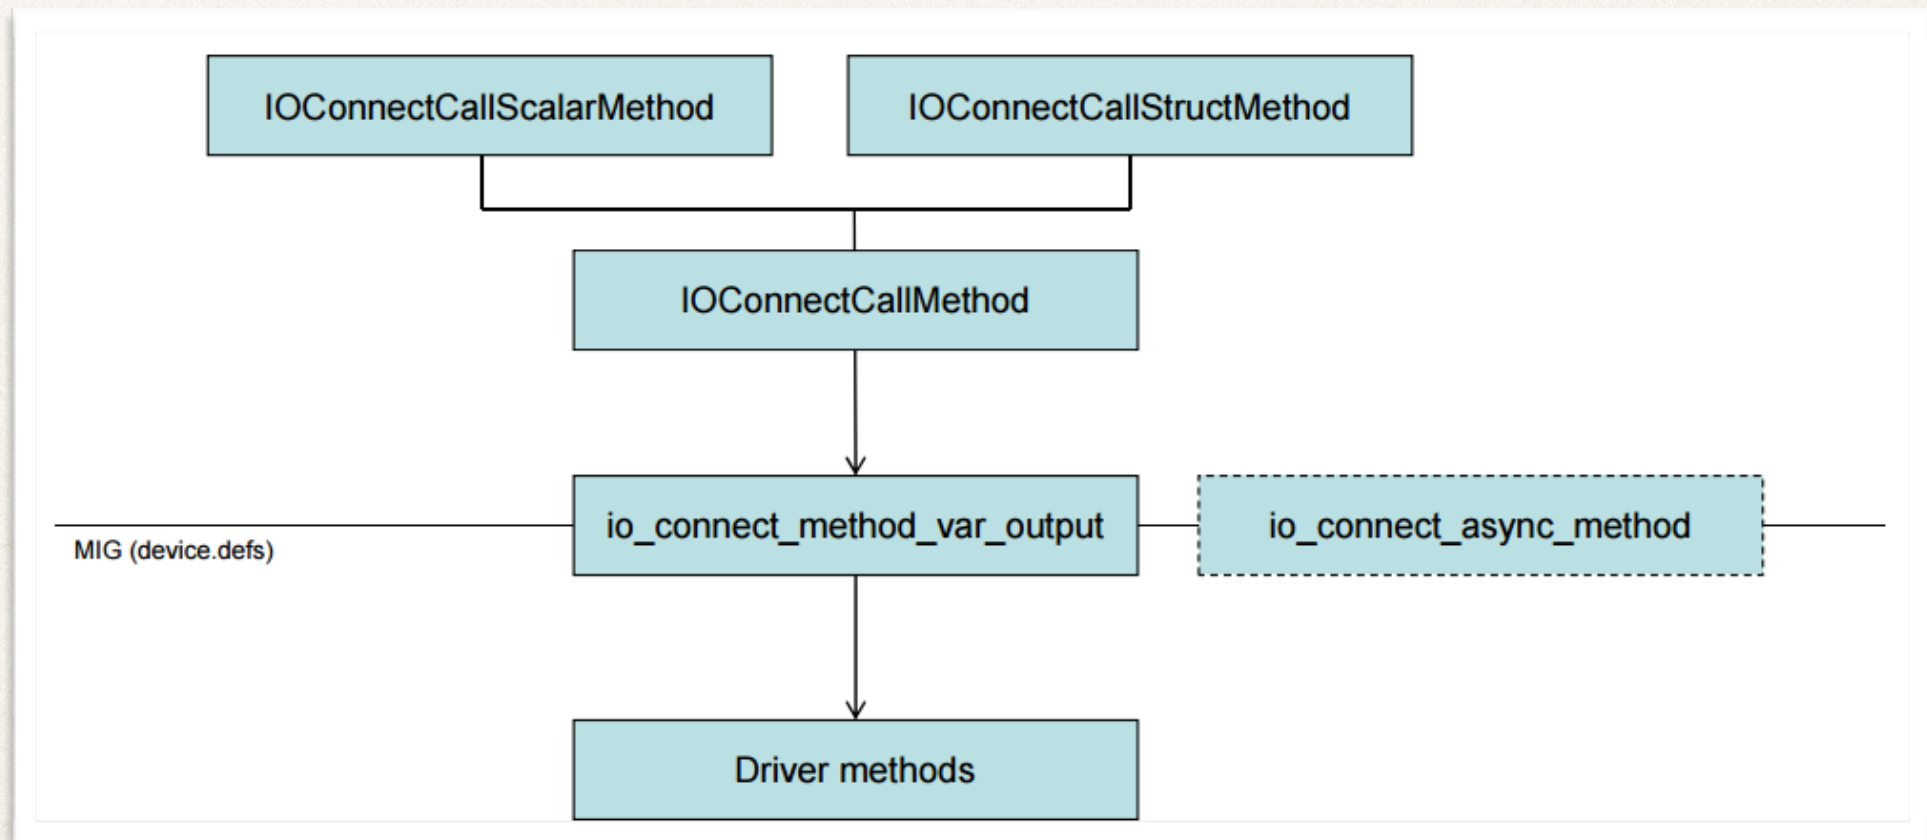


- ❖ One of my favorite issue.
- ❖ It is caused by a fundamental issue in IOKit data sharing mechanism that results in many kernel bugs
- ❖ Reported (~~burned~~) by KeenLab to Apple

Implicit Shared Memory



- ❖ IOConnectCallMethod is used to send / receive data to / from an IOUserClient object



Implicit Shared Memory

```
/* Routine io_user_client_method */
kern_return_t is_io_connect_method
(
    io_connect_t connection,
    uint32_t selector,
    io_scalar_inband64_t scalar_input,
    mach_msg_type_number_t scalar_inputCnt,
    io_struct_inband_t inband_input,
    mach_msg_type_number_t inband_inputCnt,
    mach_vm_address_t ool_input,
    mach_vm_size_t ool_input_size,
    io_struct_inband_t inband_output,
    mach_msg_type_number_t *inband_outputCnt,
    io_scalar_inband64_t scalar_output,
    mach_msg_type_number_t *scalar_outputCnt,
    mach_vm_address_t ool_output,
    mach_vm_size_t *ool_output_size
)
```

up to 16 uint64_t

up to 4096 chars

userland address

- ❖ `io_connect_method` can pass shared memory to the kernel objects through `ool_input`

IOAccelDisplayPipeUserClient2



- ❖ IOGraphicsAccelerator2::newUserClient can return a userclient object of IOAccelDisplayPipeUserClient2 (type=4)
- ❖ IOAccelDisplayPipeUserClient2::externalmethod supports many methods
- ❖ recall that io_connect_method will reach IOAccelDisplayPipeUserClient2::externalmethod

Vulnerability Analysis



❖ IOAccelDisplayPipePostCSCGammaVID::init

```
char __fastcall IOAccelDisplayPipePostCSCGammaVID::init(__int64 a1, __int64 ool_input)
{
    // [COLLAPSED LOCAL DECLARATIONS. PRESS KEYPAD CTRL-"+" TO EXPAND]

    v2 = (void *)IOMalloc(*(_QWORD *)(ool_input + 40));
    *(_QWORD *)(a1 + 56) = v2;
    if ( !v2 )
        return 0;
    *(_DWORD *)(a1 + 12) = *(_DWORD *)(ool_input + 4);
    *(_QWORD *)(a1 + 16) = *(_QWORD *)(ool_input + 8);
    *(_QWORD *)(a1 + 24) = *(_QWORD *)(ool_input + 16);
    *(_QWORD *)(a1 + 32) = *(_QWORD *)(ool_input + 24);
    v3 = *(_QWORD *)(ool_input + 32);
    *(_DWORD *)(a1 + 40) = v3;
    *(_DWORD *)(a1 + 44) = HIWORD(v3);
    size = *(_QWORD *)(ool_input + 40);
    *(_QWORD *)(a1 + 48) = size;
    memcpy(v2, (const void *)(ool_input + 48), size);
    return 1;
}
```

alloc heap with controllable size

refetch the size from shared memory

perfect heap overflow in memcpy

Vulnerability Exploitation



- ❖ Racing the double-fetch
- ❖ Thread1: send ool input to
IOAccelDisplayPipeUserClient2
- ❖ Thread2: modify the size field

Vulnerability Exploitation



- ❖ Heap Overflow occurs in memcpy
- ❖ What we can control
 - ❖ The dest buffer's size
 - ❖ The src buffer's size (through race condition)
 - ❖ The src buffer's content (in shared memory)

Kernel Vulnerability Summary



- ❖ Case 1 and 2 were fixed in 10.12.3, January 23, 2017
 - ❖ CVE-2017-2358
 - ❖ CVE-2017-2357
 - ❖ We used the two vulns in PwnFest 2016
- ❖ Case 3
 - ❖ No public CVE
 - ❖ Apple changed ool shared memory as COW

Agenda



- ❖ Introduction
- ❖ Kernel Attack Surface in Safari
- ❖ Kernel Vulnerabilities and Exploitations
- ❖ **Conclusion**

Conclusion



- ❖ Kernel exploitation is fun
- ❖ Exploiting macOS kernel inside Safari sandbox is hard, but is feasible, and will continue to be feasible

Thank You For Your Attention
Q&A

