

A Story of Exploiting macOS Sierra Kernel

Zhenquan Xu, Tielei Wang, Hao Xu of TEAM Pangu



Agenda

- ✿ New Mitigations in macOS Sierra
- ✿ Kernel Info Leak
- ✿ Kernel Code Execution
- ✿ Conclusion

macOS Sierra

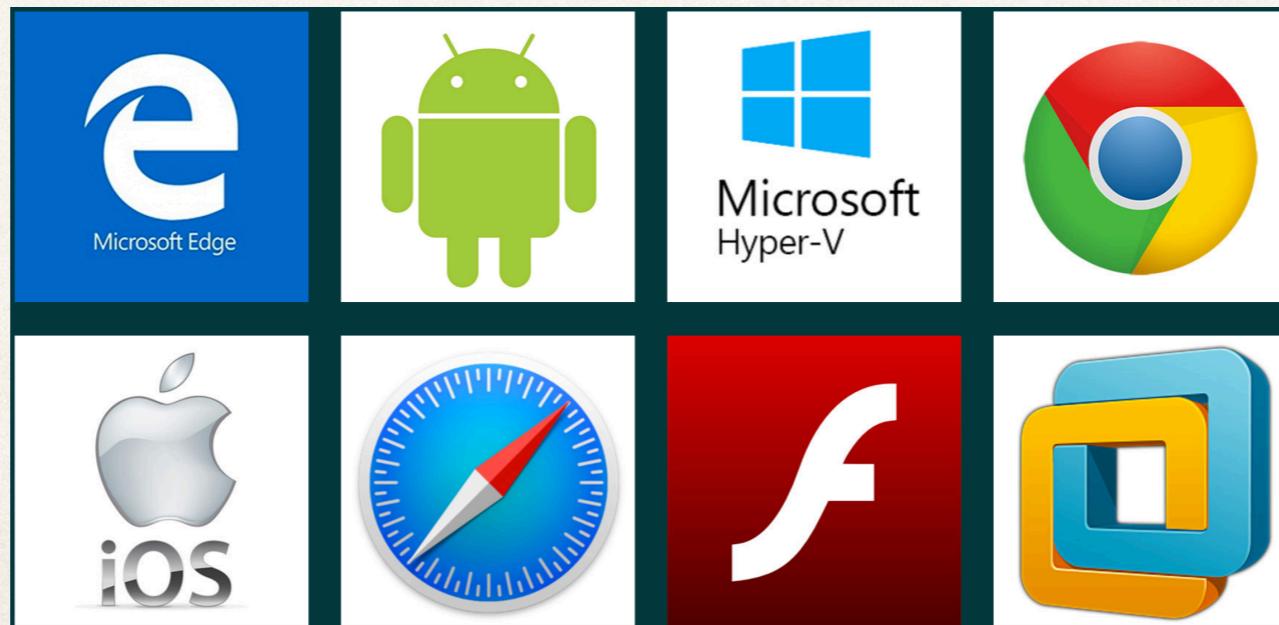
- ✿ macOS Sierra (version 10.12) is the thirteenth major release of macOS
- ✿ released to end users on September 20, 2016

PwnFest 2016

- ✿ Nov 10th ~ 11th, 2016, Seoul, Korea

WHAT IS PWNFEST

PwnFest is a bug pwning 'festival' for better security organized by POC with the help of sponsors, vendors, and judges in 2016. You can enjoy PwnFest every year.



PwnFest 2016

- ✿ Team up with Lokihardt for pwning Safari on macOS Sierra
 - ✿ Remote code execution in Safari by Lokihardt
 - ✿ Kernel code execution in Safari by us

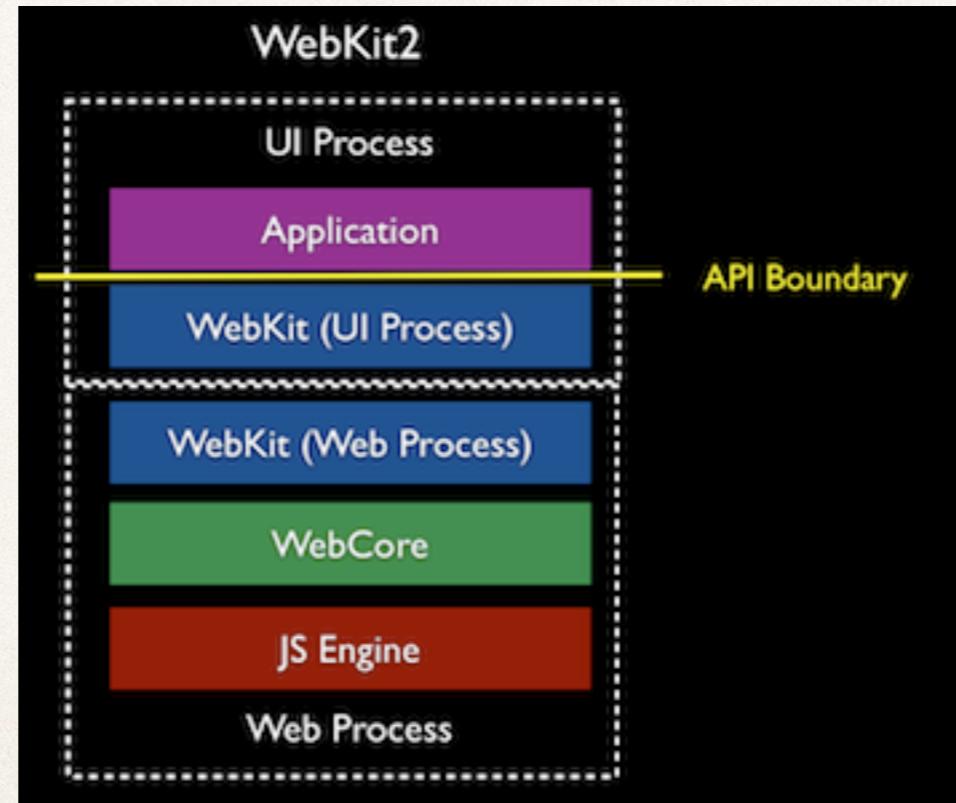


Fixed in macOS Sierra 10.12.3

- ✿ Released January 23, 2017
- ✿ CVE-2017-2358
- ✿ CVE-2017-2357

Overview of Pwn Process

- ✿ Safari process model is based on WebKit2, splitting into different families

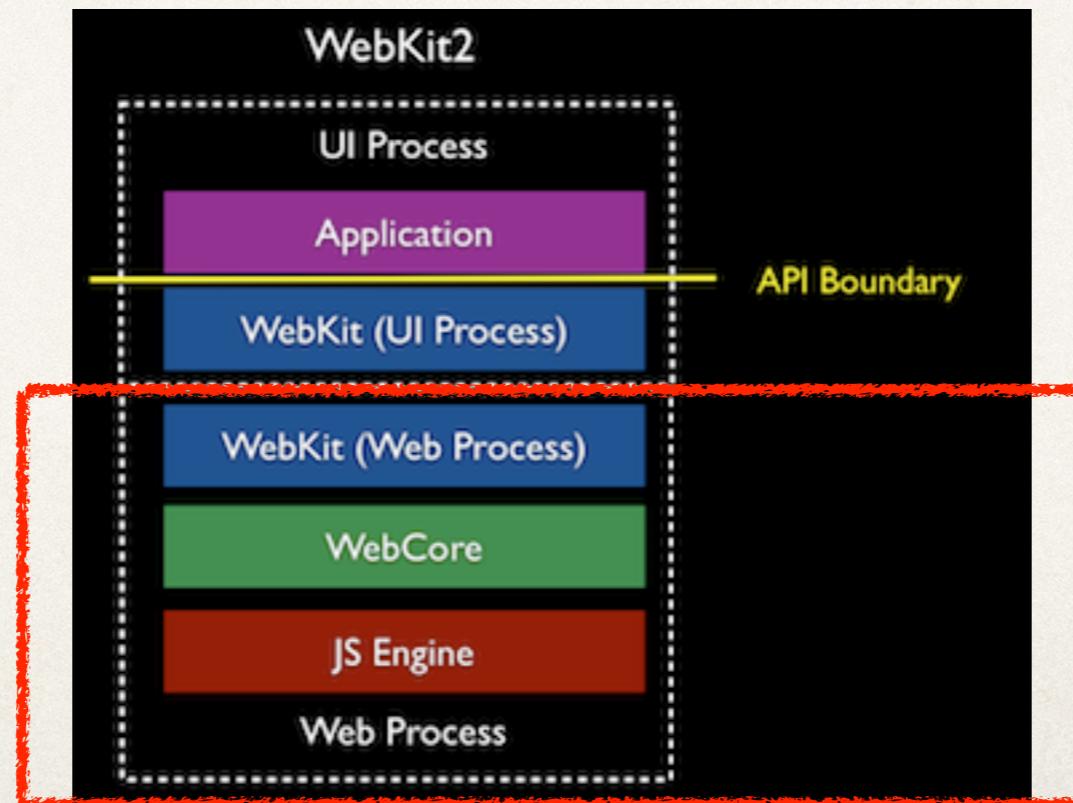


<https://googleprojectzero.blogspot.co.id/2014/11/pwn4fun-spring-2014-safari-part-ii.html>

<https://trac.webkit.org/wiki/WebKit2>

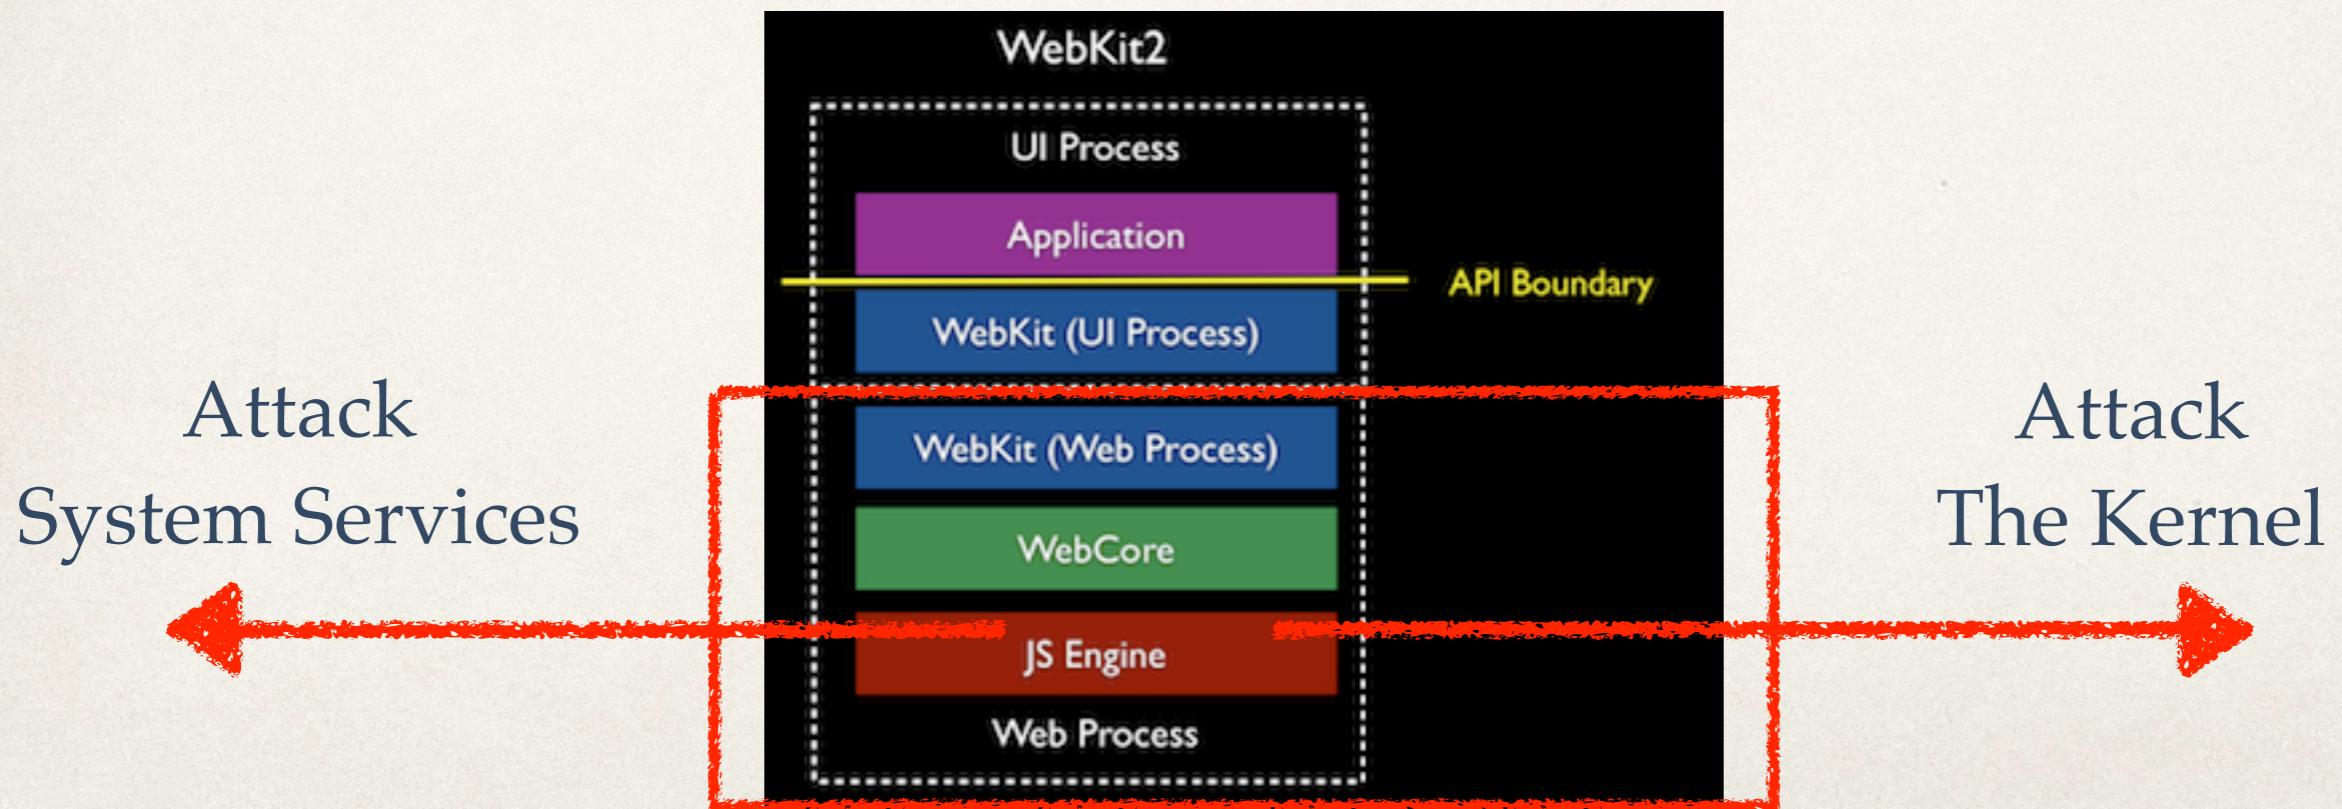
Overview of Pwn Process

- ❖ WebProcess is responsible for loading, parsing, rendering, thus is the main target
- ❖ But WebProcess is confined by sandbox



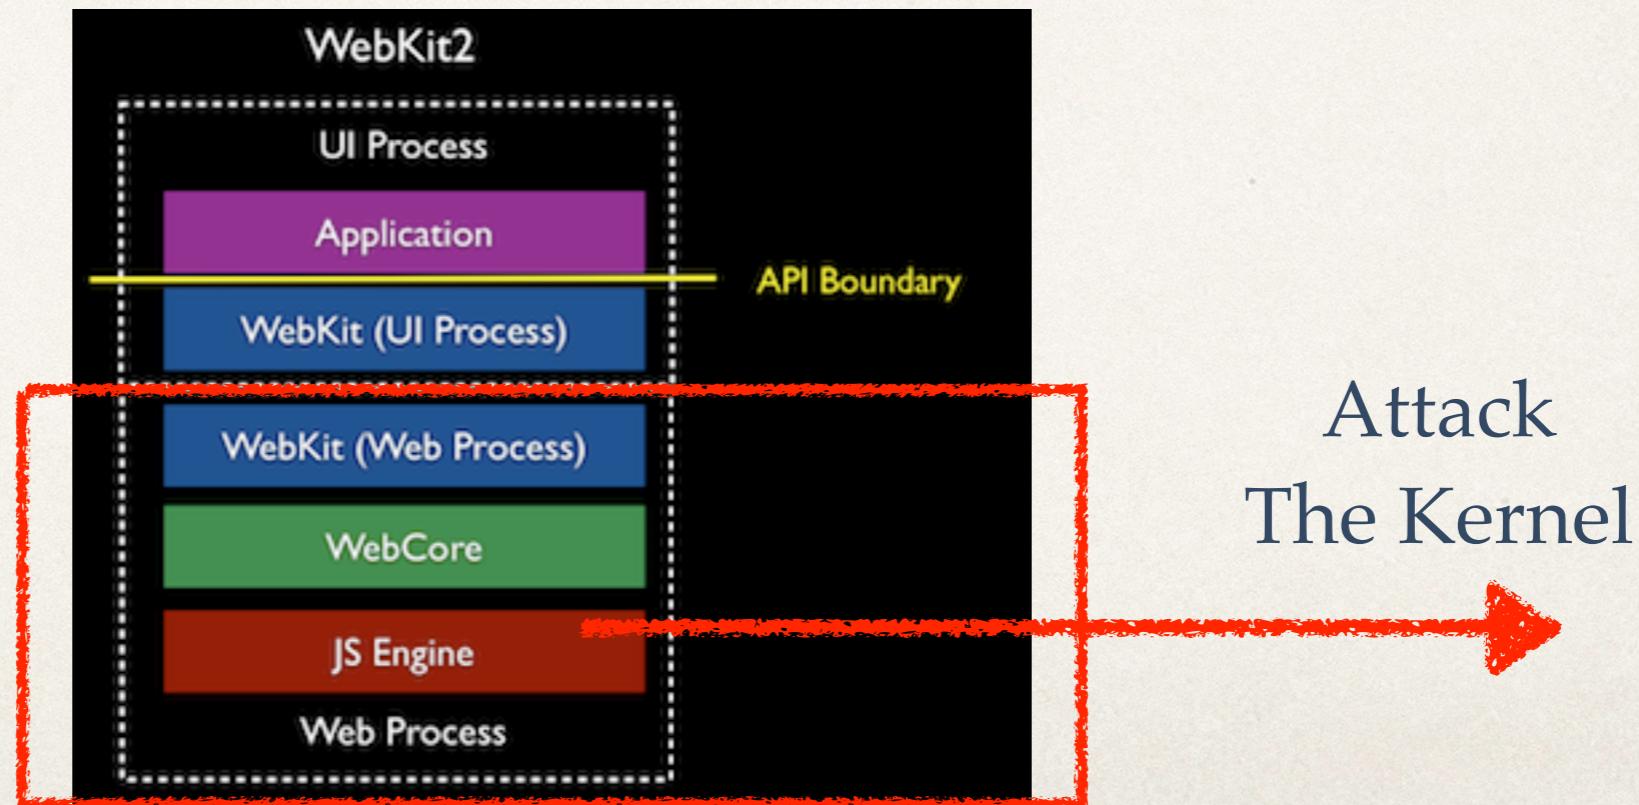
Overview of Pwn Process

- ✿ After gaining RCE in WebProcess, we need to either exploit other system services via IPC (e.g., WindowServer, fontd), or exploit kernel vulnerabilities directly



Overview of Pwn Process

- ❖ We decide to directly exploit kernel vulnerabilities



Review Kernel Attack Surfaces

- ✿ /System/Library/Frameworks/WebKit.framework/Versions/A/Resources/com.apple.WebProcess.sb
- ✿ /System/Library/Sandbox/Profiles/system.sb

```
; IOKit user clients
allow iokit-open
(iokit-user-client-class "AppleUpstreamUserClient")
(iokit-user-client-class "IOHIDParamUserClient")
(iokit-user-client-class "RootDomainUserClient")
(iokit-user-client-class "IOAudioControlUserClient")
(iokit-user-client-class "IOAudioEngineUserClient")
```

```
(allow iokit-open
    (iokit-connection "IOAccelerator")
    (iokit-registry-entry-class "IOAccelerationUserClient")
    (iokit-registry-entry-class "IOSurfaceRootUserClient")
    (iokit-registry-entry-class "IOSurfaceSendRight"))
;; CoreVideo CVCGDisplayLink
allow iokit-open
    (iokit-registry-entry-class "IOFramebufferSharedUserClient")
;; H.264 Acceleration
allow iokit-open
    (iokit-registry-entry-class "AppleSNFBUserClient")
;; QuartzCore
allow iokit-open
    (iokit-registry-entry-class "AGPMClient")
    (iokit-registry-entry-class "AppleGraphicsControlClient")
    (iokit-registry-entry-class "AppleGraphicsPolicyClient"))
```

Agenda

- ❖ New Mitigations in macOS Sierra
- ❖ Kernel Info Leak
- ❖ Kernel Code Execution
- ❖ Conclusion

New Mitigations

- ❖ No more io_service_open_extended heap spray

```
do
{
if (properties)
{
    OSObject *      obj;
    vm_offset_t     data;
    vm_map_offset_t map_data;

    if( propertiesCnt > sizeof(io_struct_inband_t))
        return( kIOReturnMessageTooLarge);

    err = vm_map_copyout( kernel_map, &map_data, (vm_map_copy_t) properties );
    res = err;
    data = CAST_DOWN(vm_offset_t, map_data);
    if (KERN_SUCCESS == err)
    {
        // must return success after vm_map_copyout() succeeds
        obj = OSUnserializeXML( (const char *) data, propertiesCnt );
        vm_deallocate( kernel_map, data, propertiesCnt );
        propertiesDict = OSDynamicCast(OSDictionary, obj);
        if (!propertiesDict)
        {
            res = kIOReturnBadArgument;
            if (obj)
                obj->release();
        }
    }
    if (kIOReturnSuccess != res)
        break;
}
```

```
if (properties) return (kIOReturnUnsupported);
#ifndef __OSX_Kernel__
{
    OSObject *      obj;
    vm_offset_t     data;
    vm_map_offset_t map_data;

    if( propertiesCnt > sizeof(io_struct_inband_t))
        return( kIOReturnMessageTooLarge);

    err = vm_map_copyout( kernel_map, &map_data, (vm_map_copy_t) properties );
    res = err;
    data = CAST_DOWN(vm_offset_t, map_data);
    if (KERN_SUCCESS == err)
    {
        // must return success after vm_map_copyout() succeeds
        obj = OSUnserializeXML( (const char *) data, propertiesCnt );
        vm_deallocate( kernel_map, data, propertiesCnt );
        propertiesDict = OSDynamicCast(OSDictionary, obj);
        if (!propertiesDict)
        {
            res = kIOReturnBadArgument;
            if (obj)
                obj->release();
        }
    }
    if (kIOReturnSuccess != res)
        break;
}
#endif
```

New Mitigations

- ✿ Limited count of IOUserClient objects
- ✿ After reaching the threshold, IOServiceOpen() returns an error

```
2017-04-07 14:11:14.446 test[51571:2757644] userclient count: 885, kr = 0x00000000 userclient port: 0x38b03
2017-04-07 14:11:14.446 test[51571:2757644] userclient count: 886, kr = 0x00000000 userclient port: 0x38c03
2017-04-07 14:11:14.446 test[51571:2757644] userclient count: 887, kr = 0x00000000 userclient port: 0x38d03
2017-04-07 14:11:14.447 test[51571:2757644] userclient count: 888, kr = 0x00000000 userclient port: 0x38e03
2017-04-07 14:11:14.447 test[51571:2757644] userclient count: 889, kr = 0x00000000 userclient port: 0x38f03
2017-04-07 14:11:29.447 test[51571:2757644] userclient count: 890, kr = 0xe00002c7 userclient port: 0x0
2017-04-07 14:11:44.447 test[51571:2757644] userclient count: 891, kr = 0xe00002c7 userclient port: 0x0
2017-04-07 14:11:59.447 test[51571:2757644] userclient count: 892, kr = 0xe00002c7 userclient port: 0x0
2017-04-07 14:12:14.448 test[51571:2757644] userclient count: 893, kr = 0xe00002c7 userclient port: 0x0
```

New Mitigations

- ✿ Changes in heap memory management
- ✿ All zones use metadata now
- ✿ Four queues of metadata:
 - ✿ any_free_foreign : pages from different zones
 - ✿ all_free : pages whose elements are all free
 - ✿ intermediate : pages contain both free and used elements
 - ✿ all_used : pages whose elements are all in use

New Mitigations

- ❖ Changes in heap memory management
 - ❖ Always poison tiny zones, periodically poison larger zones

```
/*
 * Poison the memory before it ends up on the freelist to catch
 * use-after-free and use of uninitialized memory
 *
 * Always poison tiny zones' elements (limit is 0 if -no-zp is set)
 * Also poison larger elements periodically
 */
vm_offset_t      inner_size = zone->elem_size;
uint32_t sample_factor = zp_factor + (((uint32_t)inner_size) >> zp_scale);
if (inner_size <= zp_tiny_zone_limit)
    poison = TRUE;
else if (zp_factor != 0 && sample_counter(&zone->zp_count, sample_factor) == TRUE)
    poison = TRUE;

zp_tiny_zone_limit = (vm_size_t) cpu_info.cache_line_size;
```

New Mitigations

- ❖ Changes in heap memory management
- ❖ Pointer pointing to next free element is xor-ed with the `zp_nopoison_cookie`
- ❖ No more point-to-free-vtable attack

```
/*
 * Insert this element at the head of the free list. We also xor the
 * primary pointer with the zp_nopoison_cookie to make sure a free
 * element does not provide the location of the next free element directly.
 */
*primary          = old_head ^ zp_nopoison_cookie;
```

Agenda

- ❖ New Mitigations in macOS Sierra
- ❖ **Kernel Info Leak**
- ❖ Kernel Code Execution
- ❖ Conclusion

Quick View of IOKit

- ✿ The IOKit is a collection of system frameworks, libraries, tools, and other resources for creating device drivers
- ✿ Extensive user mode API (via IOKit.framework)
 - ✿ Direct memory mapping to user mode
 - ✿ UserClient/External method calls
 - ✿ Notification and messaging

Quick View of IOKit

- ❖ IOKitLib implements non-kernel task access to common IOKit object types -
IORegistryEntry, IOService, IOIterator, etc
 - ❖ IOServiceGetMatchingService
 - ❖ Look up a registered IOService object that matches a matching dictionary
 - ❖ IOServiceOpen
 - ❖ Create a connection to an IOService, return an IOUserClient port
 - ❖ IOConnectSetNotificationPort
 - ❖ Set a port to receive notifications from an IOUserClient object
 - ❖ IOConnectCallMethod
 - ❖ Pass / get data to an IOUserClient object

Info Leak - Overview

- ✿ The vulnerability lies in IOAudioFamily (before 204.4)
- ✿ Source code:
<https://opensource.apple.com/source/IOAudioFamily/IOAudioFamily-204.4/>
- ✿ The vulnerability is an uninitialized heap issue

Vulnerability Analysis

- ❖ IOAudioControlUserClient allows userspace programs to register a notification port and send a notification message to them when certain audio events happen
- ❖ IOConnectSetNotificationPort in userspace will reach IOAudioControlUserClient::registerNotificationPort in the kernel

Vulnerability Analysis

```
IOReturn IOAudioControlUserClient::registerNotificationPort(mach_port_t port, UInt32 type, UInt32 refCon)
{
    ...
    if (notificationMessage == 0) {
        notificationMessage = (IOAudioNotificationMessage *)
            IOMallocAligned(sizeof(IOAudioNotificationMessage), sizeof(IOAudioNotificationMessage *));
        if (!notificationMessage) {
            return kIOReturnNoMemory;
        }
    }
    notificationMessage->messageHeader.msgh_bits = MACH_MSGH_BITS(MACH_MSG_TYPE_COPY_SEND, 0);
    notificationMessage->messageHeader.msgh_size = sizeof(IOAudioNotificationMessage);
    notificationMessage->messageHeader.msgh_remote_port = port;
    notificationMessage->messageHeader.msgh_local_port = MACH_PORT_NULL;
    notificationMessage->messageHeader.msgh_reserved = 0;
    notificationMessage->messageHeader.msgh_id = 0;
    notificationMessage->ref = refCon;
    ...
}

typedef struct _IOAudioNotificationMessage
{
    mach_msg_header_t messageHeader;
    UInt32 type;
    UInt32 ref;
    void * sender;
} IOAudioNotificationMessage;
```

notificationMessage is not zeroed out and the *type* and the *sender* field is not initialized

Vulnerability Analysis

- ✿ IOAudioControlUserClient allocates a notification message struct and sends it to userpspace programs via
IOAudioControlUserClient::sendChangeNotification

Info Leak - Leak to Userland

```
void IOAudioControlUserClient::sendChangeNotification(UInt32 notificationType)
{
    if (notificationMessage) {
        kern_return_t kr;

        notificationMessage->type = notificationType;
        kr = mach_msg_send_from_kernel(&notificationMessage->messageHeader,
notificationMessage->messageHeader.msgh_size);
        if ((kr != MACH_MSG_SUCCESS) && (kr != MACH_SEND_TIMED_OUT)) {
            IOLog("IOAudioControlUserClient: sendRangeChangeNotification() failed - msg_send
returned: %d\n", kr);
        }
    }
}
```

- ❖ The notificationMessage is allocated in kalloc.72 zone
- ❖ The *sender* field offset is 0x38
- ❖ We can read 8-byte data at the offset of 0x38 from any object in kalloc.72

Exploiting The Vulnerability

- ❖ Create many IOAudioControlUserClient objects through IOServiceOpen
- ❖ Register notification ports on such IOAudioControlUserClient objects via IOConnectSetNotificationPort
 - ❖ each IOAudioControlUserClient object will allocate a notificationMessage struct with uninitialized 8 bytes data
- ❖ Trigger a notification event, and receive the notification messages

However...

- ✿ Challenge 1: How to trigger the notification event
- ✿ Challenge 2: How to leak critical information such as KALSR slide

Info Leak - Sandbox Restriction

- ✿ Setting sound volume level can trigger the notification
 - ✿ i.e., set “IOAudioControlValue” via
IORegistryEntrySetCFProperties
- ✿ However, the webcontent sandbox profile does not allow to set this property

coreaudiod

- ❖ Who is allowed to set sound volume level?
- ❖ By greping IOAudioControlValue, we found /usr/sbin/coreaudiod can set sound volume level
- ❖ coreaudiod is also responsible for Mach Service com.apple.audio.coreaudiod

coreaudiod

- ❖ Webprocess can also talk with com.apple.audio.coreaudiod service
- ❖ As a result, we can instruct coreaudiod to trigger the notification event

```
;; Various services required by AppKit and other frameworks
(allow mach-lookup
  (global-name "com.apple.DiskArbitration.diskarbitrationsd")
  (global-name "com.apple.FileCoordination")
  (global-name "com.apple.FontObjectsServer")

  (global-name "com.apple.PowerManagement.control")
  (global-name "com.apple.SystemConfiguration.configd")
  (global-name "com.apple.SystemConfiguration.PPPController")
  (global-name "com.apple.audio.SystemSoundServer-OSX")
  (global-name "com.apple.audio.VDCAssistant")
  (global-name "com.apple.audio.audiohal")
  (global-name "com.apple.audio.coreaudiod")
```

How to leak KALSR slide

- ✿ We can read 8-byte data at the offset of 0x38 from a freed object in kalloc.72
- ✿ How to make the 8-byte data meaningful?

Info Leak - Leak OSSerialize Object

```
class OSSerialize : public OSObject
{
    ...
private:
    char          * data;
    unsigned int   length;
    unsigned int   capacity;
    unsigned int   capacityIncrement;
    OSArray * tags;
    bool          binary;
    bool          endCollection;
    Editor editor; // Editor is highlighted with a red box
    void * editRef;
    ...
}

typedef const OSMetaClassBase * (*Editor)(  

    void* reference,  

    OSSerialize* s,  

    OSCollection* container,  

    const OSSymbol* name,  

    const OSMetaClassBase* value);
```

- ✿ The OSSerialize objects locates in the same zone as notificationMessage object
- ✿ The member at offset 0x38 is a function pointer
- ✿ Use separate thread to spray OSSerialize objects and leak several times to ensure stability

Agenda

- ❖ New Mitigations in macOS Sierra
- ❖ Kernel Info Leak
- ❖ **Kernel Code Execution**
- ❖ Conclusion

Kernel Code Execution - Overview

- ✿ Accelerator is one of the devices that we can directly access in the webcontent sandbox
- ✿ The vulnerability lies in AMDRadeonX~~x~~000.kext
(x may vary on different platforms)
- ✿ The vulnerability is an uninitialized stack variable issue and the exploit is quite straight-forward

Kernel Code Execution - Preparation

- ❖ Two different userclients
 - ❖ AMDSIGLContext (type 1)
 - ❖ AMDAccelSharedUserClient (type 6)
- ❖ Connect these two userclients
- ❖ One method call
 - ❖ AMDRadeonX4000_AMDSIGLContext::surfaceCopy (selector 0x201)

Kernel Code Execution - surfaceCopy()

- ⌘ Local variable is not initialized to 0
- ⌘ User can control the resource id to lookup
- ⌘ LookupResource() not always initialize the local variable
- ⌘ A vcall is invoked blindly on these local variables

```
IOAccelShared2::lookupResource(v5, a2[2], &v51);
IOAccelShared2::lookupResource(v5, a2[1], (void **)&v50);
v6 = -536870206;
if ( !v51 || !v50 )
    goto LABEL_42;
v12 = (*(__int64 (**)(void))(*(_QWORD *)v51 + 368LL))();
v13 = (*(__int64 (**)(void))(*(_QWORD *)v50 + 368LL))();
```

Kernel Code Execution - Control the Stack

- ✿ The uninitialized value on kernel stack is random so we need to control the stack
- ✿ Luckily, a function (selector 7333) in the AGPM userclient come to rescue
- ✿ We are able to copy at most 4096 bytes of controlled, non-zero data onto kernel stack

```
case 7333:  
    kprintf("kAGPMSetPlimit plimit = %llu type = %s\n", *a2->scalarInput, a2->structureInput);  
    v16 = (char *)&v22 - ((a2->structureInputSize + 1 + 15LL) & 0xFFFFFFFFFFFFFFF0LL);  
    strncpy(v16, (const char *)a2->structureInput, a2->structureInputSize);
```

Kernel Code Execution - Value on Stack

- ✿ What value should we “initialize” for the uninitialized stack variable
- ✿ A pointer pointing to an object we fake on the heap
- ✿ A pointer pointing to a real object on the heap
 - ✿ It is hard to find such a candidate
 - ✿ You may have a try ;-)

Kernel Code Execution - Value on Stack

- ✿ How can we know where our fake objects locates?
- ✿ Option - 1
 - ✿ Reuse that leak to get heap address info
 - ✿ Hunt for appropriate object - HARD
 - ✿ Too many free-refill operations - UNSTABLE

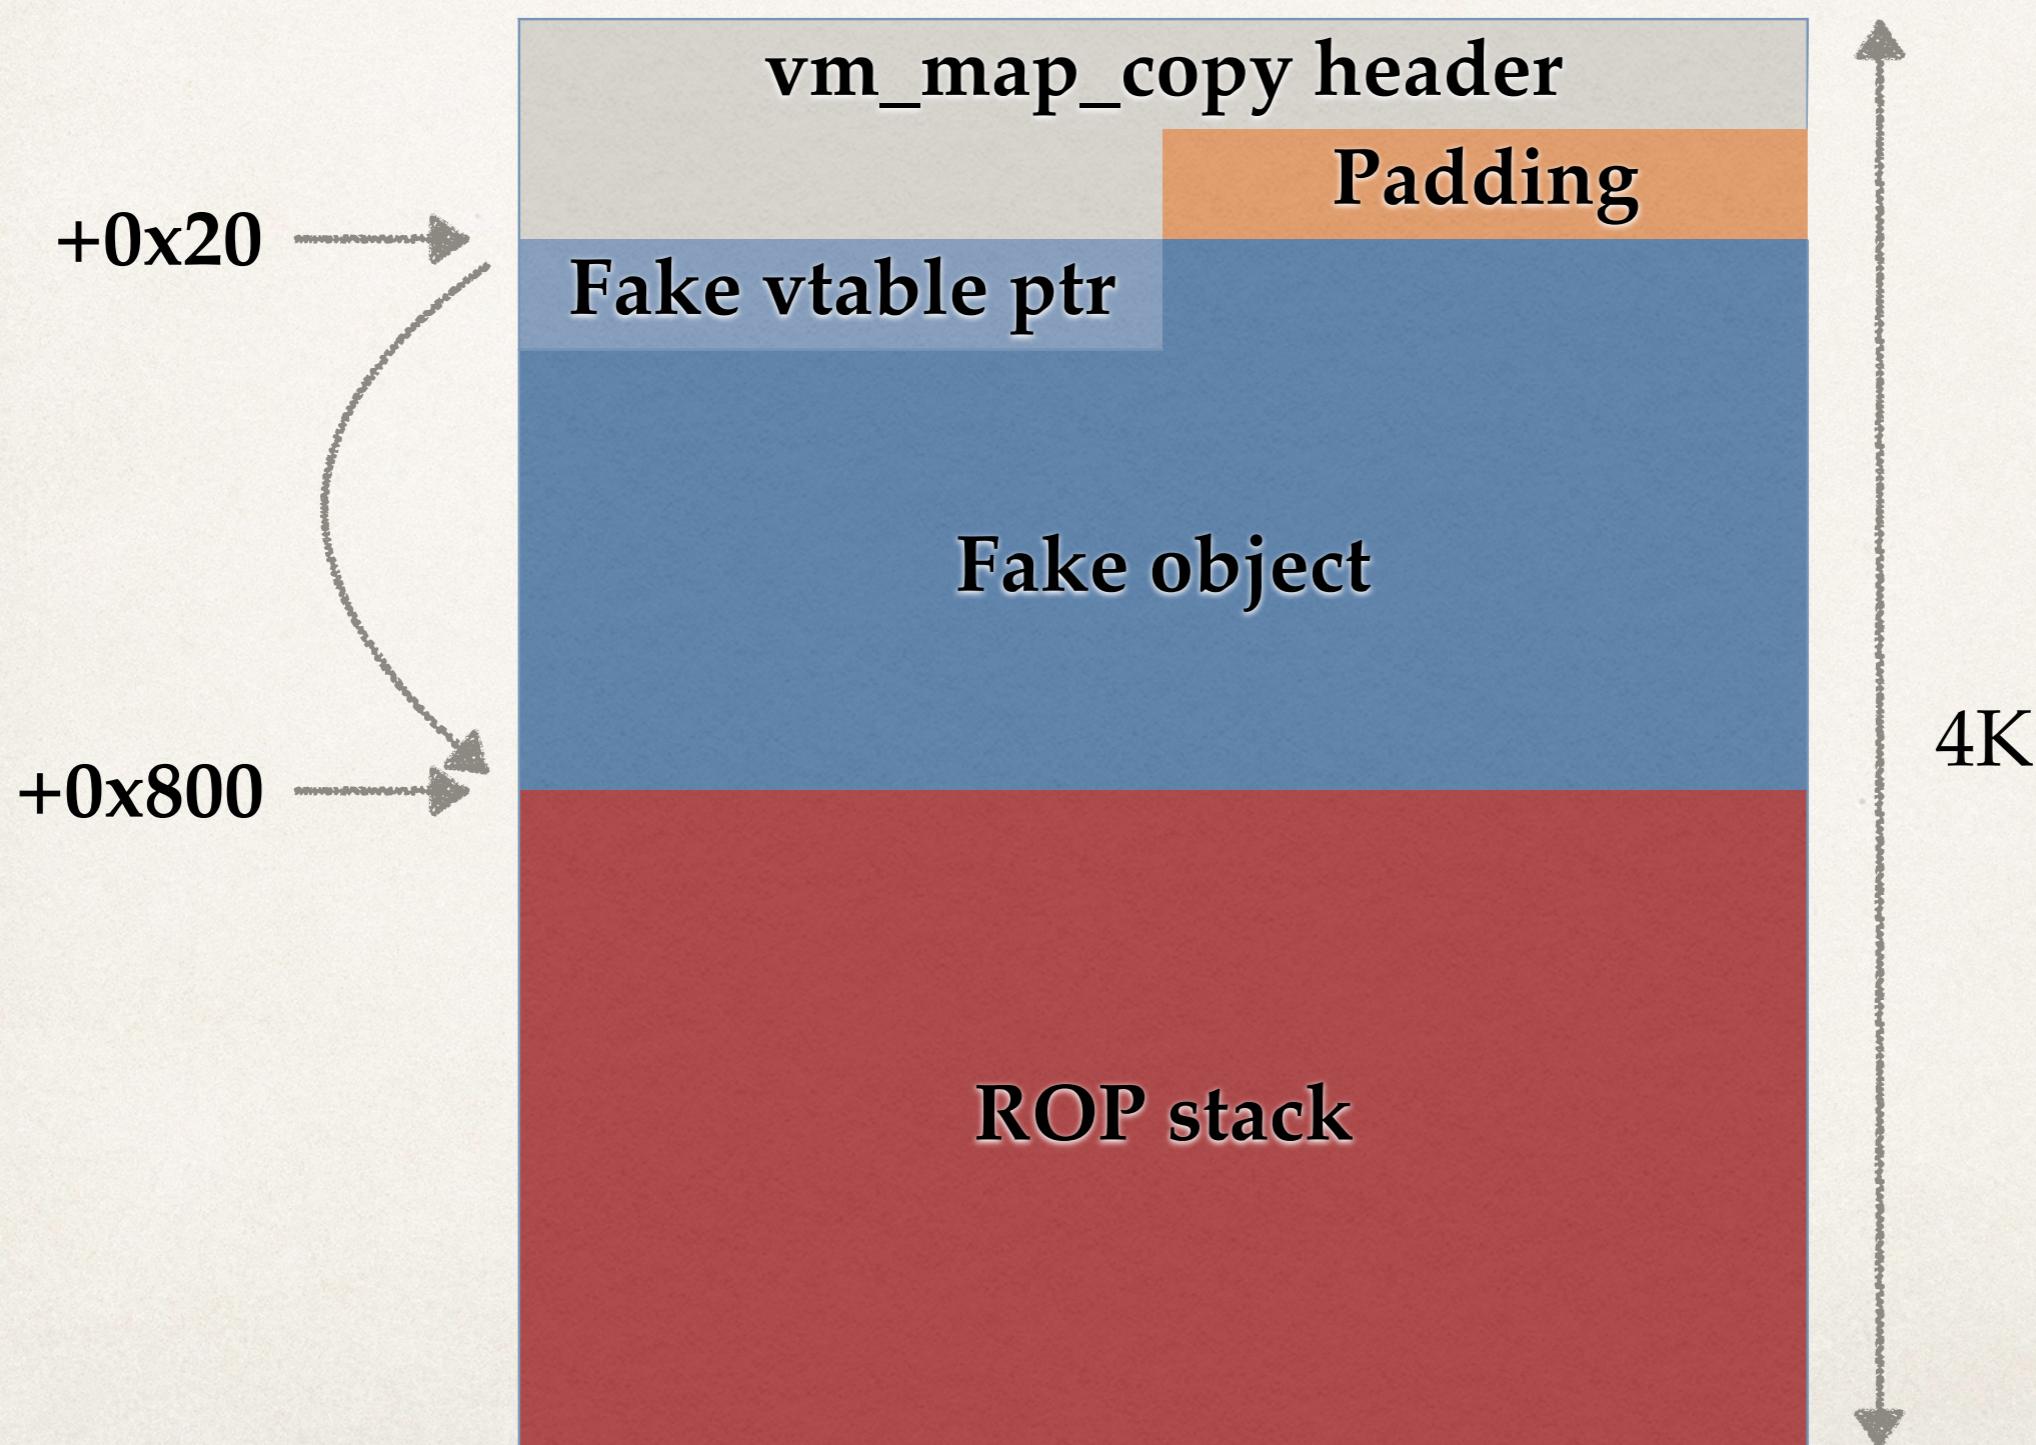
Kernel Code Execution - Value on Stack

- ✿ How can we know where our fake objects locates?
- ✿ Option - 2
 - ✿ Spray several GBs of data in the kernel
 - ✿ Heap randomization weak in the kernel
 - ✿ User-controlled data locates at fixed address

Kernel Code Execution - Spray Fake Objects

- ✿ Spray with `vm_map_copy`
 - ✿ Fast and small side effects on heap
 - ✿ Arbitrary controlled size (up to 4K) and contents
 - ✿ Uncontrolled beginning 0x18 bytes - NOT matter
 - ✿ Fake object contains a fake vtable and a fake stack

Kernel Code Execution - Spray Fake Objects



Kernel Code Execution - ROP Chain

- ✿ Save registers
- ✿ Pivot Stack
- ✿ Disable SMEP & SMAP
- ✿ Return to userland SHELLCODE
- ✿ Re-enable SMEP & SMAP
- ✿ Call `_thread_exception_return()` to exit

Kernel Code Execution - Disable SMEP&SMAP

- ✿ Bits stored in CR4 register indicate the states of CPU
- ✿ Clear 21st to disable SMEP and 22nd for SMAP
- ✿ gadgets:
 - ✿ read CR4: `mov rax, cr4 ,..., ret`
 - ✿ Write CR4: `mov cr4, rax ,..., ret`

Kernel Code Execution - Special Notes

- ✿ The kernel stack used when entering kernel mode varies from time to time
 - ✿ call sprayStack() several times to cover more stacks
- ✿ Do not forget to unlock the locks locked in AMDSIGLContext::surfaceCopy()

Agenda

- ❖ New Mitigations in macOS Sierra
- ❖ Kernel Info Leak
- ❖ Kernel Code Execution
- ❖ Conclusion

Conclusion

- ❖ New mitigations make exploit more challenging and encourage us to discover new exploit techniques
- ❖ Keep initialization in our mind when coding

Thank You For Your Attention
Q&A

